

## This is the k-nearest neighbors workbook for ECE 239AS Assignment #2

Please follow the notebook linearly to implement k-nearest neighbors.

Please print out the workbook entirely when completed.

We thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class (cs231n.stanford.edu). These are the functions in the cs231n folders and code in the jupyter notebook to preprocess and show the images. The classifiers used are based off of code prepared for CS 231n as well.

The goal of this workbook is to give you experience with the data, training and evaluating a simple classifier, k-fold cross validation, and as a Python refresher.

### Import the appropriate libraries

```
In [1]: 1 import numpy as np # for doing most of our calculations
2 import matplotlib.pyplot as plt # for plotting
3 from cs231n.data_utils import load_CIFAR10 # function to load the CIFAR-10 dataset.
4
5 # Load matplotlib images inline
6 %matplotlib inline
7
8 # These are important for reloading any code you write in external .py files.
9 # see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
10 %load_ext autoreload
11 %autoreload 2
```

```
In [3]: 1 # Set the path to the CIFAR-10 data
2 cifar10_dir = 'cifar-10-batches-py'
3 X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)
4
5 # As a sanity check, we print out the size of the training and test data.
6 print('Training data shape: ', X_train.shape)
7 print('Training labels shape: ', y_train.shape)
8 print('Test data shape: ', X_test.shape)
9 print('Test labels shape: ', y_test.shape)
```

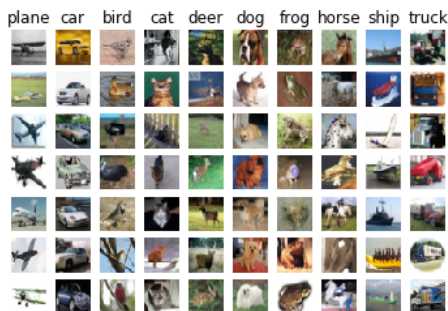
Training data shape: (50000, 32, 32, 3)

Training labels shape: (50000,)

Test data shape: (10000, 32, 32, 3)

Test labels shape: (10000,)

```
In [4]: 1 # Visualize some examples from the dataset.
2 # We show a few examples of training images from each class.
3 classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
4 num_classes = len(classes)
5 samples_per_class = 7
6 for y, cls in enumerate(classes):
7     idxs = np.flatnonzero(y_train == y)
8     idxs = np.random.choice(idxs, samples_per_class, replace=False)
9     for i, idx in enumerate(idxs):
10         plt_idx = i * num_classes + y + 1
11         plt.subplot(samples_per_class, num_classes, plt_idx)
12         plt.imshow(X_train[idx].astype('uint8'))
13         plt.axis('off')
14         if i == 0:
15             plt.title(cls)
16 plt.show()
```



```
In [5]: 1 # Subsample the data for more efficient code execution in this exercise
2 num_training = 5000
3 mask = list(range(num_training))
4 X_train = X_train[mask]
5 y_train = y_train[mask]
6
7 num_test = 500
8 mask = list(range(num_test))
9 X_test = X_test[mask]
10 y_test = y_test[mask]
11
12 # Reshape the image data into rows
13 X_train = np.reshape(X_train, (X_train.shape[0], -1))
14 X_test = np.reshape(X_test, (X_test.shape[0], -1))
15 print(X_train.shape, X_test.shape)

(5000, 3072) (500, 3072)
```

## K-nearest neighbors

In the following cells, you will build a KNN classifier and choose hyperparameters via k-fold cross-validation.

```
In [7]: 1 # Import the KNN class
2
3 from nn1 import KNN
```

```
In [8]: 1 # Declare an instance of the knn class.
2 knn = KNN()
3
4 # Train the classifier.
5 # We have implemented the training of the KNN classifier.
6 # Look at the train function in the KNN class to see what this does.
7 knn.train(X=X_train, y=y_train)
```

## Questions

- (1) Describe what is going on in the function `knn.train()`.
- (2) What are the pros and cons of this training step?

## Answers

- (1) `knn.train()` function reads and stores the training data, including images and their corresponding labels.
- (2) Pros: the implement is very simple. Cons: it's very time consuming when predicting test data.

## KNN prediction

In the following sections, you will implement the functions to calculate the distances of test points to training points, and from this information, predict the class of the KNN.

```
In [9]: 1 # Implement the function compute_distances() in the KNN class.
2 # Do not worry about the input 'norm' for now; use the default definition of the norm
3 # in the code, which is the 2-norm.
4 # You should only have to fill out the clearly marked sections.
5
6 import time
7 time_start = time.time()
8
9 dists_L2 = knn.compute_distances(X=X_test)
10
11 print('Time to run code: {}'.format(time.time()-time_start))
12 print('Frobenius norm of L2 distances: {}'.format(np.linalg.norm(dists_L2, 'fro')))
```

```
Time to run code: 31.2187762260437
Frobenius norm of L2 distances: 7906696.077040902
```

### Really slow code

Note: This probably took a while. This is because we use two for loops. We could increase the speed via vectorization, removing the for loops.

If you implemented this correctly, evaluating `np.linalg.norm(dists_L2, 'fro')` should return: ~7906696

## KNN vectorization

The above code took far too long to run. If we wanted to optimize hyperparameters, it would be time-expensive. Thus, we will speed up the code by vectorizing it, removing the for loops.

```
In [12]: 1 # Implement the function compute_L2_distances_vectorized() in the KNN class.
2 # In this function, you ought to achieve the same L2 distance but WITHOUT any for loops.
3 # Note, this is SPECIFIC for the L2 norm.
4
5 time_start = time.time()
6 dists_L2_vectorized = knn.compute_L2_distances_vectorized(X=X_test)
7 print('Time to run code: {}'.format(time.time()-time_start))
8 print('Difference in L2 distances between your KNN implementations (should be 0): {}'.format(np.linalg.norm(dists_L2 -
```

Time to run code: 0.42450380325317383

Difference in L2 distances between your KNN implementations (should be 0): 0.0

### Speedup

Depending on your computer speed, you should see a 10-100x speed up from vectorization. On our computer, the vectorized form took 0.36 seconds while the naive implementation took 38.3 seconds.

## Implementing the prediction

Now that we have functions to calculate the distances from a test point to given training points, we now implement the function that will predict the test point labels.

```
In [14]: 1 # Implement the function predict_labels in the KNN class.
2 # Calculate the training error (num_incorrect / total_samples)
3 # from running knn.predict_labels with k=1
4
5 error = 1
6
7 # ===== #
8 # YOUR CODE HERE:
9 # Calculate the error rate by calling predict_labels on the test
10 # data with k = 1. Store the error rate in the variable error.
11 # ===== #
12
13 pred = knn.predict_labels(dists_L2_vectorized)
14 count = 0
15 for i in range(len(y_test)):
16     cur_pred = pred[i]
17     cur_y = y_test[i]
18     if cur_pred != cur_y:
19         count += 1
20 error = count / y_test.shape[0]
21
22 # ===== #
23 # END YOUR CODE HERE
24 # ===== #
25
26 print(error)
```

0.726

If you implemented this correctly, the error should be: 0.726.

This means that the k-nearest neighbors classifier is right 27.4% of the time, which is not great, considering that chance levels are 10%.

## Optimizing KNN hyperparameters

In this section, we'll take the KNN classifier that you have constructed and perform cross-validation to choose a best value of  $k$ , as well as a best choice of norm.

### Create training and validation folds

First, we will create the training and validation folds for use in k-fold cross validation.

```

In [15]: 1 # Create the dataset folds for cross-validation.
          2 num_folds = 5
          3
          4 X_train_folds = []
          5 y_train_folds = []
          6
          7 # ===== #
          8 # YOUR CODE HERE:
          9 # Split the training data into num_folds (i.e., 5) folds.
         10 # X_train_folds is a list, where X_train_folds[i] contains the
         11 # data points in fold i.
         12 # y_train_folds is also a list, where y_train_folds[i] contains
         13 # the corresponding labels for the data in X_train_folds[i]
         14 # ===== #
         15
         16 idx = np.arange(num_training)
         17 np.random.shuffle(idx)
         18
         19 X_train_shuffle = X_train[idx[:]]
         20 y_train_shuffle = y_train[idx[:]]
         21 X_train_folds = np.array_split(X_train_shuffle, num_folds)
         22 y_train_folds = np.array_split(y_train_shuffle.reshape(-1, 1), num_folds)
         23
         24 # ===== #
         25 # END YOUR CODE HERE
         26 # ===== #
         27
         28

```


### Optimizing the number of nearest neighbors hyperparameter.

In this section, we select different numbers of nearest neighbors and assess which one has the lowest k-fold cross validation error.

```

In [16]: 1 time_start = time.time()
2
3 ks = [1, 2, 3, 5, 7, 10, 15, 20, 25, 30]
4
5 # ===== #
6 # YOUR CODE HERE:
7 # Calculate the cross-validation error for each k in ks, testing
8 # the trained model on each of the 5 folds. Average these errors
9 # together and make a plot of k vs. cross-validation error. Since
10 # we are assuming L2 distance here, please use the vectorized code!
11 # Otherwise, you might be waiting a long time.
12 # ===== #
13
14 error_avg = []
15 for k in ks:
16     error_sum = 0
17     for i in range(num_folds):
18         # select train and test folds
19         X_cur_validation = X_train_folds[i]
20         y_cur_validation = y_train_folds[i]
21         X_cur_train = np.concatenate(X_train_folds[:i] + X_train_folds[i + 1:])
22         y_cur_train = np.concatenate(y_train_folds[:i] + y_train_folds[i + 1:])
23
24         # create knn classifier
25         knn.train(X_cur_train, y_cur_train[:,0])
26         cur_dists = knn.compute_L2_distances_vectorized(X_cur_validation)
27         cur_preds = knn.predict_labels(cur_dists, k=k)
28
29         # calculate error
30         count = 0
31         for j in range(len(cur_preds)):
32             cur_pred = cur_preds[j]
33             cur_y = y_cur_validation[j]
34             if cur_pred != cur_y:
35                 count += 1
36         error_sum += count / X_cur_validation.shape[0]
37     error_avg.append(error_sum / num_folds)
38
39 # plot
40 plt.plot(ks, error_avg)
41 plt.title('Average Cross Validation Error vs k')
42 plt.ylabel('Average Error')
43 plt.xlabel('k')
44 plt.show()
45
46 best_error = min(error_avg)
47 best_k_idx = error_avg.index(best_error)
48 print(best_error, ks[best_k_idx])
49
50 # ===== #
51 # END YOUR CODE HERE
52 # ===== #
53
54 print('Computation time: %.2f'%(time.time()-time_start))

```



0.7192000000000001 7  
Computation time: 152.55

## Questions:

- (1) What value of  $k$  is best amongst the tested  $k$ 's?
- (2) What is the cross-validation error for this value of  $k$ ?

## Answers:

- (1)  $k = 7$

(2) The average error is 0.719


### **Optimizing the norm**

Next, we test three different norms (the 1, 2, and infinity norms) and see which distance metric results in the best cross-validation performance.

```

In [17]: 1 time_start = time.time()
2
3 L1_norm = lambda x: np.linalg.norm(x, ord=1)
4 L2_norm = lambda x: np.linalg.norm(x, ord=2)
5 Linf_norm = lambda x: np.linalg.norm(x, ord= np.inf)
6 norms = [L1_norm, L2_norm, Linf_norm]
7
8 # ===== #
9 # YOUR CODE HERE:
10 # Calculate the cross-validation error for each norm in norms, testing
11 # the trained model on each of the 5 folds. Average these errors
12 # together and make a plot of the norm used vs the cross-validation error
13 # Use the best cross-validation k from the previous part.
14 #
15 # Feel free to use the compute_distances function. We're testing just
16 # three norms, but be advised that this could still take some time.
17 # You're welcome to write a vectorized form of the L1- and Linf- norms
18 # to speed this up, but it is not necessary.
19 # ===== #
20
21 best_k = ks[best_k_idx]
22 norm_error_sum = 0
23 norm_error_avg = []
24 for n in norms:
25     for j in range(num_folds):
26         # select train and test folds
27         X_cur_validation = X_train_folds[i]
28         y_cur_validation = y_train_folds[i]
29         X_cur_train = np.concatenate(X_train_folds[:i] + X_train_folds[i + 1:])
30         y_cur_train = np.concatenate(y_train_folds[:i] + y_train_folds[i + 1:])
31
32         # create knn classifier
33         norm_knn = KNN()
34         norm_knn.train(X_cur_train, y_cur_train[:,0])
35         norm_dists = norm_knn.compute_distances(X_cur_validation, norm=n)
36         norm_preds = norm_knn.predict_labels(norm_dists, k=best_k)
37
38         # calculate error
39         count = 0
40         for j in range(len(norm_preds)):
41             cur_pred = norm_preds[j]
42             cur_y = y_cur_validation[j]
43             if cur_pred != cur_y:
44                 count += 1
45         norm_error_sum += count / X_cur_validation.shape[0]
46     norm_error_avg.append(norm_error_sum / num_folds)
47
48 # plot
49 plt.plot([1, 2, 3], norm_error_avg)
50 plt.title('Average Cross Validation Error vs norm')
51 plt.ylabel('Average Error')
52 plt.xlabel('norm')
53 plt.show()
54
55 best_error = min(norm_error_avg)
56 best_norm_idx = norm_error_avg.index(best_error)
57 print(best_error, best_norm_idx)
58
59 # ===== #
60 # END YOUR CODE HERE
61 # ===== #
62 print('Computation time: %.2f'%(time.time()-time_start))

```



0.67 0  
Computation time: 604.90

## Questions:

(1) What norm has the best cross-validation error?

(2) What is the cross-validation error for your given norm and k?

## Answers:

(1) L1 norm.

(2) 0.67

## Evaluating the model on the testing dataset.

Now, given the optimal  $k$  and norm you found in earlier parts, evaluate the testing error of the k-nearest neighbors model.

```
In [19]: 1 error = 1
2
3 # ===== #
4 # YOUR CODE HERE:
5 # Evaluate the testing error of the k-nearest neighbors classifier
6 # for your optimal hyperparameters found by 5-fold cross-validation.
7 # ===== #
8
9 knn = KNN()
10 knn.train(X=X_train, y=y_train)
11 cur_dists = knn.compute_distances(X=X_test, norm=L1_norm)
12 cur_preds = knn.predict_labels(cur_dists, k=7)
13 count_error = 0
14 for i in range(len(y_test)):
15     cur_pred = cur_preds[i]
16     cur_y = y_test[i]
17     if cur_pred != cur_y:
18         count_error += 1
19 error = count_error / y_test.shape[0]
20
21 # ===== #
22 # END YOUR CODE HERE
23 # ===== #
24
25 print('Error rate achieved: {}'.format(error))
```

Error rate achieved: 0.704

## Question:

How much did your error improve by cross-validation over naively choosing  $k=1$  and using the L2-norm?

## Answer:

0.022

```
In [ ]: 1
```



```

1 import numpy as np
2 import pdb
3
4 """
5 This code was based off of code from cs231n at Stanford University, and
6 modified for ece239as at UCLA.
7 """
8 class KNN(object):
9
10     def __init__(self):
11         pass
12
13     def train(self, X, y):
14         """
15         Inputs:
16         - X is a numpy array of size (num_examples, D)
17         - y is a numpy array of size (num_examples, )
18         """
19         self.X_train = X
20         self.y_train = y
21
22     def compute_distances(self, X, norm=None):
23         """
24         Compute the distance between each test point in X and each training point
25         in self.X_train.
26
27         Inputs:
28         - X: A numpy array of shape (num_test, D) containing test data.
29         - norm: the function with which the norm is taken.
30
31         Returns:
32         - dists: A numpy array of shape (num_test, num_train) where dists[i, j]
33           is the Euclidean distance between the ith test point and the jth
34           training point.
35         """
36         if norm is None:
37             norm = lambda x: np.sqrt(np.sum(x**2))
38             #norm = 2
39
40         num_test = X.shape[0]
41         num_train = self.X_train.shape[0]
42         dists = np.zeros((num_test, num_train))
43         for i in np.arange(num_test):
44             for j in np.arange(num_train):
45                 # ===== #
46                 # YOUR CODE HERE:
47                 #   Compute the distance between the ith test point and the jth
48                 #   training point using norm(), and store the result in dists[i, j].
49                 # ===== #
50
51                 dists[i, j] = norm(X[i] - self.X_train[j])
52
53                 # ===== #
54                 # END YOUR CODE HERE
55                 # ===== #
56
57         return dists
58
59     def compute_L2_distances_vectorized(self, X):
60         """
61         Compute the distance between each test point in X and each training point
62         in self.X_train WITHOUT using any for loops.
63
64         Inputs:
65         - X: A numpy array of shape (num_test, D) containing test data.
66
67         Returns:
68         - dists: A numpy array of shape (num_test, num_train) where dists[i, j]
69           is the Euclidean distance between the ith test point and the jth
70           training point.
71         """
72         num_test = X.shape[0]
73         num_train = self.X_train.shape[0]
74         dists = np.zeros((num_test, num_train))
75
76         # ===== #
77         # YOUR CODE HERE:
78         #   Compute the L2 distance between the ith test point and the jth
79         #   training point and store the result in dists[i, j]. You may
80

```

```

81 # NOT use a for loop (or list comprehension). You may only use
82 # numpy operations.
83 #
84 # HINT: use broadcasting. If you have a shape (N,1) array and
85 # a shape (M,) array, adding them together produces a shape (N, M)
86 # array.
87 # ===== #
88
89 test_square = np.sum(X**2, axis=1).reshape((num_test, 1))
90 train_square = np.sum(self.X_train**2, axis=1).reshape((1, num_train))
91 test_dot_train = np.dot(X, self.X_train.T)
92
93 dists = np.sqrt(dists + test_square + train_square - 2 * test_dot_train)
94
95 # ===== #
96 # END YOUR CODE HERE
97 # ===== #
98
99 return dists
100
101
102 def predict_labels(self, dists, k=1):
103     """
104     Given a matrix of distances between test points and training points,
105     predict a label for each test point.
106
107     Inputs:
108     - dists: A numpy array of shape (num_test, num_train) where dists[i, j]
109     gives the distance between the ith test point and the jth training
110     point.
111
112     Returns:
113     - y: A numpy array of shape (num_test,) containing predicted labels for
114     the test data, where y[i] is the predicted label for the test point X[i].
115     """
116     num_test = dists.shape[0]
117     y_pred = np.zeros(num_test)
118     for i in np.arange(num_test):
119         # A list of length k storing the labels of the k nearest neighbors to
120         # the ith test point.
121         closest_y = []
122         # YOUR CODE HERE:
123         # Use the distances to calculate and then store the labels of
124         # the k-nearest neighbors to the ith test point. The function
125         # numpy.argsort may be useful.
126         #
127         # After doing this, find the most common label of the k-nearest
128         # neighbors. Store the predicted label of the ith training example
129         # as y_pred[i]. Break ties by choosing the smaller label.
130         # ===== #
131
132         idx = sorted(range(dists.shape[1]), key = lambda j : dists[i,:][j])[:k]
133         closest_y = [self.y_train[i] for i in idx]
134         y_pred[i] = max(set(closest_y),key = closest_y.count)
135
136         # ===== #
137         # END YOUR CODE HERE
138         # ===== #
139
140     return y_pred
141

```

## This is the svm workbook for ECE 239AS Assignment #2

Please follow the notebook linearly to implement a linear support vector machine.

Please print out the workbook entirely when completed.

We thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class (cs231n.stanford.edu). These are the functions in the cs231n folders and includes code to preprocess and show the images. The classifiers used are based off of code prepared for CS 231n as well.

The goal of this workbook is to give you experience with training an SVM classifier via gradient descent.

### Importing libraries and data setup

```
In [1]: 1 import numpy as np # for doing most of our calculations
2 import matplotlib.pyplot as plt# for plotting
3 from cs231n.data_utils import load_CIFAR10 # function to load the CIFAR-10 dataset.
4 import pdb
5
6 # Load matplotlib images inline
7 %matplotlib inline
8
9 # These are important for reloading any code you write in external .py files.
10 # see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
11 %load_ext autoreload
12 %autoreload 2
```

```
In [2]: 1 # Set the path to the CIFAR-10 data
2 cifar10_dir = 'cifar-10-batches-py'
3 X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)
4
5 # As a sanity check, we print out the size of the training and test data.
6 print('Training data shape: ', X_train.shape)
7 print('Training labels shape: ', y_train.shape)
8 print('Test data shape: ', X_test.shape)
9 print('Test labels shape: ', y_test.shape)
```

```
Training data shape: (50000, 32, 32, 3)
Training labels shape: (50000,)
Test data shape: (10000, 32, 32, 3)
Test labels shape: (10000,)
```

```
In [3]: 1 # Visualize some examples from the dataset.
2 # We show a few examples of training images from each class.
3 classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
4 num_classes = len(classes)
5 samples_per_class = 7
6 for y, cls in enumerate(classes):
7     idxs = np.flatnonzero(y_train == y)
8     idxs = np.random.choice(idxs, samples_per_class, replace=False)
9     for i, idx in enumerate(idxs):
10         plt_idx = i * num_classes + y + 1
11         plt.subplot(samples_per_class, num_classes, plt_idx)
12         plt.imshow(X_train[idx].astype('uint8'))
13         plt.axis('off')
14         if i == 0:
15             plt.title(cls)
16 plt.show()
```



```
In [4]: 1 # Split the data into train, val, and test sets. In addition we will
2 # create a small development set as a subset of the training data;
3 # we can use this for development so our code runs faster.
4 num_training = 49000
5 num_validation = 1000
6 num_test = 1000
7 num_dev = 500
8
9 # Our validation set will be num_validation points from the original
10 # training set.
11 mask = range(num_training, num_training + num_validation)
12 X_val = X_train[mask]
13 y_val = y_train[mask]
14
15 # Our training set will be the first num_train points from the original
16 # training set.
17 mask = range(num_training)
18 X_train = X_train[mask]
19 y_train = y_train[mask]
20
21 # We will also make a development set, which is a small subset of
22 # the training set.
23 mask = np.random.choice(num_training, num_dev, replace=False)
24 X_dev = X_train[mask]
25 y_dev = y_train[mask]
26
27 # We use the first num_test points of the original test set as our
28 # test set.
29 mask = range(num_test)
30 X_test = X_test[mask]
31 y_test = y_test[mask]
32
33 print('Train data shape: ', X_train.shape)
34 print('Train labels shape: ', y_train.shape)
35 print('Validation data shape: ', X_val.shape)
36 print('Validation labels shape: ', y_val.shape)
37 print('Test data shape: ', X_test.shape)
38 print('Test labels shape: ', y_test.shape)
39 print('Dev data shape: ', X_dev.shape)
40 print('Dev labels shape: ', y_dev.shape)
```


```
Train data shape: (49000, 32, 32, 3)
Train labels shape: (49000,)
Validation data shape: (1000, 32, 32, 3)
Validation labels shape: (1000,)
Test data shape: (1000, 32, 32, 3)
Test labels shape: (1000,)
Dev data shape: (500, 32, 32, 3)
Dev labels shape: (500,)
```

```
In [5]: 1 # Preprocessing: reshape the image data into rows
2 X_train = np.reshape(X_train, (X_train.shape[0], -1))
3 X_val = np.reshape(X_val, (X_val.shape[0], -1))
4 X_test = np.reshape(X_test, (X_test.shape[0], -1))
5 X_dev = np.reshape(X_dev, (X_dev.shape[0], -1))
6
7 # As a sanity check, print out the shapes of the data
8 print('Training data shape: ', X_train.shape)
9 print('Validation data shape: ', X_val.shape)
10 print('Test data shape: ', X_test.shape)
11 print('dev data shape: ', X_dev.shape)
```

```
Training data shape: (49000, 3072)
Validation data shape: (1000, 3072)
Test data shape: (1000, 3072)
dev data shape: (500, 3072)
```

```
In [6]: 1 # Preprocessing: subtract the mean image
2 # first: compute the image mean based on the training data
3 mean_image = np.mean(X_train, axis=0)
4 print(mean_image[:10]) # print a few of the elements
5 plt.figure(figsize=(4,4))
6 plt.imshow(mean_image.reshape((32,32,3)).astype('uint8')) # visualize the mean image
7 plt.show()
```

```
[130.64189796 135.98173469 132.47391837 130.05569388 135.34804082
131.75402041 130.96055102 136.14328571 132.47636735 131.48467347]
```



```
In [7]: 1 # second: subtract the mean image from train and test data
2 X_train -= mean_image
3 X_val -= mean_image
4 X_test -= mean_image
5 X_dev -= mean_image
```

```
In [8]: 1 # third: append the bias dimension of ones (i.e. bias trick) so that our SVM
2 # only has to worry about optimizing a single weight matrix W.
3 X_train = np.hstack([X_train, np.ones((X_train.shape[0], 1))])
4 X_val = np.hstack([X_val, np.ones((X_val.shape[0], 1))])
5 X_test = np.hstack([X_test, np.ones((X_test.shape[0], 1))])
6 X_dev = np.hstack([X_dev, np.ones((X_dev.shape[0], 1))])
7
8 print(X_train.shape, X_val.shape, X_test.shape, X_dev.shape)
```

```
(49000, 3073) (1000, 3073) (1000, 3073) (500, 3073)
```

## Question:

(1) For the SVM, we perform mean-subtraction on the data. However, for the KNN notebook, we did not. Why?

## Answer:

(1) In SVM, we perform normalization in order to scale all data to have similar influence on distance matrix. For KNN, the normalization is on all data and it will reduce the distances for all. It will not influence the result.

## Training an SVM

The following cells will take you through building an SVM. You will implement its loss function, then subsequently train it with gradient descent. Finally, you will choose the learning rate of gradient descent to optimize its classification performance.

```
In [9]: 1 from nndl.svm import SVM
```

```
In [10]: 1 # Declare an instance of the SVM class.
2 # Weights are initialized to a random value.
3 # Note, to keep people's initial solutions consistent, we are going to use a random seed.
4
5 np.random.seed(1)
6
7 num_classes = len(np.unique(y_train))
8 num_features = X_train.shape[1]
9
10 svm = SVM(dims=[num_classes, num_features])
```

### SVM loss

```
In [11]: 1  ## Implement the loss function for in the SVM class(nndl/svm.py), svm.loss()
2
3  loss = svm.loss(X_train, y_train)
4  print('The training set loss is {}'.format(loss))
5
6  # If you implemented the loss correctly, it should be 15569.98
```

The training set loss is 15569.977915410236.

### SVM gradient

```
In [12]: 1  ## Calculate the gradient of the SVM class.
2  # For convenience, we'll write one function that computes the loss
3  # and gradient together. Please modify svm.loss_and_grad(X, y).
4  # You may copy and paste your loss code from svm.loss() here, and then
5  # use the appropriate intermediate values to calculate the gradient.
6
7  loss, grad = svm.loss_and_grad(X_dev, y_dev)
8
9  # Compare your gradient to a numerical gradient check.
10 # You should see relative gradient errors on the order of 1e-07 or less if you implemented the gradient correctly.
11 svm.grad_check_sparse(X_dev, y_dev, grad)
```

```
numerical: -3.949738 analytic: -3.949738, relative error: 2.893617e-08
numerical: 6.563685 analytic: 6.563686, relative error: 2.521352e-08
numerical: -0.935538 analytic: -0.935538, relative error: 1.732962e-08
numerical: 9.288176 analytic: 9.288176, relative error: 3.749327e-10
numerical: 4.790768 analytic: 4.790767, relative error: 6.218272e-08
numerical: -0.886145 analytic: -0.886144, relative error: 2.204496e-07
numerical: 12.689140 analytic: 12.689141, relative error: 3.754695e-09
numerical: -13.026562 analytic: -13.026562, relative error: 1.101139e-09
numerical: 6.809312 analytic: 6.809311, relative error: 1.982750e-08
numerical: -25.761661 analytic: -25.761661, relative error: 1.366211e-09
```

## A vectorized version of SVM

To speed things up, we will vectorize the loss and gradient calculations. This will be helpful for stochastic gradient descent.

```
In [13]: 1  import time
```

```
In [14]: 1  ## Implement svm.fast_loss_and_grad which calculates the loss and gradient
2  # WITHOUT using any for loops.
3
4  # Standard loss and gradient
5  tic = time.time()
6  loss, grad = svm.loss_and_grad(X_dev, y_dev)
7  toc = time.time()
8  print('Normal loss / grad_norm: {} / {} computed in {}'.format(loss, np.linalg.norm(grad, 'fro'), toc - tic))
9
10 tic = time.time()
11 loss_vectorized, grad_vectorized = svm.fast_loss_and_grad(X_dev, y_dev)
12 toc = time.time()
13 print('Vectorized loss / grad: {} / {} computed in {}'.format(loss_vectorized, np.linalg.norm(grad_vectorized, 'fro'),
14
15 # The losses should match but your vectorized implementation should be much faster.
16 print('difference in loss / grad: {} / {}'.format(loss - loss_vectorized, np.linalg.norm(grad - grad_vectorized)))
17
18 # You should notice a speedup with the same output, i.e., differences on the order of 1e-12
```


```
Normal loss / grad_norm: 16330.202893832919 / 2131.1522324879757 computed in 0.028795242309570312s
Vectorized loss / grad: 16330.202893832915 / 2131.152232487976 computed in 0.009291887283325195s
difference in loss / grad: 3.637978807091713e-12 / 7.3279039256198e-12
```

## Stochastic gradient descent

We now implement stochastic gradient descent. This uses the same principles of gradient descent we discussed in class, however, it calculates the gradient by only using examples from a subset of the training set (so each gradient calculation is faster).

```
In [15]: 1 # Implement svm.train() by filling in the code to extract a batch of data
2 # and perform the gradient step.
3
4 tic = time.time()
5 loss_hist = svm.train(X_train, y_train, learning_rate=5e-4,
6                       num_iters=1500, verbose=True)
7 toc = time.time()
8 print('That took {}s'.format(toc - tic))
9
10 plt.plot(loss_hist)
11 plt.xlabel('Iteration number')
12 plt.ylabel('Loss value')
13 plt.show()
```

```
iteration 0 / 1500: loss 16557.38000190916
iteration 100 / 1500: loss 4701.089451272714
iteration 200 / 1500: loss 4017.333137942788
iteration 300 / 1500: loss 3681.9226471953625
iteration 400 / 1500: loss 2732.6164373988995
iteration 500 / 1500: loss 2786.6378424645054
iteration 600 / 1500: loss 2837.0357842782673
iteration 700 / 1500: loss 2206.2348687399317
iteration 800 / 1500: loss 2269.0388241169803
iteration 900 / 1500: loss 2543.23781538592
iteration 1000 / 1500: loss 2566.6921357268275
iteration 1100 / 1500: loss 2182.068905905164
iteration 1200 / 1500: loss 1861.1182244250458
iteration 1300 / 1500: loss 1982.9013858528251
iteration 1400 / 1500: loss 1927.520415858212
That took 3.7440080642700195s
```



**Evaluate the performance of the trained SVM on the validation data.**

```
In [16]: 1 ## Implement svm.predict() and use it to compute the training and testing error.
2
3 y_train_pred = svm.predict(X_train)
4 print('training accuracy: {}'.format(np.mean(np.equal(y_train, y_train_pred), )))
5 y_val_pred = svm.predict(X_val)
6 print('validation accuracy: {}'.format(np.mean(np.equal(y_val, y_val_pred), )))
```

```
training accuracy: 0.28530612244897957
validation accuracy: 0.3
```

## Optimize the SVM

Note, to make things faster and simpler, we won't do k-fold cross-validation, but will only optimize the hyperparameters on the validation dataset (X\_val, y\_val).

```

In [18]: 1 # ===== #
2 # YOUR CODE HERE:
3 # Train the SVM with different learning rates and evaluate on the
4 # validation data.
5 # Report:
6 # - The best learning rate of the ones you tested.
7 # - The best VALIDATION accuracy corresponding to the best VALIDATION error.
8 #
9 # Select the SVM that achieved the best validation error and report
10 # its error rate on the test set.
11 # Note: You do not need to modify SVM class for this section
12 # ===== #
13 rates = [1e-5, 5e-4, 1e-4, 5e-3, 1e-3, 5e-2, 1e-2, 5e-1, 1e-1]
14 accuracies = []
15 for rate in rates:
16     svm.train(X_train, y_train, learning_rate=rate, num_iters=1500, verbose=False)
17     pred = svm.predict(X_val)
18     accuracy = np.sum(y_val == pred) / len(y_val)
19     accuracies.append(accuracy)
20 best_idx = np.argmax(accuracies)
21
22 print("The best learning rate: ", rates[best_idx])
23 print("The best validation accuracy: ", accuracies[best_idx])
24 print("The best validation error: ", 1 - accuracies[best_idx])
25
26 svm.train(X_train, y_train, learning_rate=rates[best_idx], num_iters=1500, verbose=False)
27 pred = svm.predict(X_test)
28 error_rate = 1 - (np.sum(y_test == pred) / len(y_test))
29 print("Error rate on test set: ", error_rate)
30 # ===== #
31 # END YOUR CODE HERE
32 # ===== #
33

```

```

The best learning rate: 0.01
The best validation accuracy: 0.312
The best validation error: 0.688
Error rate on test set: 0.6890000000000001

```

```

In [ ]: 1

```



```

1 import numpy as np
2 import pdb
3
4 """
5 This code was based off of code from cs231n at Stanford University, and
6 modified for ece239as at UCLA.
7 """
8 class SVM(object):
9
10     def __init__(self, dims=[10, 3073]):
11         self.init_weights(dims=dims)
12
13     def init_weights(self, dims):
14         """
15         Initializes the weight matrix of the SVM. Note that it has shape (C, D)
16         where C is the number of classes and D is the feature size.
17         """
18         self.W = np.random.normal(size=dims)
19
20     def loss(self, X, y):
21         """
22         Calculates the SVM loss.
23
24         Inputs have dimension D, there are C classes, and we operate on
25         minibatches of N examples.
26
27         Inputs:
28         - X: A numpy array of shape (N, D) containing a minibatch of data.
29         - y: A numpy array of shape (N,) containing training labels; y[i] = c
30             means that X[i] has label c, where 0 <= c < C.
31
32         Returns a tuple of:
33         - loss as single float
34         """
35         # compute the loss and the gradient
36         num_classes = self.W.shape[0]
37         num_train = X.shape[0]
38         loss = 0.0
39
40         for i in np.arange(num_train):
41             # ===== #
42             # YOUR CODE HERE:
43             # Calculate the normalized SVM loss, and store it as 'loss'.
44             # (That is, calculate the sum of the losses of all the training
45             # set margins, and then normalize the loss by the number of
46             # training examples.)
47             # ===== #
48             score = np.dot(self.W, X[i])
49             cur_y = y[i]
50             difference = score + 1 - score[cur_y]
51             difference[cur_y] = 0
52             difference = np.maximum(difference, 0)
53             cur_loss = np.sum(difference)
54             loss += cur_loss
55         loss /= num_train
56
57         # ===== #
58         # END YOUR CODE HERE
59         # ===== #
60
61         return loss
62
63     def loss_and_grad(self, X, y):
64         """
65         Same as self.loss(X, y), except that it also returns the gradient.
66
67         Output: grad -- a matrix of the same dimensions as W containing
68             the gradient of the loss with respect to W.
69         """
70
71         # compute the loss and the gradient
72         num_classes = self.W.shape[0]
73         num_train = X.shape[0]
74         loss = 0.0
75         grad = np.zeros_like(self.W)
76
77         for i in np.arange(num_train):
78             # ===== #
79             # YOUR CODE HERE:
80             # Calculate the SVM loss and the gradient. Store the gradient in
81             # the variable grad.

```

```

82 # ===== #
83 score = np.dot(self.W, X[i])
84 cur_y = y[i]
85 difference = score + 1 - score[cur_y]
86 difference[cur_y] = 0
87 difference = np.maximum(difference, 0)
88 cur_loss = np.sum(difference)
89 loss += cur_loss
90 count_positive = 0
91 for j in range(len(difference)):
92     if difference[j] > 0:
93         grad[j] += X[i]
94         count_positive += 1
95 grad[cur_y] -= count_positive * X[i]
96
97 # ===== #
98 # END YOUR CODE HERE
99 # ===== #
100
101 loss /= num_train
102 grad /= num_train
103
104 return loss, grad
105
106 def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
107     """
108     sample a few random elements and only return numerical
109     in these dimensions.
110     """
111
112     for i in np.arange(num_checks):
113         ix = tuple([np.random.randint(m) for m in self.W.shape])
114
115         oldval = self.W[ix]
116         self.W[ix] = oldval + h # increment by h
117         fxph = self.loss(X, y)
118         self.W[ix] = oldval - h # decrement by h
119         fxmh = self.loss(X, y) # evaluate f(x - h)
120         self.W[ix] = oldval # reset
121
122         grad_numerical = (fxph - fxmh) / (2 * h)
123         grad_analytic = your_grad[ix]
124         rel_error = abs(grad_numerical - grad_analytic) / (abs(grad_numerical)
+ abs(grad_analytic))
125         print('numerical: %f analytic: %f, relative error: %e' %
(grad_numerical, grad_analytic, rel_error))
126
127 def fast_loss_and_grad(self, X, y):
128     """
129     A vectorized implementation of loss_and_grad. It shares the same
130     inputs and outputs as loss_and_grad.
131     """
132     loss = 0.0
133     grad = np.zeros(self.W.shape) # initialize the gradient as zero
134
135     # ===== #
136     # YOUR CODE HERE:
137     # Calculate the SVM loss WITHOUT any for loops.
138     # ===== #
139     score = np.dot(self.W, X.T)
140     cur_y_idx = y, range(len(y))
141     cur_y_score = score[cur_y_idx]
142     difference = score - cur_y_score + np.ones(score.shape)
143     difference[cur_y_idx] = 0
144     difference = np.maximum(difference, 0)
145     loss = np.sum(difference) / X.shape[0]
146
147     # ===== #
148     # END YOUR CODE HERE
149     # ===== #
150
151
152     # ===== #
153     # YOUR CODE HERE:
154     # Calculate the SVM grad WITHOUT any for loops.
155     # ===== #
156     difference[difference > 0] = 1
157     difference[cur_y_idx] = -1 * np.sum(difference, axis=0)
158     grad = np.dot(difference, X) / X.shape[0]
159
160     # ===== #
161     # END YOUR CODE HERE
162     # ===== #
163
164     return loss, grad

```

```

164
165 def train(self, X, y, learning_rate=1e-3, num_iters=100,
166           batch_size=200, verbose=False):
167     """
168     Train this linear classifier using stochastic gradient descent.
169
170     Inputs:
171     - X: A numpy array of shape (N, D) containing training data; there are N
172         training samples each of dimension D.
173     - y: A numpy array of shape (N,) containing training labels; y[i] = c
174         means that X[i] has label 0 ≤ c < C for C classes.
175     - learning_rate: (float) learning rate for optimization.
176     - num_iters: (integer) number of steps to take when optimizing
177     - batch_size: (integer) number of training examples to use at each step.
178     - verbose: (boolean) If true, print progress during optimization.
179
180     Outputs:
181     A list containing the value of the loss function at each training
182     iteration.
183     """
184     num_train, dim = X.shape
185     num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is
186     number of classes
187
188     self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes the
189     weights of self.W
190
191     # Run stochastic gradient descent to optimize W
192     loss_history = []
193
194     for it in np.arange(num_iters):
195         X_batch = None
196         y_batch = None
197
198         # ===== #
199         # YOUR CODE HERE:
200         # Sample batch_size elements from the training data for use in
201         # gradient descent. After sampling,
202         # - X_batch should have shape: (dim, batch_size)
203         # - y_batch should have shape: (batch_size,)
204         # The indices should be randomly generated to reduce correlations
205         # in the dataset. Use np.random.choice. It's okay to sample with
206         # replacement.
207         # ===== #
208         indices = np.random.choice(X.shape[0], batch_size)
209         X_batch = X[indices]
210         y_batch = y[indices]
211         # ===== #
212         # END YOUR CODE HERE
213         # ===== #
214
215         # evaluate loss and gradient
216         loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
217         loss_history.append(loss)
218
219         # ===== #
220         # YOUR CODE HERE:
221         # Update the parameters, self.W, with a gradient step
222         # ===== #
223         self.W += -learning_rate*grad
224         # ===== #
225         # END YOUR CODE HERE
226         # ===== #
227
228         if verbose and it % 100 == 0:
229             print('iteration {} / {}: loss {}'.format(it, num_iters, loss))
230
231     return loss_history
232
233 def predict(self, X):
234     """
235     Inputs:
236     - X: N x D array of training data. Each row is a D-dimensional point.
237
238     Returns:
239     - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
240         array of length N, and each element is an integer giving the predicted
241         class.
242     """
243     y_pred = np.zeros(X.shape[0])
244
245     # ===== #
246     # YOUR CODE HERE:
247     # Predict the labels given the training data with the parameter self.W.

```

```
245 # ===== #
246 res = np.dot(self.W, X.T)
247 y_pred = np.argmax(res, axis=0)
248 # ===== #
249 # END YOUR CODE HERE
250 # ===== #
251
252 return y_pred
253
254
```

## This is the softmax workbook for ECE 239AS Assignment #2

Please follow the notebook linearly to implement a softmax classifier.

Please print out the workbook entirely when completed.

We thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class ([cs231n.stanford.edu](http://cs231n.stanford.edu)). These are the functions in the cs231n folders and code in the jupyter notebook to preprocess and show the images. The classifiers used are based off of code prepared for CS 231n as well.

The goal of this workbook is to give you experience with training a softmax classifier.

```
In [1]: 1 import random
        2 import numpy as np
        3 from cs231n.data_utils import load_CIFAR10
        4 import matplotlib.pyplot as plt
        5
        6 %matplotlib inline
        7 %load_ext autoreload
        8 %autoreload 2
```

```

In [2]: 1 def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000, num_dev=500):
2         """
3         Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
4         it for the linear classifier. These are the same steps as we used for the
5         SVM, but condensed to a single function.
6         """
7         # Load the raw CIFAR-10 data
8         cifar10_dir = 'cifar-10-batches-py'
9         X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)
10
11        # subsample the data
12        mask = list(range(num_training, num_training + num_validation))
13        X_val = X_train[mask]
14        y_val = y_train[mask]
15        mask = list(range(num_training))
16        X_train = X_train[mask]
17        y_train = y_train[mask]
18        mask = list(range(num_test))
19        X_test = X_test[mask]
20        y_test = y_test[mask]
21        mask = np.random.choice(num_training, num_dev, replace=False)
22        X_dev = X_train[mask]
23        y_dev = y_train[mask]
24
25        # Preprocessing: reshape the image data into rows
26        X_train = np.reshape(X_train, (X_train.shape[0], -1))
27        X_val = np.reshape(X_val, (X_val.shape[0], -1))
28        X_test = np.reshape(X_test, (X_test.shape[0], -1))
29        X_dev = np.reshape(X_dev, (X_dev.shape[0], -1))
30
31        # Normalize the data: subtract the mean image
32        mean_image = np.mean(X_train, axis = 0)
33        X_train -= mean_image
34        X_val -= mean_image
35        X_test -= mean_image
36        X_dev -= mean_image
37
38        # add bias dimension and transform into columns
39        X_train = np.hstack([X_train, np.ones((X_train.shape[0], 1))])
40        X_val = np.hstack([X_val, np.ones((X_val.shape[0], 1))])
41        X_test = np.hstack([X_test, np.ones((X_test.shape[0], 1))])
42        X_dev = np.hstack([X_dev, np.ones((X_dev.shape[0], 1))])
43
44        return X_train, y_train, X_val, y_val, X_test, y_test, X_dev, y_dev
45
46        # Invoke the above function to get our data.
47        X_train, y_train, X_val, y_val, X_test, y_test, X_dev, y_dev = get_CIFAR10_data()
48        print('Train data shape: ', X_train.shape)
49        print('Train labels shape: ', y_train.shape)
50        print('Validation data shape: ', X_val.shape)
51        print('Validation labels shape: ', y_val.shape)
52        print('Test data shape: ', X_test.shape)
53        print('Test labels shape: ', y_test.shape)
54        print('dev data shape: ', X_dev.shape)
55        print('dev labels shape: ', y_dev.shape)

```

```

Train data shape: (49000, 3073)
Train labels shape: (49000,)
Validation data shape: (1000, 3073)
Validation labels shape: (1000,)
Test data shape: (1000, 3073)
Test labels shape: (1000,)
dev data shape: (500, 3073)
dev labels shape: (500,)

```

## Training a softmax classifier.

The following cells will take you through building a softmax classifier. You will implement its loss function, then subsequently train it with gradient descent. Finally, you will choose the learning rate of gradient descent to optimize its classification performance.

```

In [3]: 1 from nnrl import Softmax

```

```

In [4]: 1 # Declare an instance of the Softmax class.
2 # Weights are initialized to a random value.
3 # Note, to keep people's first solutions consistent, we are going to use a random seed.
4
5 np.random.seed(1)
6
7 num_classes = len(np.unique(y_train))
8 num_features = X_train.shape[1]
9
10 softmax = Softmax(dims=[num_classes, num_features])

```

### Softmax loss

```
In [5]: 1 ## Implement the loss function of the softmax using a for loop over
        2 # the number of examples
        3
        4 loss = softmax.loss(X_train, y_train)
```

```
In [6]: 1 print(loss)

2.3277607028048863
```

### Question:

You'll notice the loss returned by the softmax is about 2.3 (if implemented correctly). Why does this value make sense?

### Answer:

Because at the beginning, all scores are zero. The loss =  $\log(\text{class number}) = \log(10) = 2.3$ .

### Softmax gradient

```
In [7]: 1 ## Calculate the gradient of the softmax loss in the Softmax class.
        2 # For convenience, we'll write one function that computes the loss
        3 # and gradient together, softmax.loss_and_grad(X, y)
        4 # You may copy and paste your loss code from softmax.loss() here, and then
        5 # use the appropriate intermediate values to calculate the gradient.
        6
        7 loss, grad = softmax.loss_and_grad(X_dev, y_dev)
        8
        9 # Compare your gradient to a gradient check we wrote.
       10 # You should see relative gradient errors on the order of 1e-07 or less if you implemented the gradient correctly.
       11 softmax.grad_check_sparse(X_dev, y_dev, grad)
```

```
numerical: 0.028704 analytic: 0.028704, relative error: 1.532463e-07
numerical: 1.575618 analytic: 1.575618, relative error: 2.149405e-08
numerical: -1.361092 analytic: -1.361092, relative error: 1.412920e-08
numerical: 1.771825 analytic: 1.771825, relative error: 6.019498e-09
numerical: 0.693808 analytic: 0.693808, relative error: 9.775131e-08
numerical: 0.951248 analytic: 0.951248, relative error: 5.075736e-08
numerical: -0.634508 analytic: -0.634508, relative error: 7.046441e-08
numerical: -0.150207 analytic: -0.150207, relative error: 7.570901e-08
numerical: 1.558823 analytic: 1.558823, relative error: 3.511548e-09
numerical: -1.900958 analytic: -1.900958, relative error: 2.879331e-08
```

## A vectorized version of Softmax

To speed things up, we will vectorize the loss and gradient calculations. This will be helpful for stochastic gradient descent.

```
In [8]: 1 import time
```

```
In [9]: 1 ## Implement softmax.fast_loss_and_grad which calculates the loss and gradient
        2 # WITHOUT using any for loops.
        3
        4 # Standard loss and gradient
        5 tic = time.time()
        6 loss, grad = softmax.loss_and_grad(X_dev, y_dev)
        7 toc = time.time()
        8 print('Normal loss / grad_norm: {} / {} computed in {}s'.format(loss, np.linalg.norm(grad, 'fro'), toc - tic))
        9
       10 tic = time.time()
       11 loss_vectorized, grad_vectorized = softmax.fast_loss_and_grad(X_dev, y_dev)
       12 toc = time.time()
       13 print('Vectorized loss / grad: {} / {} computed in {}s'.format(loss_vectorized, np.linalg.norm(grad_vectorized, 'fro'),
       14
       15 # The losses should match but your vectorized implementation should be much faster.
       16 print('difference in loss / grad: {} / {} '.format(loss - loss_vectorized, np.linalg.norm(grad - grad_vectorized)))
       17
       18 # You should notice a speedup with the same output.
```

```
Normal loss / grad_norm: 2.2978300107980734 / 313.5219408763691 computed in 0.04951214790344238s
Vectorized loss / grad: 2.2978300107980765 / 313.5219408763691 computed in 0.004825115203857422s
difference in loss / grad: -3.1086244689504383e-15 / 2.835334683406196e-13
```

## Stochastic gradient descent

We now implement stochastic gradient descent. This uses the same principles of gradient descent we discussed in class, however, it calculates the gradient by only using examples from a subset of the training set (so each gradient calculation is faster).

## Question:


How should the softmax gradient descent training step differ from the svm training step, if at all?

## Answer:

The overall processes of gradient descent won't change. The only difference is their loss function and values of gradient.

```
In [10]: 1 # Implement softmax.train() by filling in the code to extract a batch of data
2 # and perform the gradient step.
3 import time
4
5
6 tic = time.time()
7 loss_hist = softmax.train(X_train, y_train, learning_rate=1e-7,
8                           num_iters=1500, verbose=True)
9 toc = time.time()
10 print('That took {}s'.format(toc - tic))
11
12 plt.plot(loss_hist)
13 plt.xlabel('Iteration number')
14 plt.ylabel('Loss value')
15 plt.show()
```

```
iteration 0 / 1500: loss 2.3365926606637544
iteration 100 / 1500: loss 2.0557222613850827
iteration 200 / 1500: loss 2.0357745120662813
iteration 300 / 1500: loss 1.9813348165609888
iteration 400 / 1500: loss 1.9583142443981612
iteration 500 / 1500: loss 1.862265307354135
iteration 600 / 1500: loss 1.8532611454359382
iteration 700 / 1500: loss 1.8353062223725827
iteration 800 / 1500: loss 1.829389246882764
iteration 900 / 1500: loss 1.8992158530357484
iteration 1000 / 1500: loss 1.97835035402523
iteration 1100 / 1500: loss 1.8470797913532633
iteration 1200 / 1500: loss 1.8411450268664082
iteration 1300 / 1500: loss 1.79104024957921
iteration 1400 / 1500: loss 1.8705803029382257
That took 3.935007095336914s
```



Evaluate the performance of the trained softmax classifier on the validation data.

```
In [11]: 1 ## Implement softmax.predict() and use it to compute the training and testing error.
2
3 y_train_pred = softmax.predict(X_train)
4 print('training accuracy: {}'.format(np.mean(np.equal(y_train,y_train_pred), )))
5 y_val_pred = softmax.predict(X_val)
6 print('validation accuracy: {}'.format(np.mean(np.equal(y_val, y_val_pred)), ))
```

```
training accuracy: 0.3811428571428571
validation accuracy: 0.398
```

## Optimize the softmax classifier

You may copy and paste your optimization code from the SVM here.

```
In [12]: 1 np.finfo(float).eps
```

```
Out[12]: 2.220446049250313e-16
```



```

In [13]: 1 # ===== #
2 # YOUR CODE HERE:
3 #   Train the Softmax classifier with different learning rates and
4 #   evaluate on the validation data.
5 #   Report:
6 #       - The best learning rate of the ones you tested.
7 #       - The best validation accuracy corresponding to the best validation error.
8 #
9 #   Select the SVM that achieved the best validation error and report
10 #   its error rate on the test set.
11 # ===== #
12 rates = [1e-5, 5e-4, 1e-4, 5e-3, 1e-3, 5e-2, 1e-2, 5e-1, 1e-1]
13 accuracies = []
14 for rate in rates:
15     softmax.train(X_train, y_train, learning_rate=rate, num_iters=1500, verbose=False)
16     pred = softmax.predict(X_val)
17     accuracy = np.sum(y_val == pred) / len(y_val)
18     accuracies.append(accuracy)
19 best_idx = np.argmax(accuracies)
20
21 print("The best learning rate: ", rates[best_idx])
22 print("The best validation accuracy: ", accuracies[best_idx])
23 print("The best validation error: ", 1 - accuracies[best_idx])
24
25 softmax.train(X_train, y_train, learning_rate=rates[best_idx], num_iters=1500, verbose=False)
26 pred = softmax.predict(X_test)
27 error_rate = 1 - (np.sum(y_test == pred) / len(y_test))
28 print("Error rate on test set: ", error_rate)
29 # ===== #
30 # END YOUR CODE HERE
31 # ===== #
32

```

```

The best learning rate: 1e-05
The best validation accuracy: 0.349
The best validation error: 0.651
Error rate on test set: 0.6990000000000001

```

```

In [ ]: 1

```

```

1 import numpy as np
2
3 class Softmax(object):
4
5     def __init__(self, dims=[10, 3073]):
6         self.init_weights(dims=dims)
7
8     def init_weights(self, dims):
9         """
10        Initializes the weight matrix of the Softmax classifier.
11        Note that it has shape (C, D) where C is the number of
12        classes and D is the feature size.
13        """
14        self.W = np.random.normal(size=dims) * 0.0001
15
16    def loss(self, X, y):
17        """
18        Calculates the softmax loss.
19
20        Inputs have dimension D, there are C classes, and we operate on
minibatches
21        of N examples.
22
23        Inputs:
24        - X: A numpy array of shape (N, D) containing a minibatch of data.
25        - y: A numpy array of shape (N,) containing training labels; y[i] = c
means
26        that X[i] has label c, where 0 <= c < C.
27
28        Returns a tuple of:
29        - loss as single float
30        """
31
32        # Initialize the loss to zero.
33        loss = 0.0
34
35        # ===== #
36        # YOUR CODE HERE:
37        # Calculate the normalized softmax loss. Store it as the variable
loss.
38        # (That is, calculate the sum of the losses of all the training
39        # set margins, and then normalize the loss by the number of
40        # training examples.)
41        # ===== #
42
43        scores = np.dot(self.W, X.T)
44        for i in range(scores.shape[1]):
45            score = scores[:, i]
46            score -= np.max(score)
47            cur_class_score = score[y[i]]
48            loss += np.log(np.sum(np.exp(score)))
49            loss -= cur_class_score
50        loss /= X.shape[0]
51
52        # ===== #
53        # END YOUR CODE HERE
54        # ===== #
55
56        return loss
57
58    def loss_and_grad(self, X, y):
59        """
60        Same as self.loss(X, y), except that it also returns the gradient.
61
62        Output: grad -- a matrix of the same dimensions as W containing
63        the gradient of the loss with respect to W.
64        """
65
66        # Initialize the loss and gradient to zero.
67        loss = 0.0
68        grad = np.zeros_like(self.W)
69
70        # ===== #
71        # YOUR CODE HERE:
72        # Calculate the softmax loss and the gradient. Store the gradient
73        # as the variable grad.
74        # ===== #
75
76        scores = np.dot(self.W, X.T)
77        num_train = X.shape[0]
78        num_classes = self.W.shape[0]
79        for i in range(num_train):
80            score = scores[:, i]
81            score -= np.max(score)

```

```

82     cur_class_score = score[y[i]]
83     sum_exp = np.sum(np.exp(score))
84
85     loss += np.log(sum_exp)
86     loss -= cur_class_score
87
88     for j in range(num_classes):
89         grad[j] += (np.exp(score[j]) / sum_exp) * X[i]
90     grad[y[i]] -= X[i]
91     loss /= num_train
92     grad /= num_train
93
94     # ===== #
95     # END YOUR CODE HERE
96     # ===== #
97
98     return loss, grad
99
100 def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
101     """
102     sample a few random elements and only return numerical
103     in these dimensions.
104     """
105
106     for i in np.arange(num_checks):
107         ix = tuple([np.random.randint(m) for m in self.W.shape])
108
109         oldval = self.W[ix]
110         self.W[ix] = oldval + h # increment by h
111         fxph = self.loss(X, y)
112         self.W[ix] = oldval - h # decrement by h
113         fxmh = self.loss(X,y) # evaluate f(x - h)
114         self.W[ix] = oldval # reset
115
116         grad_numerical = (fxph - fxmh) / (2 * h)
117         grad_analytic = your_grad[ix]
118         rel_error = abs(grad_numerical - grad_analytic) / (abs(grad_numerical)
+ abs(grad_analytic))
119         print('numerical: %f analytic: %f, relative error: %e' %
(grad_numerical, grad_analytic, rel_error))
120
121 def fast_loss_and_grad(self, X, y):
122     """
123     A vectorized implementation of loss_and_grad. It shares the same
124     inputs and ouptuts as loss_and_grad.
125     """
126     loss = 0.0
127     grad = np.zeros(self.W.shape) # initialize the gradient as zero
128
129     # ===== #
130     # YOUR CODE HERE:
131     # Calculate the softmax loss and gradient WITHOUT any for loops.
132     # ===== #
133
134     num_train = X.shape[0]
135     scores = np.dot(self.W, X.T)
136     scores -= np.max(scores, axis=0, keepdims=True)
137     exp_scores = np.exp(scores)
138     probs = exp_scores / np.sum(exp_scores, axis=0, keepdims=True)
139     corres_probs = probs[y, range(num_train)]
140     log_probs = -np.log(corres_probs.clip(min=np.finfo(float).eps))
141     loss = np.sum(log_probs) / num_train
142
143     probs[y, range(num_train)] -= 1
144     grad = np.dot(probs, X)
145     grad /= num_train
146
147     # ===== #
148     # END YOUR CODE HERE
149     # ===== #
150
151     return loss, grad
152
153 def train(self, X, y, learning_rate=1e-3, num_iters=100,
154         batch_size=200, verbose=False):
155     """
156     Train this linear classifier using stochastic gradient descent.
157
158     Inputs:
159     - X: A numpy array of shape (N, D) containing training data; there are N
160         training samples each of dimension D.
161     - y: A numpy array of shape (N,) containing training labels; y[i] = c
162         means that X[i] has label 0 <= c < C for C classes.
163     - learning_rate: (float) learning rate for optimization.

```

```

164 - num_iters: (integer) number of steps to take when optimizing
165 - batch_size: (integer) number of training examples to use at each step.
166 - verbose: (boolean) If true, print progress during optimization.
167
168 Outputs:
169 A list containing the value of the loss function at each training
iteration.
170 """
171 num_train, dim = X.shape
172 num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is
number of classes
173
174 self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes the
weights of self.W
175
176 # Run stochastic gradient descent to optimize W
177 loss_history = []
178
179 for it in np.arange(num_iters):
180     X_batch = None
181     y_batch = None
182
183     # ===== #
184     # YOUR CODE HERE:
185     # Sample batch_size elements from the training data for use in
186     # gradient descent. After sampling,
187     # - X_batch should have shape: (dim, batch_size)
188     # - y_batch should have shape: (batch_size,)
189     # The indices should be randomly generated to reduce correlations
190     # in the dataset. Use np.random.choice. It's okay to sample with
191     # replacement.
192     # ===== #
193     indices = np.random.choice(X.shape[0], batch_size)
194     X_batch = X[indices]
195     y_batch = y[indices]
196     # ===== #
197     # END YOUR CODE HERE
198     # ===== #
199
200     # evaluate loss and gradient
201     loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
202     loss_history.append(loss)
203
204     # ===== #
205     # YOUR CODE HERE:
206     # Update the parameters, self.W, with a gradient step
207     # ===== #
208
209     self.W += -learning_rate * grad
210
211     # ===== #
212     # END YOUR CODE HERE
213     # ===== #
214
215     if verbose and it % 100 == 0:
216         print('iteration {} / {}: loss {}'.format(it, num_iters, loss))
217
218     return loss_history
219
220 def predict(self, X):
221     """
222     Inputs:
223     - X: N x D array of training data. Each row is a D-dimensional point.
224
225     Returns:
226     - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
227       array of length N, and each element is an integer giving the predicted
228       class.
229     """
230     y_pred = np.zeros(X.shape[0])
231     # ===== #
232     # YOUR CODE HERE:
233     # Predict the labels given the training data.
234     # ===== #
235     scores = np.dot(self.W, X.T)
236     y_pred = np.argmax(scores, axis = 0)
237     # ===== #
238     # END YOUR CODE HERE
239     # ===== #
240
241     return y_pred
242
243

```