


1. Solution:

(a) From the question, we know that Wx stores the results of projection of x to each row of W . So $y = Wx$ is to encode the information of x . Then we have $\hat{x} = W^T y$ which is to decode the information of x . So $L = \frac{1}{2} \|W^T Wx - x\|^2 = \frac{1}{2} \|W^T y - x\|^2 = \frac{1}{2} \|\hat{x} - x\|^2$ is the sum of difference between the reconstructed x and the original x . Thus, this minimization finds a W that ought to preserve information about x .

(b) The computational graph is as below:




(c) Let two paths are L_1 and L_2 . Then we have their gradients as $\nabla_W L_1$ and $\nabla_W L_2$. Then we can get the final gradient as $\nabla_W L = \nabla_W L_1 + \nabla_W L_2$.

(d) From the computational graph, we can get:

$$\begin{aligned} \nabla_W L &= \nabla_W L_1 + \nabla_W L_2 \\ &= W(W^T Wx - x)x^T + (Wx)(W^T Wx - x)^T. \end{aligned}$$

2. Solution:

(a) The computational graph for L_1 is as below:




(b) From the computational graph above, we have

$$\frac{\partial L}{\partial X} = -\alpha D(\alpha X X^T + \beta^T I)^{-1} X$$

$\frac{\partial L}{\partial x} = 2x \frac{\partial L}{\partial k}$ (in the next graph, ~~the~~ the node transform represents ~~the~~ the transformation from x to k).

(c) The computational graph for L_2 is as below:



(d) From the computational graph above, we have

$$\frac{\partial l_2}{\partial \chi} = \alpha K^T (\gamma \gamma^T) K^T \chi = \alpha (\alpha \chi \chi^T + \beta^T I) \gamma^T (\gamma \gamma^T) (\alpha \chi \chi^T + \beta^T I)^T \chi$$

(e) From the results of $\frac{\partial L_1}{\partial x}$ and $\frac{\partial L_2}{\partial x}$, we can get

$$\frac{\partial L}{\partial x} = \frac{\partial L_1}{\partial x} + \frac{\partial L_2}{\partial x}$$

$$= -\alpha D(\alpha XX^T + \beta^{-1} I)^{-1} X + \alpha(\alpha XX^T + \beta^{-1} I)^{-1} (Y Y^T) (\alpha XX^T + \beta^{-1} I)^{-1} X$$

$$= -\alpha D K^T x + \partial K^T (Y Y^T) K^T x$$

This is the 2-layer neural network workbook for ECE 239AS Assignment #3

Please follow the notebook linearly to implement a two layer neural network.

Please print out the workbook entirely when completed.

We thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class (cs231n.stanford.edu). These are the functions in the cs231n folders and code in the jupyter notebook to preprocess and show the images. The classifiers used are based off of code prepared for CS 231n as well.

The goal of this workbook is to give you experience with training a two layer neural network.

```
In [1]: 1 import random
2 import numpy as np
3 from cs231n.data_utils import load_CIFAR10
4 import matplotlib.pyplot as plt
5
6 %matplotlib inline
7 %load_ext autoreload
8 %autoreload 2
9
10 def rel_error(x, y):
11     """ returns relative error """
12     return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

Toy example

Before loading CIFAR-10, there will be a toy example to test your implementation of the forward and backward pass

```
In [2]: 1 from nnlib.neural_net import TwoLayerNet

In [3]: 1 # Create a small net and some toy data to check your implementations.
2 # Note that we set the random seed for repeatable experiments.
3
4 input_size = 4
5 hidden_size = 10
6 num_classes = 3
7 num_inputs = 5
8
9 def init_toy_model():
10     np.random.seed(0)
11     return TwoLayerNet(input_size, hidden_size, num_classes, std=1e-1)
12
13 def init_toy_data():
14     np.random.seed(1)
15     X = 10 * np.random.randn(num_inputs, input_size)
16     y = np.array([0, 1, 2, 2, 1])
17     return X, y
18
19 net = init_toy_model()
20 X, y = init_toy_data()
```

Compute forward pass scores

```
In [5]: 1  ## Implement the forward pass of the neural network.
2
3  # Note, there is a statement if y is None: return scores, which is why
4  # the following call will calculate the scores.
5  scores = net.loss(X)
6  print('Your scores:')
7  print(scores)
8  print()
9  print('correct scores:')
10 correct_scores = np.asarray([
11     [-1.07260209,  0.05083871, -0.87253915],
12     [-2.02778743, -0.10832494, -1.52641362],
13     [-0.74225908,  0.15259725, -0.39578548],
14     [-0.38172726,  0.10835902, -0.17328274],
15     [-0.64417314, -0.18886813, -0.41106892]])
16 print(correct_scores)
17 print()
18
19 # The difference should be very small. We get < 1e-7
20 print('Difference between your scores and correct scores:')
21 print(np.sum(np.abs(scores - correct_scores)))
```

```
Your scores:
[[-1.07260209  0.05083871 -0.87253915]
 [-2.02778743 -0.10832494 -1.52641362]
 [-0.74225908  0.15259725 -0.39578548]
 [-0.38172726  0.10835902 -0.17328274]
 [-0.64417314 -0.18886813 -0.41106892]]
```

```
correct scores:
[[-1.07260209  0.05083871 -0.87253915]
 [-2.02778743 -0.10832494 -1.52641362]
 [-0.74225908  0.15259725 -0.39578548]
 [-0.38172726  0.10835902 -0.17328274]
 [-0.64417314 -0.18886813 -0.41106892]]
```

```
Difference between your scores and correct scores:
3.381231233889892e-08
```

Forward pass loss

```
In [6]: 1  loss, _ = net.loss(X, y, reg=0.05)
2  correct_loss = 1.071696123862817
3
4  # should be very small, we get < 1e-12
5  print('Difference between your loss and correct loss:')
6  print(np.sum(np.abs(loss - correct_loss)))
```

```
Difference between your loss and correct loss:
0.0
```

```
In [7]: 1  print(loss)

1.071696123862817
```

Backward pass

Implements the backwards pass of the neural network. Check your gradients with the gradient check utilities provided.

```
In [9]: 1  from cs231n.gradient_check import eval_numerical_gradient
2
3  # Use numeric gradient checking to check your implementation of the backward pass.
4  # If your implementation is correct, the difference between the numeric and
5  # analytic gradients should be less than 1e-8 for each of W1, W2, b1, and b2.
6
7  loss, grads = net.loss(X, y, reg=0.05)
8
9  # these should all be less than 1e-8 or so
10 for param_name in grads:
11     f = lambda W: net.loss(X, y, reg=0.05)[0]
12     param_grad_num = eval_numerical_gradient(f, net.params[param_name], verbose=False)
13     print('{} max relative error: {}'.format(param_name, rel_error(param_grad_num, grads[param_name])))
```


```
W2 max relative error: 2.9632227682005116e-10
b2 max relative error: 1.2482660547101085e-09
W1 max relative error: 1.2832874456864775e-09
b1 max relative error: 3.1726806716844575e-09
```

Training the network

Implement `neural_net.train()` to train the network via stochastic gradient descent, much like the softmax and SVM.

```
In [10]: 1 net = init_toy_model()
2 stats = net.train(X, y, X, y,
3                 learning_rate=1e-1, reg=5e-6,
4                 num_iters=100, verbose=False)
5
6 print('Final training loss: ', stats['loss_history'][-1])
7
8 # plot the loss history
9 plt.plot(stats['loss_history'])
10 plt.xlabel('iteration')
11 plt.ylabel('training loss')
12 plt.title('Training Loss history')
13 plt.show()
```

Final training loss: 0.014498902952971663



Classify CIFAR-10

Do classification on the CIFAR-10 dataset.


```

In [11]: 1 from cs231n.data_utils import load_CIFAR10
2
3 def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000):
4     """
5     Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
6     it for the two-layer neural net classifier. These are the same steps as
7     we used for the SVM, but condensed to a single function.
8     """
9     # Load the raw CIFAR-10 data
10    cifar10_dir = 'cifar-10-batches-py'
11    X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)
12
13    # Subsample the data
14    mask = list(range(num_training, num_training + num_validation))
15    X_val = X_train[mask]
16    y_val = y_train[mask]
17    mask = list(range(num_training))
18    X_train = X_train[mask]
19    y_train = y_train[mask]
20    mask = list(range(num_test))
21    X_test = X_test[mask]
22    y_test = y_test[mask]
23
24    # Normalize the data: subtract the mean image
25    mean_image = np.mean(X_train, axis=0)
26    X_train -= mean_image
27    X_val -= mean_image
28    X_test -= mean_image
29
30    # Reshape data to rows
31    X_train = X_train.reshape(num_training, -1)
32    X_val = X_val.reshape(num_validation, -1)
33    X_test = X_test.reshape(num_test, -1)
34
35    return X_train, y_train, X_val, y_val, X_test, y_test
36
37
38 # Invoke the above function to get our data.
39 X_train, y_train, X_val, y_val, X_test, y_test = get_CIFAR10_data()
40 print('Train data shape: ', X_train.shape)
41 print('Train labels shape: ', y_train.shape)
42 print('Validation data shape: ', X_val.shape)
43 print('Validation labels shape: ', y_val.shape)
44 print('Test data shape: ', X_test.shape)
45 print('Test labels shape: ', y_test.shape)

```

```

Train data shape: (49000, 3072)
Train labels shape: (49000,)
Validation data shape: (1000, 3072)
Validation labels shape: (1000,)
Test data shape: (1000, 3072)
Test labels shape: (1000,)

```

Running SGD

If your implementation is correct, you should see a validation accuracy of around 28-29%.

```
In [12]: 1 input_size = 32 * 32 * 3
2 hidden_size = 50
3 num_classes = 10
4 net = TwoLayerNet(input_size, hidden_size, num_classes)
5
6 # Train the network
7 stats = net.train(X_train, y_train, X_val, y_val,
8                 num_iters=1000, batch_size=200,
9                 learning_rate=1e-4, learning_rate_decay=0.95,
10                reg=0.25, verbose=True)
11
12 # Predict on the validation set
13 val_acc = (net.predict(X_val) == y_val).mean()
14 print('Validation accuracy: ', val_acc)
15
16 # Save this net as the variable subopt_net for later comparison.
17 subopt_net = net
```

```
iteration 0 / 1000: loss 2.302757518613176
iteration 100 / 1000: loss 2.302120159207236
iteration 200 / 1000: loss 2.2956136007408703
iteration 300 / 1000: loss 2.2518259043164135
iteration 400 / 1000: loss 2.188995235046776
iteration 500 / 1000: loss 2.1162527791897747
iteration 600 / 1000: loss 2.064670827698217
iteration 700 / 1000: loss 1.9901688623083942
iteration 800 / 1000: loss 2.002827640124685
iteration 900 / 1000: loss 1.9465176817856495
Validation accuracy: 0.283
```

Questions:



The training accuracy isn't great.

- (1) What are some of the reasons why this is the case? Take the following cell to do some analyses and then report your answers in the cell following the one below.
- (2) How should you fix the problems you identified in (1)?

```
In [13]: 1 stats['train_acc_history']
```

```
Out[13]: [0.095, 0.15, 0.25, 0.25, 0.315]
```

```
In [19]: 1 # ===== #
2 # YOUR CODE HERE:
3 # Do some debugging to gain some insight into why the optimization
4 # isn't great.
5 # ===== #
6
7 # Plot the loss function and train / validation accuracies
8
9 plt.plot(stats['loss_history'])
10 plt.title('Training Loss History')
11 plt.xlabel('Iteration')
12 plt.ylabel('Training Loss')
13 plt.show()
14
15 plt.plot(stats['train_acc_history'])
16 plt.plot(stats['val_acc_history'])
17 plt.title('Classification Accuracy History')
18 plt.xlabel('Iteration')
19 plt.ylabel('Accuracy')
20 plt.show()
21
22 # ===== #
23 # END YOUR CODE HERE
24 # ===== #
```



Answers:

- (1) On the one hand, from the training loss history, we can see that at the beginning, the loss is roughly stable and didn't decrease. After about 200 rounds of iteration, the loss decreased gradually. So we can guess that the learning rate is not large enough. On the other hand, from the accuracy history, we find that the curves are still increasing and not converge. So we need to do more rounds of iteration to make it converge.
- (2) First, I will increase the value of learning rate. Second, I will do more rounds of iteration to make the accuracy curves converge.

Optimize the neural network

Use the following part of the Jupyter notebook to optimize your hyperparameters on the validation set. Store your nets as `best_net`.


```

In [22]: 1 best_net = None # store the best model into this
2
3 # ===== #
4 # YOUR CODE HERE:
5 # Optimize over your hyperparameters to arrive at the best neural
6 # network. You should be able to get over 50% validation accuracy.
7 # For this part of the notebook, we will give credit based on the
8 # accuracy you get. Your score on this question will be multiplied by:
9 # min(floor((X - 28%) / %22, 1)
10 # where if you get 50% or higher validation accuracy, you get full
11 # points.
12 #
13 # Note, you need to use the same network structure (keep hidden_size = 50)!
14 # ===== #
15
16 best_acc = -1
17 learning_rates = [1e-1, 5e-2, 3e-2, 1e-2, 5e-3, 3e-3, 1e-3]
18 results = {}
19 for lr in learning_rates:
20     net = TwoLayerNet(input_size, hidden_size, num_classes)
21     stats = net.train(X_train, y_train, X_val, y_val,
22                      num_iters=5000, batch_size=200,
23                      learning_rate=lr, learning_rate_decay=0.95,
24                      reg=0.55)
25
26     y_train_pred = net.predict(X_train)
27     acc_train = np.mean(y_train == y_train_pred)
28     y_val_pred = net.predict(X_val)
29     acc_val = np.mean(y_val == y_val_pred)
30
31     results[lr] = (acc_train, acc_val)
32
33     if best_acc < acc_val:
34         best_acc = acc_val
35         best_net = net
36
37 # Print out results
38 for lr in sorted(results):
39     train_acc, val_acc = results[lr]
40     print ('Learning rates: %e Train accuracy: %f Validation accuracy: %f' % (lr, train_acc, val_acc))
41
42 print ('Best validation accuracy: %f' % best_acc)
43
44 # Plot the loss function and train / validation accuracies
45 plt.plot(stats['loss_history'])
46 plt.title('Training Loss History')
47 plt.xlabel('Iteration')
48 plt.ylabel('Training Loss')
49 plt.show()
50
51 plt.plot(stats['train_acc_history'])
52 plt.plot(stats['val_acc_history'])
53 plt.title('Classification Accuracy History')
54 plt.xlabel('Iteration')
55 plt.ylabel('Accuracy')
56 plt.show()
57
58 # ===== #
59 # END YOUR CODE HERE
60 # ===== #
61 best_net = net

```

```

/Users/hannah_wang/Desktop/hw3/code/nndl/neural_net.py:120: RuntimeWarning: overflow encountered in exp
softmax_loss = np.sum(-np.log(np.exp(scores[np.arange(N), y]) / np.sum(np.exp(scores), axis = 1)))
/Users/hannah_wang/Desktop/hw3/code/nndl/neural_net.py:120: RuntimeWarning: invalid value encountered in true_divide
softmax_loss = np.sum(-np.log(np.exp(scores[np.arange(N), y]) / np.sum(np.exp(scores), axis = 1)))
/Users/hannah_wang/Desktop/hw3/code/nndl/neural_net.py:120: RuntimeWarning: divide by zero encountered in log
softmax_loss = np.sum(-np.log(np.exp(scores[np.arange(N), y]) / np.sum(np.exp(scores), axis = 1)))
/Users/hannah_wang/Desktop/hw3/code/nndl/neural_net.py:139: RuntimeWarning: overflow encountered in exp
score_exp = np.exp(scores)
/Users/hannah_wang/Desktop/hw3/code/nndl/neural_net.py:140: RuntimeWarning: invalid value encountered in true_divide
prob = score_exp / np.sum(score_exp, axis=1, keepdims=True)
/Users/hannah_wang/Desktop/hw3/code/nndl/neural_net.py:92: RuntimeWarning: invalid value encountered in maximum
relu = lambda x: np.maximum(0, x)
/Users/hannah_wang/Desktop/hw3/code/nndl/neural_net.py:143: RuntimeWarning: invalid value encountered in maximum
dldW2 = np.maximum(0, W1.dot(X.T) + b1.reshape([W1.shape[0], 1]))
/Users/hannah_wang/Desktop/hw3/code/nndl/neural_net.py:148: RuntimeWarning: invalid value encountered in greater
dlda = (W1.dot(X.T) > 0) * daydh.dot(dlday.T)
/Users/hannah_wang/Desktop/hw3/code/nndl/neural_net.py:266: RuntimeWarning: invalid value encountered in maximum
relu = np.maximum(0, ll)



```

```



Learning rates: 1.000000e-03 Train accuracy: 0.569878 Validation accuracy: 0.509000
Learning rates: 3.000000e-03 Train accuracy: 0.530102 Validation accuracy: 0.489000
Learning rates: 5.000000e-03 Train accuracy: 0.100265 Validation accuracy: 0.087000
Learning rates: 1.000000e-02 Train accuracy: 0.100265 Validation accuracy: 0.087000
Learning rates: 3.000000e-02 Train accuracy: 0.100265 Validation accuracy: 0.087000
Learning rates: 5.000000e-02 Train accuracy: 0.100265 Validation accuracy: 0.087000

```

Learning rates: 1.000000e-01 Train accuracy: 0.100265 Validation accuracy: 0.087000
Best validation accuracy: 0.509000



```
In [23]: 1 from cs231n.vis_utils import visualize_grid
2
3 # Visualize the weights of the network
4
5 def show_net_weights(net):
6     W1 = net.params['W1']
7     W1 = W1.T.reshape(32, 32, 3, -1).transpose(3, 0, 1, 2)
8     plt.imshow(visualize_grid(W1, padding=3).astype('uint8'))
9     plt.gca().axis('off')
10    plt.show()
11
12 show_net_weights(subopt_net)
13 show_net_weights(best_net)
```



Question:

(1) What differences do you see in the weights between the suboptimal net and the best net you arrived at?

Answer:

(1) The suboptimal net contains much more noises and we can hardly distinguish the differences among different images. While in the best net, we have more features to distinguish images.

Evaluate on test set

```
In [24]: 1 test_acc = (best_net.predict(X_test) == y_test).mean()  
         2 print('Test accuracy: ', test_acc)
```

Test accuracy: 0.507

```
In [ ]: 1
```

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 """
5 This code was originally written for CS 231n at Stanford University
6 (cs231n.stanford.edu). It has been modified in various areas for use in the
7 ECE 239AS class at UCLA. This includes the descriptions of what code to
8 implement as well as some slight potential changes in variable names to be
9 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung
10 for
11 permission to use this code. To see the original version, please visit
12 cs231n.stanford.edu.
13 """
14 class TwoLayerNet(object):
15     """
16     A two-layer fully-connected neural network. The net has an input dimension
17     of N, a hidden layer dimension of H, and performs classification over C
18     classes.
19     We train the network with a softmax loss function and L2 regularization on
20     the weight matrices. The network uses a ReLU nonlinearity after the first fully
21     connected layer.
22
23     In other words, the network has the following architecture:
24     input - fully connected layer - ReLU - fully connected layer - softmax
25
26     The outputs of the second fully-connected layer are the scores for each
27     class.
28     """
29     def __init__(self, input_size, hidden_size, output_size, std=1e-4):
30         """
31         Initialize the model. Weights are initialized to small random values and
32         biases are initialized to zero. Weights and biases are stored in the
33         variable self.params, which is a dictionary with the following keys:
34
35         W1: First layer weights; has shape (H, D)
36         b1: First layer biases; has shape (H,)
37         W2: Second layer weights; has shape (C, H)
38         b2: Second layer biases; has shape (C,)
39
40         Inputs:
41         - input_size: The dimension D of the input data.
42         - hidden_size: The number of neurons H in the hidden layer.
43         - output_size: The number of classes C.
44         """
45         self.params = {}
46         self.params['W1'] = std * np.random.randn(hidden_size, input_size)
47         self.params['b1'] = np.zeros(hidden_size)
48         self.params['W2'] = std * np.random.randn(output_size, hidden_size)
49         self.params['b2'] = np.zeros(output_size)
50
51     def loss(self, X, y=None, reg=0.0):
52         """
53         Compute the loss and gradients for a two layer fully connected neural
54         network.
55
56         Inputs:
57         - X: Input data of shape (N, D). Each X[i] is a training sample.
58         - y: Vector of training labels. y[i] is the label for X[i], and each y[i]
59 is
60         an integer in the range 0 <= y[i] < C. This parameter is optional; if
61 it
62         is not passed then we only return scores, and if it is passed then we
63         instead return the loss and gradients.
64         - reg: Regularization strength.
65
66         Returns:
67         If y is None, return a matrix scores of shape (N, C) where scores[i, c]
68 is
69         the score for class c on input X[i].
70
71         If y is not None, instead return a tuple of:
72         - loss: Loss (data loss and regularization loss) for this batch of
73 training
74         samples.
75         - grads: Dictionary mapping parameter names to gradients of those
76 parameters
77         with respect to the loss function; has the same keys as self.params.
78         """

```

```

75 # Unpack variables from the params dictionary
76 W1, b1 = self.params['W1'], self.params['b1']
77 W2, b2 = self.params['W2'], self.params['b2']
78 N, D = X.shape
79
80 # Compute the forward pass
81 scores = None
82
83 # ===== #
84 # YOUR CODE HERE:
85 # Calculate the output scores of the neural network. The result
86 # should be (C, N). As stated in the description for this class,
87 # there should not be a ReLU layer after the second FC layer.
88 # The output of the second FC layer is the output scores. Do not
89 # use a for loop in your implementation.
90 # ===== #
91
92 relu = lambda x: np.maximum(0, x)
93 h1 = np.dot(X, W1.T) + b1 # H * N
94 output = relu(h1)
95 h2 = np.dot(output, W2.T) + b2
96 scores = h2
97
98 # ===== #
99 # END YOUR CODE HERE
100 # ===== #
101
102
103 # If the targets are not given then jump out, we're done
104 if y is None:
105     return scores
106
107 # Compute the loss
108 loss = None
109
110 # ===== #
111 # YOUR CODE HERE:
112 # Calculate the loss of the neural network. This includes the
113 # softmax loss and the L2 regularization for W1 and W2. Store the
114 # total loss in the variable loss. Multiply the regularization
115 # loss by 0.5 (in addition to the factor reg).
116 # ===== #
117
118 # scores is num_examples by num_classes
119
120 softmax_loss = np.sum(-np.log(np.exp(scores[np.arange(N), y]) /
np.sum(np.exp(scores), axis = 1)))
121 # penalize by frobenius norm
122 reg_loss = 0.5 * reg * (np.linalg.norm(W1, 'fro')**2 + np.linalg.norm(W2,
'fro')**2)
123 loss = softmax_loss / N + reg_loss
124
125 # ===== #
126 # END YOUR CODE HERE
127 # ===== #
128
129 grads = {}
130
131 # ===== #
132 # YOUR CODE HERE:
133 # Implement the backward pass. Compute the derivatives of the
134 # weights and the biases. Store the results in the grads
135 # dictionary. e.g., grads['W1'] should store the gradient for
136 # W1, and be of the same size as W1.
137 # ===== #
138
139 score_exp = np.exp(scores)
140 prob = score_exp / np.sum(score_exp, axis=1, keepdims=True)
141 prob[np.arange(N), y] -= 1
142 dlday = prob / N
143 dldW2 = np.maximum(0, W1.dot(X.T) + b1.reshape([W1.shape[0], 1]))
144 grads['W2'] = dlday.T.dot(dldW2.T) + reg * W2
145 grads['b2'] = np.sum(dlday, axis=0, keepdims=True)
146
147 daydh = W2.T
148 dlda = (W1.dot(X.T) > 0) * daydh.dot(dlday.T)
149 grads['W1'] = dlda.dot(X) + reg * W1
150 grads['b1'] = np.sum(dlda, axis=1, keepdims=True).T
151
152 # ===== #
153 # END YOUR CODE HERE
154 # ===== #
155
156 return loss, grads

```

```

157
158 def train(self, X, y, X_val, y_val,
159           learning_rate=1e-3, learning_rate_decay=0.95,
160           reg=1e-5, num_iters=100,
161           batch_size=200, verbose=False):
162     """
163     Train this neural network using stochastic gradient descent.
164
165     Inputs:
166     - X: A numpy array of shape (N, D) giving training data.
167     - y: A numpy array of shape (N,) giving training labels; y[i] = c means
that X[i] has label c, where 0 ≤ c < C.
168     - X_val: A numpy array of shape (N_val, D) giving validation data.
169     - y_val: A numpy array of shape (N_val,) giving validation labels.
170     - learning_rate: Scalar giving learning rate for optimization.
171     - learning_rate_decay: Scalar giving factor used to decay the learning
rate
172     after each epoch.
173     - reg: Scalar giving regularization strength.
174     - num_iters: Number of steps to take when optimizing.
175     - batch_size: Number of training examples to use per step.
176     - verbose: boolean; if true print progress during optimization.
177     """
178     num_train = X.shape[0]
179     iterations_per_epoch = max(num_train / batch_size, 1)
180
181     # Use SGD to optimize the parameters in self.model
182     loss_history = []
183     train_acc_history = []
184     val_acc_history = []
185
186     for it in np.arange(num_iters):
187         X_batch = None
188         y_batch = None
189
190         # ===== #
191         # YOUR CODE HERE:
192         # Create a minibatch by sampling batch_size samples randomly.
193         # ===== #
194
195         idx = np.random.choice(num_train, batch_size)
196         X_batch = X[idx]
197         y_batch = y[idx]
198
199         # ===== #
200         # END YOUR CODE HERE
201         # ===== #
202
203         # Compute loss and gradients using the current minibatch
204         loss, grads = self.loss(X_batch, y=y_batch, reg=reg)
205         loss_history.append(loss)
206
207         # ===== #
208         # YOUR CODE HERE:
209         # Perform a gradient descent step using the minibatch to update
210         # all parameters (i.e., W1, W2, b1, and b2).
211         # ===== #
212
213         self.params['W1'] = self.params['W1'] - grads['W1'] * learning_rate
214         self.params['b1'] = self.params['b1'] - grads['b1'] * learning_rate
215         self.params['W2'] = self.params['W2'] - grads['W2'] * learning_rate
216         self.params['b2'] = self.params['b2'] - grads['b2'] * learning_rate
217
218         # ===== #
219         # END YOUR CODE HERE
220         # ===== #
221
222         if verbose and it % 100 == 0:
223             print('iteration {} / {}: loss {}'.format(it, num_iters, loss))
224
225         # Every epoch, check train and val accuracy and decay learning rate.
226         if it % iterations_per_epoch == 0:
227             # Check accuracy
228             train_acc = (self.predict(X_batch) == y_batch).mean()
229             val_acc = (self.predict(X_val) == y_val).mean()
230             train_acc_history.append(train_acc)
231             val_acc_history.append(val_acc)
232
233             # Decay learning rate
234             learning_rate *= learning_rate_decay
235
236     return {
237         'loss_history': loss_history,
238

```

```

239     'train_acc_history': train_acc_history,
240     'val_acc_history': val_acc_history,
241 }
242
243 def predict(self, X):
244     """
245     Use the trained weights of this two-layer network to predict labels for
246     data points. For each data point we predict scores for each of the C
247     classes, and assign each data point to the class with the highest score.
248
249     Inputs:
250     - X: A numpy array of shape (N, D) giving N D-dimensional data points to
251         classify.
252
253     Returns:
254     - y_pred: A numpy array of shape (N,) giving predicted labels for each of
255         the elements of X. For all i, y_pred[i] = c means that X[i] is
256         predicted
257         to have class c, where 0 <= c < C.
258     """
259     y_pred = None
260
261     # ===== #
262     # YOUR CODE HERE:
263     # Predict the class given the input data.
264     # ===== #
265
266     l1 = np.dot(X, self.params['W1'].T) + self.params['b1']
267     relu = np.maximum(0, l1)
268     scores = np.dot(relu, self.params['W2'].T) + self.params['b2']
269     y_pred = np.argmax(scores, axis=1)
270
271     # ===== #
272     # END YOUR CODE HERE
273     # ===== #
274
275     return y_pred
276
277

```


Fully connected networks

In the previous notebook, you implemented a simple two-layer neural network class. However, this class is not modular. If you wanted to change the number of layers, you would need to write a new loss and gradient function. If you wanted to optimize the network with different optimizers, you'd need to write new training functions. If you wanted to incorporate regularizations, you'd have to modify the loss and gradient function.

Instead of having to modify functions each time, for the rest of the class, we'll work in a more modular framework where we define forward and backward layers that calculate losses and gradients respectively. Since the forward and backward layers share intermediate values that are useful for calculating both the loss and the gradient, we'll also have these function return "caches" which store useful intermediate values.

The goal is that through this modular design, we can build different sized neural networks for various applications.

In this HW #3, we'll define the basic architecture, and in HW #4, we'll build on this framework to implement different optimizers and regularizations (like BatchNorm and Dropout).

CS231n has built a solid API for building these modular frameworks and training them, and we will use their very well implemented framework as opposed to "reinventing the wheel." This includes using their Solver, various utility functions, and their layer structure. This also includes `nndl.fc_net`, `nndl.layers`, and `nndl.layer_utils`. As in prior assignments, we thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class (cs231n.stanford.edu).

Modular layers

This notebook will build modular layers in the following manner. First, there will be a forward pass for a given layer with inputs (`x`) and return the output of that layer (`out`) as well as cached variables (`cache`) that will be used to calculate the gradient in the backward pass.

```
def layer_forward(x, w):
    """ Receive inputs x and weights w """
    # Do some computations ...
    z = # ... some intermediate value
    # Do some more computations ...
    out = # the output

    cache = (x, w, z, out) # Values we need to compute gradients

    return out, cache
```

The backward pass will receive upstream derivatives and the `cache` object, and will return gradients with respect to the inputs and weights, like this:

```
def layer_backward(dout, cache):
    """
    Receive derivative of loss with respect to outputs and cache,
    and compute derivative with respect to inputs.
    """
    # Unpack cache values
    x, w, z, out = cache

    # Use values in cache to compute derivatives
    dx = # Derivative of loss with respect to x
    dw = # Derivative of loss with respect to w

    return dx, dw
```

```
In [1]: 1 ## Import and setups
        2
        3 import time
        4 import numpy as np
        5 import matplotlib.pyplot as plt
        6 from nndl.fc_net import *
        7 from cs231n.data_utils import get_CIFAR10_data
        8 from cs231n.gradient_check import eval_numerical_gradient, eval_numerical_gradient_array
        9 from cs231n.solver import Solver
        10
        11 %matplotlib inline
        12 plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
        13 plt.rcParams['image.interpolation'] = 'nearest'
        14 plt.rcParams['image.cmap'] = 'gray'
        15
        16 # for auto-reloading external modules
        17 # see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
        18 %load_ext autoreload
        19 %autoreload 2
        20
        21 def rel_error(x, y):
        22     """ returns relative error """
        23     return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

```
In [2]: 1 # Load the (preprocessed) CIFAR10 data.
2
3 data = get_CIFAR10_data()
4 for k in data.keys():
5     print('{}: {}'.format(k, data[k].shape))
```

```
X_train: (49000, 3, 32, 32)
y_train: (49000,)
X_val: (1000, 3, 32, 32)
y_val: (1000,)
X_test: (1000, 3, 32, 32)
y_test: (1000,)
```

Linear layers

In this section, we'll implement the forward and backward pass for the linear layers.

The linear layer forward pass is the function `affine_forward` in `nn1/layers.py` and the backward pass is `affine_backward`.

After you have implemented these, test your implementation by running the cell below.

Affine layer forward pass

Implement `affine_forward` and then test your code by running the following cell.

```
In [3]: 1 # Test the affine_forward function
2
3 num_inputs = 2
4 input_shape = (4, 5, 6)
5 output_dim = 3
6
7 input_size = num_inputs * np.prod(input_shape)
8 weight_size = output_dim * np.prod(input_shape)
9
10 x = np.linspace(-0.1, 0.5, num=input_size).reshape(num_inputs, *input_shape)
11 w = np.linspace(-0.2, 0.3, num=weight_size).reshape(np.prod(input_shape), output_dim)
12 b = np.linspace(-0.3, 0.1, num=output_dim)
13
14 out, _ = affine_forward(x, w, b)
15 correct_out = np.array([[ 1.49834967,  1.70660132,  1.91485297],
16                        [ 3.25553199,  3.5141327,  3.77273342]])
17
18 # Compare your output with ours. The error should be around 1e-9.
19 print('Testing affine_forward function:')
20 print('difference: {}'.format(rel_error(out, correct_out)))
```

```
Testing affine_forward function:
difference: 9.769849468192957e-10
```

Affine layer backward pass

Implement `affine_backward` and then test your code by running the following cell.

```
In [4]: 1 # Test the affine_backward function
2
3 x = np.random.randn(10, 2, 3)
4 w = np.random.randn(6, 5)
5 b = np.random.randn(5)
6 dout = np.random.randn(10, 5)
7
8 dx_num = eval_numerical_gradient_array(lambda x: affine_forward(x, w, b)[0], x, dout)
9 dw_num = eval_numerical_gradient_array(lambda w: affine_forward(x, w, b)[0], w, dout)
10 db_num = eval_numerical_gradient_array(lambda b: affine_forward(x, w, b)[0], b, dout)
11
12 _, cache = affine_forward(x, w, b)
13 dx, dw, db = affine_backward(dout, cache)
14
15 # The error should be around 1e-10
16 print('Testing affine_backward function:')
17 print('dx error: {}'.format(rel_error(dx_num, dx)))
18 print('dw error: {}'.format(rel_error(dw_num, dw)))
19 print('db error: {}'.format(rel_error(db_num, db)))
```

```
Testing affine_backward function:
dx error: 3.010879138834073e-10
dw error: 2.2214518603736647e-11
db error: 1.9746463154109852e-11
```

Activation layers

In this section you'll implement the ReLU activation.

ReLU forward pass

Implement the `relu_forward` function in `nndl/layers.py` and then test your code by running the following cell.

```
In [5]: 1 # Test the relu_forward function
2
3 x = np.linspace(-0.5, 0.5, num=12).reshape(3, 4)
4
5 out, _ = relu_forward(x)
6 correct_out = np.array([[ 0.,          0.,          0.,          0.,          ],
7                          [ 0.,          0.,          0.04545455, 0.13636364, ],
8                          [ 0.22727273, 0.31818182, 0.40909091, 0.5,          ]])
9
10 # Compare your output with ours. The error should be around 1e-8
11 print('Testing relu_forward function:')
12 print('difference: {}'.format(rel_error(out, correct_out)))
```

Testing relu_forward function:
difference: 4.999999798022158e-08

ReLU backward pass

Implement the `relu_backward` function in `nndl/layers.py` and then test your code by running the following cell.

```
In [6]: 1 x = np.random.randn(10, 10)
2 dout = np.random.randn(*x.shape)
3
4 dx_num = eval_numerical_gradient_array(lambda x: relu_forward(x)[0], x, dout)
5
6 _, cache = relu_forward(x)
7 dx = relu_backward(dout, cache)
8
9 # The error should be around 1e-12
10 print('Testing relu_backward function:')
11 print('dx error: {}'.format(rel_error(dx_num, dx)))
```

Testing relu_backward function:
dx error: 3.2756100263351556e-12

Combining the affine and ReLU layers

Often times, an affine layer will be followed by a ReLU layer. So let's make one that puts them together. Layers that are combined are stored in `nndl/layer_utils.py`.

Affine-ReLU layers

We've implemented `affine_relu_forward()` and `affine_relu_backward` in `nndl/layer_utils.py`. Take a look at them to make sure you understand what's going on. Then run the following cell to ensure its implemented correctly.

```
In [7]: 1 from nndl.layer_utils import affine_relu_forward, affine_relu_backward
2
3 x = np.random.randn(2, 3, 4)
4 w = np.random.randn(12, 10)
5 b = np.random.randn(10)
6 dout = np.random.randn(2, 10)
7
8 out, cache = affine_relu_forward(x, w, b)
9 dx, dw, db = affine_relu_backward(dout, cache)
10
11 dx_num = eval_numerical_gradient_array(lambda x: affine_relu_forward(x, w, b)[0], x, dout)
12 dw_num = eval_numerical_gradient_array(lambda w: affine_relu_forward(x, w, b)[0], w, dout)
13 db_num = eval_numerical_gradient_array(lambda b: affine_relu_forward(x, w, b)[0], b, dout)
14
15 print('Testing affine_relu_forward and affine_relu_backward:')
16 print('dx error: {}'.format(rel_error(dx_num, dx)))
17 print('dw error: {}'.format(rel_error(dw_num, dw)))
18 print('db error: {}'.format(rel_error(db_num, db)))
```

Testing affine_relu_forward and affine_relu_backward:
dx error: 6.307295906050895e-10
dw error: 8.914200071695358e-11
db error: 7.826729437422599e-12

Softmax and SVM losses

You've already implemented these, so we have written these in `layers.py`. The following code will ensure they are working correctly.

```

In [8]: 1 num_classes, num_inputs = 10, 50
2 x = 0.001 * np.random.randn(num_inputs, num_classes)
3 y = np.random.randint(num_classes, size=num_inputs)
4
5 dx_num = eval_numerical_gradient(lambda x: svm_loss(x, y)[0], x, verbose=False)
6 loss, dx = svm_loss(x, y)
7
8 # Test svm_loss function. Loss should be around 9 and dx error should be 1e-9
9 print('Testing svm_loss:')
10 print('loss: {}'.format(loss))
11 print('dx error: {}'.format(rel_error(dx_num, dx)))
12
13 dx_num = eval_numerical_gradient(lambda x: softmax_loss(x, y)[0], x, verbose=False)
14 loss, dx = softmax_loss(x, y)
15
16 # Test softmax_loss function. Loss should be 2.3 and dx error should be 1e-8
17 print('\nTesting softmax_loss:')
18 print('loss: {}'.format(loss))
19 print('dx error: {}'.format(rel_error(dx_num, dx)))

```

```

Testing svm_loss:
loss: 9.001620019513892
dx error: 8.182894472887002e-10

```

```

Testing softmax_loss:
loss: 2.302747596966573
dx error: 8.426201588202499e-09

```

Implementation of a two-layer NN

In `nndl/fc_net.py`, implement the class `TwoLayerNet` which uses the layers you made here. When you have finished, the following cell will test your implementation.

```

In [9]: 1 N, D, H, C = 3, 5, 50, 7
2 X = np.random.randn(N, D)
3 y = np.random.randint(C, size=N)
4
5 std = 1e-2
6 model = TwoLayerNet(input_dim=D, hidden_dims=H, num_classes=C, weight_scale=std)
7
8 print('Testing initialization ... ')
9 W1_std = abs(model.params['W1'].std() - std)
10 b1 = model.params['b1']
11 W2_std = abs(model.params['W2'].std() - std)
12 b2 = model.params['b2']
13 assert W1_std < std / 10, 'First layer weights do not seem right'
14 assert np.all(b1 == 0), 'First layer biases do not seem right'
15 assert W2_std < std / 10, 'Second layer weights do not seem right'
16 assert np.all(b2 == 0), 'Second layer biases do not seem right'
17
18 print('Testing test-time forward pass ... ')
19 model.params['W1'] = np.linspace(-0.7, 0.3, num=D*H).reshape(D, H)
20 model.params['b1'] = np.linspace(-0.1, 0.9, num=H)
21 model.params['W2'] = np.linspace(-0.3, 0.4, num=H*C).reshape(H, C)
22 model.params['b2'] = np.linspace(-0.9, 0.1, num=C)
23 X = np.linspace(-5.5, 4.5, num=N*D).reshape(D, N).T
24 scores = model.loss(X)
25 correct_scores = np.asarray(
26     [[11.53165108, 12.2917344, 13.05181771, 13.81190102, 14.57198434, 15.33206765, 16.09215096],
27      [12.05769098, 12.74614105, 13.43459113, 14.1230412, 14.81149128, 15.49994135, 16.18839143],
28      [12.58373087, 13.20054771, 13.81736455, 14.43418138, 15.05099822, 15.66781506, 16.2846319 ]])
29 scores_diff = np.abs(scores - correct_scores).sum()
30 assert scores_diff < 1e-6, 'Problem with test-time forward pass'
31
32 print('Testing training loss (no regularization)')
33 y = np.asarray([0, 5, 1])
34 loss, grads = model.loss(X, y)
35 correct_loss = 3.4702243556
36 assert abs(loss - correct_loss) < 1e-10, 'Problem with training-time loss'
37
38 model.reg = 1.0
39 loss, grads = model.loss(X, y)
40 correct_loss = 26.5948426952
41 assert abs(loss - correct_loss) < 1e-10, 'Problem with regularization loss'
42
43 for reg in [0.0, 0.7]:
44     print('Running numeric gradient check with reg = {}'.format(reg))
45     model.reg = reg
46     loss, grads = model.loss(X, y)
47
48     for name in sorted(grads):
49         f = lambda _: model.loss(X, y)[0]
50         grad_num = eval_numerical_gradient(f, model.params[name], verbose=False)
51         print('{} relative error: {}'.format(name, rel_error(grad_num, grads[name])))

```

```

Testing initialization ...
Testing test-time forward pass ...
Testing training loss (no regularization)
Running numeric gradient check with reg = 0.0
W1 relative error: 1.8336562786695002e-08
W2 relative error: 3.201560569143183e-10
b1 relative error: 9.828320891054654e-09
b2 relative error: 1.0
Running numeric gradient check with reg = 0.7
W1 relative error: 2.5279152310200606e-07
W2 relative error: 7.976652806155026e-08
b1 relative error: 1.564680430849598e-08
b2 relative error: 1.0

```

Solver

We will now use the `cs231n Solver` class to train these networks. Familiarize yourself with the API in `cs231n/solver.py`. After you have done so, declare an instance of a `TwoLayerNet` with 200 units and then train it with the `Solver`. Choose parameters so that your validation accuracy is at least 40%.

```

In [10]: 1 model = TwoLayerNet()
2 solver = None
3
4 # ===== #
5 # YOUR CODE HERE:
6 #   Declare an instance of a TwoLayerNet and then train
7 #   it with the Solver. Choose hyperparameters so that your validation
8 #   accuracy is at least 40%. We won't have you optimize this further
9 #   since you did it in the previous notebook.
10 #
11 # ===== #
12
13 solver = Solver(model, data, update_rule='sgd',
14                 optim_config={
15                     'learning_rate': 1e-3,
16                 },
17                 lr_decay=0.95,
18                 num_epochs=10,
19                 batch_size=100,
20                 print_every=500
21             )
22 solver.train()
23
24 # ===== #
25 # END YOUR CODE HERE
26 # ===== #


```

```

(Iteration 1 / 4900) loss: 2.303671
(Epoch 0 / 10) train acc: 0.125000; val_acc: 0.128000
(Epoch 1 / 10) train acc: 0.423000; val_acc: 0.440000
(Iteration 501 / 4900) loss: 1.362025
(Epoch 2 / 10) train acc: 0.501000; val_acc: 0.473000
(Iteration 1001 / 4900) loss: 1.475222
(Epoch 3 / 10) train acc: 0.529000; val_acc: 0.495000
(Iteration 1501 / 4900) loss: 1.363798
(Epoch 4 / 10) train acc: 0.568000; val_acc: 0.506000
(Iteration 2001 / 4900) loss: 1.059229
(Epoch 5 / 10) train acc: 0.496000; val_acc: 0.494000
(Iteration 2501 / 4900) loss: 1.183275
(Epoch 6 / 10) train acc: 0.565000; val_acc: 0.480000
(Iteration 3001 / 4900) loss: 1.448210
(Epoch 7 / 10) train acc: 0.590000; val_acc: 0.522000
(Iteration 3501 / 4900) loss: 1.182251
(Epoch 8 / 10) train acc: 0.616000; val_acc: 0.506000
(Iteration 4001 / 4900) loss: 1.218148
(Epoch 9 / 10) train acc: 0.608000; val_acc: 0.501000
(Iteration 4501 / 4900) loss: 0.964236
(Epoch 10 / 10) train acc: 0.618000; val_acc: 0.507000

```

```
In [11]: 1 # Run this cell to visualize training loss and train / val accuracy
2
3 plt.subplot(2, 1, 1)
4 plt.title('Training loss')
5 plt.plot(solver.loss_history, 'o')
6 plt.xlabel('Iteration')
7
8 plt.subplot(2, 1, 2)
9 plt.title('Accuracy')
10 plt.plot(solver.train_acc_history, '-o', label='train')
11 plt.plot(solver.val_acc_history, '-o', label='val')
12 plt.plot([0.5] * len(solver.val_acc_history), 'k--')
13 plt.xlabel('Epoch')
14 plt.legend(loc='lower right')
15 plt.gcf().set_size_inches(15, 12)
16 plt.show()
```



Multilayer Neural Network

Now, we implement a multi-layer neural network.

Read through the `FullyConnectedNet` class in the file `nnd1/fc_net.py`.

Implement the initialization, the forward pass, and the backward pass. There will be lines for batchnorm and dropout layers and caches; ignore these all for now. That'll be in assignment #4.


```

In [12]: 1 N, D, H1, H2, C = 2, 15, 20, 30, 10
          2 X = np.random.randn(N, D)
          3 y = np.random.randint(C, size=(N,))
          4
          5 for reg in [0, 3.14]:
          6     print('Running check with reg = {}'.format(reg))
          7     model = FullyConnectedNet([H1, H2], input_dim=D, num_classes=C,
          8                               reg=reg, weight_scale=5e-2, dtype=np.float64)
          9
          10    loss, grads = model.loss(X, y)
          11    print('Initial loss: {}'.format(loss))
          12
          13    for name in sorted(grads):
          14        f = lambda _: model.loss(X, y)[0]
          15        grad_num = eval_numerical_gradient(f, model.params[name], verbose=False, h=1e-5)
          16        print('{} relative error: {}'.format(name, rel_error(grad_num, grads[name])))

```

```

Running check with reg = 0
Initial loss: 2.2946741792511536
W1 relative error: 5.375684907457941e-07
W2 relative error: 3.0460898865864483e-07
W3 relative error: 4.017042458877704e-08
b1 relative error: 3.8902340093824456e-09
b2 relative error: 5.754520470019443e-09
b3 relative error: 9.008473410971608e-11
Running check with reg = 3.14
Initial loss: 6.954350043636696
W1 relative error: 1.7058425334773003e-07
W2 relative error: 2.0127817584974917e-08
W3 relative error: 2.2882986739267522e-08
b1 relative error: 1.509307878869427e-07
b2 relative error: 1.264727808089801e-08
b3 relative error: 2.036894054330387e-10

```

```


In [14]: 1 # Use the three layer neural network to overfit a small dataset.
2
3 num_train = 50
4 small_data = {
5     'X_train': data['X_train'][:num_train],
6     'y_train': data['y_train'][:num_train],
7     'X_val': data['X_val'],
8     'y_val': data['y_val'],
9 }
10
11
12 ##### !!!!!
13 # Play around with the weight_scale and learning_rate so that you can overfit a small dataset.
14 # Your training accuracy should be 1.0 to receive full credit on this part.
15 weight_scale = 2e-2
16 learning_rate = 3e-3
17
18 model = FullyConnectedNet([100, 100],
19                             weight_scale=weight_scale, dtype=np.float64)
20 solver = Solver(model, small_data,
21                 print_every=10, num_epochs=20, batch_size=25,
22                 update_rule='sgd',
23                 optim_config={
24                     'learning_rate': learning_rate,
25                 })
26 solver.train()
27
28
29 plt.plot(solver.loss_history, 'o')
30 plt.title('Training loss history')
31 plt.xlabel('Iteration')
32 plt.ylabel('Training loss')
33 plt.show()

```

```

(Iteration 1 / 40) loss: 2.987310
(Epoch 0 / 20) train acc: 0.340000; val_acc: 0.137000
(Epoch 1 / 20) train acc: 0.320000; val_acc: 0.125000
(Epoch 2 / 20) train acc: 0.520000; val_acc: 0.124000
(Epoch 3 / 20) train acc: 0.580000; val_acc: 0.134000
(Epoch 4 / 20) train acc: 0.840000; val_acc: 0.156000
(Epoch 5 / 20) train acc: 0.880000; val_acc: 0.163000
(Iteration 11 / 40) loss: 1.034940
(Epoch 6 / 20) train acc: 0.980000; val_acc: 0.142000
(Epoch 7 / 20) train acc: 0.960000; val_acc: 0.129000
(Epoch 8 / 20) train acc: 0.980000; val_acc: 0.154000
(Epoch 9 / 20) train acc: 1.000000; val_acc: 0.156000
(Epoch 10 / 20) train acc: 1.000000; val_acc: 0.162000
(Iteration 21 / 40) loss: 0.180548
(Epoch 11 / 20) train acc: 1.000000; val_acc: 0.163000
(Epoch 12 / 20) train acc: 1.000000; val_acc: 0.173000
(Epoch 13 / 20) train acc: 1.000000; val_acc: 0.166000
(Epoch 14 / 20) train acc: 1.000000; val_acc: 0.176000
(Epoch 15 / 20) train acc: 1.000000; val_acc: 0.165000
(Iteration 31 / 40) loss: 0.104919
(Epoch 16 / 20) train acc: 1.000000; val_acc: 0.167000
(Epoch 17 / 20) train acc: 1.000000; val_acc: 0.170000
(Epoch 18 / 20) train acc: 1.000000; val_acc: 0.167000
(Epoch 19 / 20) train acc: 1.000000; val_acc: 0.162000
(Epoch 20 / 20) train acc: 1.000000; val_acc: 0.171000

```



```
In [ ]: 1
```

```
In [ ]: 1
```

```

1 import numpy as np
2 import pdb
3
4 """
5 This code was originally written for CS 231n at Stanford University
6 (cs231n.stanford.edu). It has been modified in various areas for use in the
7 ECE 239AS class at UCLA. This includes the descriptions of what code to
8 implement as well as some slight potential changes in variable names to be
9 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung
10 for
11 permission to use this code. To see the original version, please visit
12 cs231n.stanford.edu.
13 """
14
15 def affine_forward(x, w, b):
16     """
17     Computes the forward pass for an affine (fully-connected) layer.
18
19     The input x has shape (N, d_1, ..., d_k) and contains a minibatch of N
20     examples, where each example x[i] has shape (d_1, ..., d_k). We will
21     reshape each input into a vector of dimension D = d_1 * ... * d_k, and
22     then transform it to an output vector of dimension M.
23
24     Inputs:
25     - x: A numpy array containing input data, of shape (N, d_1, ..., d_k)
26     - w: A numpy array of weights, of shape (D, M)
27     - b: A numpy array of biases, of shape (M,)
28
29     Returns a tuple of:
30     - out: output, of shape (N, M)
31     - cache: (x, w, b)
32     """
33
34     # ===== #
35     # YOUR CODE HERE:
36     # Calculate the output of the forward pass. Notice the dimensions
37     # of w are D x M, which is the transpose of what we did in earlier
38     # assignments.
39     # ===== #
40
41     x_reshape = x.reshape((x.shape[0], -1)) # N * D
42     out = x_reshape.dot(w) + b.reshape((1, b.shape[0])) # N * M
43
44     # ===== #
45     # END YOUR CODE HERE
46     # ===== #
47
48     cache = (x, w, b)
49     return out, cache
50
51 def affine_backward(dout, cache):
52     """
53     Computes the backward pass for an affine layer.
54
55     Inputs:
56     - dout: Upstream derivative, of shape (N, M)
57     - cache: Tuple of:
58       - x: Input data, of shape (N, d_1, ... d_k)
59       - w: Weights, of shape (D, M)
60
61     Returns a tuple of:
62     - dx: Gradient with respect to x, of shape (N, d_1, ..., d_k)
63     - dw: Gradient with respect to w, of shape (D, M)
64     - db: Gradient with respect to b, of shape (M,)
65     """
66
67     x, w, b = cache
68     dx, dw, db = None, None, None
69
70     # ===== #
71     # YOUR CODE HERE:
72     # Calculate the gradients for the backward pass.
73     # ===== #
74
75     x_reshape = np.reshape(x, (x.shape[0], -1))
76     dx_reshape = dout.dot(w.T)
77     dx = np.reshape(dx_reshape, x.shape) # N * D
78     dw = x_reshape.T.dot(dout) # D * M
79     db = np.sum(dout.T, axis=1, keepdims=True).T # M * 1
80
81     # ===== #
82     # END YOUR CODE HERE
83     # ===== #

```

```

84
85     return dx, dw, db
86
87 def relu_forward(x):
88     """
89     Computes the forward pass for a layer of rectified linear units (ReLU).
90
91     Input:
92     - x: Inputs, of any shape
93
94     Returns a tuple of:
95     - out: Output, of the same shape as x
96     - cache: x
97     """
98     # ===== #
99     # YOUR CODE HERE:
100    #     Implement the ReLU forward pass.
101    # ===== #
102
103    out = np.maximum(0, x)
104
105    # ===== #
106    # END YOUR CODE HERE
107    # ===== #
108
109    cache = x
110    return out, cache
111
112
113 def relu_backward(dout, cache):
114     """
115     Computes the backward pass for a layer of rectified linear units (ReLU).
116
117     Input:
118     - dout: Upstream derivatives, of any shape
119     - cache: Input x, of same shape as dout
120
121     Returns:
122     - dx: Gradient with respect to x
123     """
124     x = cache
125
126     # ===== #
127     # YOUR CODE HERE:
128     #     Implement the ReLU backward pass
129     # ===== #
130
131     dx = (x > 0) * (dout)
132
133     # ===== #
134     # END YOUR CODE HERE
135     # ===== #
136
137     return dx
138
139 def svm_loss(x, y):
140     """
141     Computes the loss and gradient using for multiclass SVM classification.
142
143     Inputs:
144     - x: Input data, of shape (N, C) where x[i, j] is the score for the jth
145     class
146     for the ith input.
147     - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
148     0 <= y[i] < C
149
150     Returns a tuple of:
151     - loss: Scalar giving the loss
152     - dx: Gradient of the loss with respect to x
153     """
154     N = x.shape[0]
155     correct_class_scores = x[np.arange(N), y]
156     margins = np.maximum(0, x - correct_class_scores[:, np.newaxis] + 1.0)
157     margins[np.arange(N), y] = 0
158     loss = np.sum(margins) / N
159     num_pos = np.sum(margins > 0, axis=1)
160     dx = np.zeros_like(x)
161     dx[margins > 0] = 1
162     dx[np.arange(N), y] -= num_pos
163     dx /= N
164     return loss, dx
165
166 def softmax_loss(x, y):

```

```

167 """
168 Computes the loss and gradient for softmax classification.
169
170 Inputs:
171 - x: Input data, of shape (N, C) where x[i, j] is the score for the jth
class
172     for the ith input.
173 - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
174     0 <= y[i] < C
175
176 Returns a tuple of:
177 - loss: Scalar giving the loss
178 - dx: Gradient of the loss with respect to x
179 """
180
181 probs = np.exp(x - np.max(x, axis=1, keepdims=True))
182 probs /= np.sum(probs, axis=1, keepdims=True)
183 N = x.shape[0]
184 loss = -np.sum(np.log(probs[np.arange(N), y])) / N
185 dx = probs.copy()
186 dx[np.arange(N), y] -= 1
187 dx /= N
188 return loss, dx
189

```

```

1 import numpy as np
2
3 from .layers import *
4 from .layer_utils import *
5
6 """
7 This code was originally written for CS 231n at Stanford University
8 (cs231n.stanford.edu). It has been modified in various areas for use in the
9 ECE 239AS class at UCLA. This includes the descriptions of what code to
10 implement as well as some slight potential changes in variable names to be
11 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung
12 for
13 permission to use this code. To see the original version, please visit
14 cs231n.stanford.edu.
15 """
16 class TwoLayerNet(object):
17     """
18     A two-layer fully-connected neural network with ReLU nonlinearity and
19     softmax loss that uses a modular layer design. We assume an input dimension
20     of D, a hidden dimension of H, and perform classification over C classes.
21
22     The architecture should be affine - relu - affine - softmax.
23
24     Note that this class does not implement gradient descent; instead, it
25     will interact with a separate Solver object that is responsible for running
26     optimization.
27
28     The learnable parameters of the model are stored in the dictionary
29     self.params that maps parameter names to numpy arrays.
30     """
31
32     def __init__(self, input_dim=3*32*32, hidden_dims=100, num_classes=10,
33                 dropout=0, weight_scale=1e-3, reg=0.0):
34         """
35         Initialize a new network.
36
37         Inputs:
38         - input_dim: An integer giving the size of the input
39         - hidden_dims: An integer giving the size of the hidden layer
40         - num_classes: An integer giving the number of classes to classify
41         - dropout: Scalar between 0 and 1 giving dropout strength.
42         - weight_scale: Scalar giving the standard deviation for random
43           initialization of the weights.
44         - reg: Scalar giving L2 regularization strength.
45         """
46         self.params = {}
47         self.reg = reg
48
49         # ===== #
50         # YOUR CODE HERE:
51         # Initialize W1, W2, b1, and b2. Store these as self.params['W1'],
52         # self.params['W2'], self.params['b1'] and self.params['b2']. The
53         # biases are initialized to zero and the weights are initialized
54         # so that each parameter has mean 0 and standard deviation
55         # weight_scale.
56         # The dimensions of W1 should be (input_dim, hidden_dim) and the
57         # dimensions of W2 should be (hidden_dims, num_classes)
58         # ===== #
59         self.params['W1'] = np.random.randn(input_dim, hidden_dims) *
60 weight_scale
61 self.params['W2'] = np.random.randn(hidden_dims, num_classes) *
62 weight_scale
63 self.params['b1'] = np.zeros((hidden_dims, 1))
64 self.params['b2'] = np.zeros((num_classes, 1))
65
66 # ===== #
67 # END YOUR CODE HERE
68 # ===== #
69
70 def loss(self, X, y=None):
71     """
72     Compute loss and gradient for a minibatch of data.
73
74     Inputs:
75     - X: Array of input data of shape (N, d_1, ..., d_k)
76     - y: Array of labels, of shape (N,). y[i] gives the label for X[i].
77
78     Returns:
79     If y is None, then run a test-time forward pass of the model and return:
80     - scores: Array of shape (N, C) giving classification scores, where
81       scores[i, c] is the classification score for X[i] and class c.

```

```

81     If y is not None, then run a training-time forward and backward pass and
82     return a tuple of:
83     - loss: Scalar value giving the loss
84     - grads: Dictionary with the same keys as self.params, mapping parameter
85       names to gradients of the loss with respect to those parameters.
86     """
87     scores = None
88
89     # ===== #
90     # YOUR CODE HERE:
91     # Implement the forward pass of the two-layer neural network. Store
92     # the class scores as the variable 'scores'. Be sure to use the layers
93     # you prior implemented.
94     # ===== #
95
96     out1, cache1 = affine_forward(X, self.params['W1'], self.params['b1'])
97     out2, cache2 = relu_forward(out1)
98     scores, cache3 = affine_forward(out2, self.params['W2'],
self.params['b2'])
99
100    # ===== #
101    # END YOUR CODE HERE
102    # ===== #
103
104    # If y is None then we are in test mode so just return scores
105    if y is None:
106        return scores
107
108    loss, grads = 0, {}
109    # ===== #
110    # YOUR CODE HERE:
111    # Implement the backward pass of the two-layer neural net. Store
112    # the loss as the variable 'loss' and store the gradients in the
113    # 'grads' dictionary. For the grads dictionary, grads['W1'] holds
114    # the gradient for W1, grads['b1'] holds the gradient for b1, etc.
115    # i.e., grads[k] holds the gradient for self.params[k].
116    #
117    # Add L2 regularization, where there is an added cost  $0.5 * \text{self.reg} * W^2$ 
118    # for each W. Be sure to include the 0.5 multiplying factor to
119    # match our implementation.
120    #
121    # And be sure to use the layers you prior implemented.
122    # ===== #
123
124    loss, dx = softmax_loss(scores, y)
125    loss += 0.5 * self.reg * (np.linalg.norm(self.params['W1'], 'fro')**2 +
np.linalg.norm(self.params['W2'], 'fro')**2)
126
127    dh1, dw2, db2 = affine_backward(dx, cache3)
128    da1 = relu_backward(dh1, cache2)
129    dx2, dw1, db1 = affine_backward(da1, cache1)
130
131    grads['W1'] = dw1 + self.reg * self.params['W1']
132    grads['b1'] = db1.T
133    grads['W2'] = dw2 + self.reg * self.params['W2']
134    grads['b2'] = db2.T
135
136    # ===== #
137    # END YOUR CODE HERE
138    # ===== #
139
140    return loss, grads
141
142
143 class FullyConnectedNet(object):
144     """
145     A fully-connected neural network with an arbitrary number of hidden layers,
146     ReLU nonlinearities, and a softmax loss function. This will also implement
147     dropout and batch normalization as options. For a network with L layers,
148     the architecture will be
149
150     {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax
151
152     where batch normalization and dropout are optional, and the {...} block is
153     repeated L - 1 times.
154
155     Similar to the TwoLayerNet above, learnable parameters are stored in the
156     self.params dictionary and will be learned using the Solver class.
157     """
158
159     def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
160                   dropout=0, use_batchnorm=False, reg=0.0,
161                   weight_scale=1e-2, dtype=np.float32, seed=None):
162         """

```



```

163 Initialize a new FullyConnectedNet.
164
165 Inputs:
166 - hidden_dims: A list of integers giving the size of each hidden layer.
167 - input_dim: An integer giving the size of the input.
168 - num_classes: An integer giving the number of classes to classify.
169 - dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0
then
170     the network should not use dropout at all.
171 - use_batchnorm: Whether or not the network should use batch
normalization.
172 - reg: Scalar giving L2 regularization strength.
173 - weight_scale: Scalar giving the standard deviation for random
174 initialization of the weights.
175 - dtype: A numpy datatype object; all computations will be performed
using
176     this datatype. float32 is faster but less accurate, so you should use
177     float64 for numeric gradient checking.
178 - seed: If not None, then pass this random seed to the dropout layers.
This
179     will make the dropout layers deterministic so we can gradient check the
180     model.
181 """
182 self.use_batchnorm = use_batchnorm
183 self.use_dropout = dropout > 0
184 self.reg = reg
185 self.num_layers = 1 + len(hidden_dims)
186 self.dtype = dtype
187 self.params = {}
188
189 # ===== #
190 # YOUR CODE HERE:
191 # Initialize all parameters of the network in the self.params
dictionary.
192 # The weights and biases of layer 1 are W1 and b1; and in general the
193 # weights and biases of layer i are Wi and bi. The
194 # biases are initialized to zero and the weights are initialized
195 # so that each parameter has mean 0 and standard deviation
weight_scale.
196 # ===== #
197
198 dims = [input_dim] + hidden_dims + [num_classes]
199 for i in range(self.num_layers):
200     self.params['W' + str(i + 1)] = np.random.normal(0, weight_scale, size=
(dims[i], dims[i + 1]))
201     self.params['b' + str(i + 1)] = np.zeros((dims[i + 1], 1))
202
203 # ===== #
204 # END YOUR CODE HERE
205 # ===== #
206
207 # When using dropout we need to pass a dropout_param dictionary to each
208 # dropout layer so that the layer knows the dropout probability and the
mode
209 # (train / test). You can pass the same dropout_param to each dropout
layer.
210 self.dropout_param = {}
211 if self.use_dropout:
212     self.dropout_param = {'mode': 'train', 'p': dropout}
213     if seed is not None:
214         self.dropout_param['seed'] = seed
215
216 # With batch normalization we need to keep track of running means and
217 # variances, so we need to pass a special bn_param object to each batch
218 # normalization layer. You should pass self.bn_params[0] to the forward
pass
219 # of the first batch normalization layer, self.bn_params[1] to the
forward
220 # pass of the second batch normalization layer, etc.
221 self.bn_params = []
222 if self.use_batchnorm:
223     self.bn_params = [{'mode': 'train'} for i in np.arange(self.num_layers
- 1)]
224
225 # Cast all parameters to the correct datatype
226 for k, v in self.params.items():
227     self.params[k] = v.astype(dtype)
228
229
230 def loss(self, X, y=None):
231     """
232     Compute loss and gradient for the fully-connected net.
233
234     Input / output: Same as TwoLayerNet above.

```

```

235 """
236 X = X.astype(self.dtype)
237 mode = 'test' if y is None else 'train'
238
239 # Set train/test mode for batchnorm params and dropout param since they
240 # behave differently during training and testing.
241 if self.dropout_param is not None:
242     self.dropout_param['mode'] = mode
243 if self.use_batchnorm:
244     for bn_param in self.bn_params:
245         bn_param[mode] = mode
246
247 scores = None
248
249 # ===== #
250 # YOUR CODE HERE:
251 # Implement the forward pass of the FC net and store the output
252 # scores as the variable "scores".
253 # ===== #
254
255 param = {}
256 h = {}
257 h[0] = [X]
258
259 for i in range(self.num_layers):
260     param[i + 1] = affine_forward(h[i][0], self.params['W' + str(i + 1)],
self.params['b' + str(i + 1)])
261     h[i + 1] = relu_forward(param[i + 1][0])
262
263 scores = param[self.num_layers][0]
264
265 # ===== #
266 # END YOUR CODE HERE
267 # ===== #
268
269 # If test mode return early
270 if mode == 'test':
271     return scores
272
273 loss, grads = 0.0, {}
274 # ===== #
275 # YOUR CODE HERE:
276 # Implement the backwards pass of the FC net and store the gradients
277 # in the grads dict, so that grads[k] is the gradient of self.params[k]
278 # Be sure your L2 regularization includes a 0.5 factor.
279 # ===== #
280
281 loss, dx = softmax_loss(scores, y)
282 weights = [self.params['W' + str(i + 1)] for i in range(self.num_layers)]
283 loss += 0.5 * self.reg * sum([np.linalg.norm(weight, 'fro')**2 for weight
in weights])
284
285 das = {}
286 dhs = {}
287 dws = {}
288 dbs = {}
289 das[self.num_layers] = dx
290
291 for i in range(self.num_layers)[::-1]:
292     dh, dw, db = affine_backward(das[i + 1], param[i + 1][1])
293     dhs[i] = dh
294     dws[i + 1] = dw
295     dbs[i + 1] = db
296     if i != 0:
297         das[i] = relu_backward(dhs[i], h[i][1])
298
299 for i in range(self.num_layers):
300     grads['W' + str(i + 1)] = dws[i + 1] + self.reg * self.params['W' +
str(i + 1)]
301     grads['b' + str(i + 1)] = dbs[i + 1].T
302
303 # ===== #
304 # END YOUR CODE HERE
305 # ===== #
306 return loss, grads
307

```