

## Optimization for Fully Connected Networks

In this notebook, we will implement different optimization rules for gradient descent. We have provided starter code; however, you will need to copy and paste your code from your implementation of the modular fully connected nets in HW #3 to build upon this.

If you did not complete `affine_forward` and `backwards` passes, or `relu_forward` and `backward` passes from HW #3 correctly, you may use another classmate's implementation of these functions for this assignment, or contact us at [ece239as.w18@gmail.com](mailto:ece239as.w18@gmail.com).

CS231n has built a solid API for building these modular frameworks and training them, and we will use their very well implemented framework as opposed to "reinventing the wheel." This includes using their Solver, various utility functions, and their layer structure. This also includes `nndl.fc_net`, `nndl.layers`, and `nndl.layer_utils`. As in prior assignments, we thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class ([cs231n.stanford.edu](http://cs231n.stanford.edu)).

```
In [1]: 1  ## Import and setups
2
3  import time
4  import numpy as np
5  import matplotlib.pyplot as plt
6  from nndl.fc_net import *
7  from cs231n.data_utils import get_CIFAR10_data
8  from cs231n.gradient_check import eval_numerical_gradient, eval_numerical_gradient_array
9  from cs231n.solver import Solver
10
11  %matplotlib inline
12  plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
13  plt.rcParams['image.interpolation'] = 'nearest'
14  plt.rcParams['image.cmap'] = 'gray'
15
16  # for auto-reloading external modules
17  # see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
18  %load_ext autoreload
19  %autoreload 2
20
21  def rel_error(x, y):
22      """ returns relative error """
23      return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

```
In [2]: 1  # Load the (preprocessed) CIFAR10 data.
2
3  data = get_CIFAR10_data()
4  for k in data.keys():
5      print('{}: {}'.format(k, data[k].shape))
```

```
X_train: (49000, 3, 32, 32)
y_train: (49000,)
X_val: (1000, 3, 32, 32)
y_val: (1000,)
X_test: (1000, 3, 32, 32)
y_test: (1000,)
```

## Building upon your HW #3 implementation

Copy and paste the following functions from your HW #3 implementation of a modular FC net:

- `affine_forward` in `nndl/layers.py`
- `affine_backward` in `nndl/layers.py`
- `relu_forward` in `nndl/layers.py`
- `relu_backward` in `nndl/layers.py`
- `affine_relu_forward` in `nndl/layer_utils.py`
- `affine_relu_backward` in `nndl/layer_utils.py`
- The `FullyConnectedNet` class in `nndl/fc_net.py`

**Test all functions you copy and pasted**

```
In [5]: 1 from nndl.layer_tests import *
2
3 affine_forward_test(); print('\n')
4 affine_backward_test(); print('\n')
5 relu_forward_test(); print('\n')
6 relu_backward_test(); print('\n')
7 affine_relu_test(); print('\n')
8 fc_net_test()
```

If affine\_forward function is working, difference should be less than 1e-9:  
 difference: 9.769849468192957e-10

If affine\_backward is working, error should be less than 1e-9::  
 dx error: 3.666950908980358e-10  
 dw error: 8.201783430817534e-10  
 db error: 1.6263680637581037e-11

If relu\_forward function is working, difference should be around 1e-8:  
 difference: 4.999999798022158e-08

If relu\_backward function is working, error should be less than 1e-9:  
 dx error: 3.2755836707832798e-12

If affine\_relu\_forward and affine\_relu\_backward are working, error should be less than 1e-9::  
 dx error: 5.543539459666791e-11  
 dw error: 2.363803633156294e-10  
 db error: 7.826635255953652e-12

Running check with reg = 0  
 Initial loss: 2.308535257408285  
 W1 relative error: 7.372336841535986e-07  
 W2 relative error: 3.66674068452486e-07  
 W3 relative error: 1.343554782509405e-07  
 b1 relative error: 1.4252530105730905e-08  
 b2 relative error: 5.616642824420512e-09  
 b3 relative error: 1.1821438210204362e-10  
 Running check with reg = 3.14  
 Initial loss: 7.3016750389748655  
 W1 relative error: 0.004730420797436833  
 W2 relative error: 0.33008883827369184  
 W3 relative error: 5.44330801643248e-08  
 b1 relative error: 1.742017280608633e-08  
 b2 relative error: 0.9602337605136528  
 b3 relative error: 4.064532921805321e-10

## Training a larger model

In general, proceeding with vanilla stochastic gradient descent to optimize models may be fraught with problems and limitations, as discussed in class. Thus, we implement optimizers that improve on SGD.

### SGD + momentum

In the following section, implement SGD with momentum. Read the `nndl/optim.py` API, which is provided by CS231n, and be sure you understand it. After, implement `sgd_momentum` in `nndl/optim.py`. Test your implementation of `sgd_momentum` by running the cell below.

```
In [6]: 1 from nndl.optim import sgd_momentum
2
3 N, D = 4, 5
4 w = np.linspace(-0.4, 0.6, num=N*D).reshape(N, D)
5 dw = np.linspace(-0.6, 0.4, num=N*D).reshape(N, D)
6 v = np.linspace(0.6, 0.9, num=N*D).reshape(N, D)
7
8 config = {'learning_rate': 1e-3, 'velocity': v}
9 next_w, _ = sgd_momentum(w, dw, config=config)
10
11 expected_next_w = np.asarray([
12     [ 0.1406,      0.20738947,  0.27417895,  0.34096842,  0.40775789],
13     [ 0.47454737,  0.54133684,  0.60812632,  0.67491579,  0.74170526],
14     [ 0.80849474,  0.87528421,  0.94207368,  1.00886316,  1.07565263],
15     [ 1.14244211,  1.20923158,  1.27602105,  1.34281053,  1.4096      ]])
16 expected_velocity = np.asarray([
17     [ 0.5406,      0.55475789,  0.56891579,  0.58307368,  0.59723158],
18     [ 0.61138947,  0.62554737,  0.63970526,  0.65386316,  0.66802105],
19     [ 0.68217895,  0.69633684,  0.71049474,  0.72465263,  0.73881053],
20     [ 0.75296842,  0.76712632,  0.78128421,  0.79544211,  0.8096      ]])
21
22 print('next_w error: {}'.format(rel_error(next_w, expected_next_w)))
23 print('velocity error: {}'.format(rel_error(expected_velocity, config['velocity'])))
```

```
next_w error: 8.882347033505819e-09
velocity error: 4.269287743278663e-09
```

## SGD + Nesterov momentum

Implement `sgd_nesterov_momentum` in `nndl/optim.py`.

```
In [7]: 1 from nndl.optim import sgd_nesterov_momentum
2
3 N, D = 4, 5
4 w = np.linspace(-0.4, 0.6, num=N*D).reshape(N, D)
5 dw = np.linspace(-0.6, 0.4, num=N*D).reshape(N, D)
6 v = np.linspace(0.6, 0.9, num=N*D).reshape(N, D)
7
8 config = {'learning_rate': 1e-3, 'velocity': v}
9 next_w, _ = sgd_nesterov_momentum(w, dw, config=config)
10
11 expected_next_w = np.asarray([
12     [0.08714,      0.15246105,  0.21778211,  0.28310316,  0.34842421],
13     [0.41374526,  0.47906632,  0.54438737,  0.60970842,  0.67502947],
14     [0.74035053,  0.80567158,  0.87099263,  0.93631368,  1.00163474],
15     [1.06695579,  1.13227684,  1.19759789,  1.26291895,  1.32824      ]])
16 expected_velocity = np.asarray([
17     [ 0.5406,      0.55475789,  0.56891579,  0.58307368,  0.59723158],
18     [ 0.61138947,  0.62554737,  0.63970526,  0.65386316,  0.66802105],
19     [ 0.68217895,  0.69633684,  0.71049474,  0.72465263,  0.73881053],
20     [ 0.75296842,  0.76712632,  0.78128421,  0.79544211,  0.8096      ]])
21
22 print('next_w error: {}'.format(rel_error(next_w, expected_next_w)))
23 print('velocity error: {}'.format(rel_error(expected_velocity, config['velocity'])))
```

```
next_w error: 1.0875186845081027e-08
velocity error: 4.269287743278663e-09
```

## Evaluating SGD, SGD+Momentum, and SGD+NesterovMomentum

Run the following cell to train a 6 layer FC net with SGD, SGD+momentum, and SGD+Nesterov momentum. You should see that SGD+momentum achieves a better loss than SGD, and that SGD+Nesterov momentum achieves a slightly better loss (and training accuracy) than SGD+momentum.

```


In [8]: 1 num_train = 4000
2 small_data = {
3     'X_train': data['X_train'][:num_train],
4     'y_train': data['y_train'][:num_train],
5     'X_val': data['X_val'],
6     'y_val': data['y_val'],
7 }
8
9 solvers = {}
10
11 for update_rule in ['sgd', 'sgd_momentum', 'sgd_nesterov_momentum']:
12     print('Optimizing with {}'.format(update_rule))
13     model = FullyConnectedNet([100, 100, 100, 100, 100], weight_scale=5e-2)
14
15     solver = Solver(model, small_data,
16                     num_epochs=5, batch_size=100,
17                     update_rule=update_rule,
18                     optim_config={
19                         'learning_rate': 1e-2,
20                     },
21                     verbose=False)
22     solvers[update_rule] = solver
23     solver.train()
24     print
25
26 plt.subplot(3, 1, 1)
27 plt.title('Training loss')
28 plt.xlabel('Iteration')
29
30 plt.subplot(3, 1, 2)
31 plt.title('Training accuracy')
32 plt.xlabel('Epoch')
33
34 plt.subplot(3, 1, 3)
35 plt.title('Validation accuracy')
36 plt.xlabel('Epoch')
37
38 for update_rule, solver in solvers.items():
39     plt.subplot(3, 1, 1)
40     plt.plot(solver.loss_history, 'o', label=update_rule)
41
42     plt.subplot(3, 1, 2)
43     plt.plot(solver.train_acc_history, '-o', label=update_rule)
44
45     plt.subplot(3, 1, 3)
46     plt.plot(solver.val_acc_history, '-o', label=update_rule)
47
48 for i in [1, 2, 3]:
49     plt.subplot(3, 1, i)
50     plt.legend(loc='upper center', ncol=4)
51 plt.gcf().set_size_inches(15, 15)
52 plt.show()

```

Optimizing with sgd  
 Optimizing with sgd\_momentum  
 Optimizing with sgd\_nesterov\_momentum

/Users/hannah\_wang/anaconda3/lib/python3.6/site-packages/matplotlib/figure.py:98: MatplotlibDeprecationWarning:  
 Adding an axes using the same arguments as a previous axes currently reuses the earlier instance. In a future version,  
 a new instance will always be created and returned. Meanwhile, this warning can be suppressed, and the future behavior  
 ensured, by passing a unique label to each axes instance.

"Adding an axes using the same arguments as a previous axes "



## RMSProp

Now we go to techniques that adapt the gradient. Implement `rmsprop` in `nndl/optim.py`. Test your implementation by running the cell below.

```
In [9]: 1 from nndl.optim import rmsprop
2
3 N, D = 4, 5
4 w = np.linspace(-0.4, 0.6, num=N*D).reshape(N, D)
5 dw = np.linspace(-0.6, 0.4, num=N*D).reshape(N, D)
6 a = np.linspace(0.6, 0.9, num=N*D).reshape(N, D)
7
8 config = {'learning_rate': 1e-2, 'a': a}
9 next_w, _ = rmsprop(w, dw, config=config)
10
11 expected_next_w = np.asarray([
12     [-0.39223849, -0.34037513, -0.28849239, -0.23659121, -0.18467247],
13     [-0.132737, -0.08078555, -0.02881884, 0.02316247, 0.07515774],
14     [0.12716641, 0.17918792, 0.23122175, 0.28326742, 0.33532447],
15     [0.38739248, 0.43947102, 0.49155973, 0.54365823, 0.59576619]])
16 expected_cache = np.asarray([
17     [0.5976, 0.6126277, 0.6277108, 0.64284931, 0.65804321],
18     [0.67329252, 0.68859723, 0.70395734, 0.71937285, 0.73484377],
19     [0.75037008, 0.7659518, 0.78158892, 0.79728144, 0.81302936],
20     [0.82883269, 0.84469141, 0.86060554, 0.87657507, 0.8926 ]])
21
22 print('next_w error: {}'.format(rel_error(expected_next_w, next_w)))
23 print('cache error: {}'.format(rel_error(expected_cache, config['a'])))
```

```
next_w error: 9.524687511038133e-08
cache error: 2.6477955807156126e-09
```

## Adaptive moments

Now, implement adam in nndl/optim.py . Test your implementation by running the cell below.

```
In [10]: 1 # Test Adam implementation; you should see errors around 1e-7 or less
2 from nndl.optim import adam
3
4 N, D = 4, 5
5 w = np.linspace(-0.4, 0.6, num=N*D).reshape(N, D)
6 dw = np.linspace(-0.6, 0.4, num=N*D).reshape(N, D)
7 v = np.linspace(0.6, 0.9, num=N*D).reshape(N, D)
8 a = np.linspace(0.7, 0.5, num=N*D).reshape(N, D)
9
10 config = {'learning_rate': 1e-2, 'v': v, 'a': a, 't': 5}
11 next_w, _ = adam(w, dw, config=config)
12
13 expected_next_w = np.asarray([
14     [-0.40094747, -0.34836187, -0.29577703, -0.24319299, -0.19060977],
15     [-0.1380274, -0.08544591, -0.03286534, 0.01971428, 0.0722929],
16     [0.1248705, 0.17744702, 0.23002243, 0.28259667, 0.33516969],
17     [0.38774145, 0.44031188, 0.49288093, 0.54544852, 0.59801459]])
18 expected_a = np.asarray([
19     [0.69966, 0.68908382, 0.67851319, 0.66794809, 0.65738853],
20     [0.64683452, 0.63628604, 0.6257431, 0.61520571, 0.60467385],
21     [0.59414753, 0.58362676, 0.57311152, 0.56260183, 0.55209767],
22     [0.54159906, 0.53110598, 0.52061845, 0.51013645, 0.49966, ]])
23 expected_v = np.asarray([
24     [0.48, 0.49947368, 0.51894737, 0.53842105, 0.55789474],
25     [0.57736842, 0.59684211, 0.61631579, 0.63578947, 0.65526316],
26     [0.67473684, 0.69421053, 0.71368421, 0.73315789, 0.75263158],
27     [0.77210526, 0.79157895, 0.81105263, 0.83052632, 0.85, ]])
28
29 print('next_w error: {}'.format(rel_error(expected_next_w, next_w)))
30 print('a error: {}'.format(rel_error(expected_a, config['a'])))
31 print('v error: {}'.format(rel_error(expected_v, config['v'])))
```

next\_w error: 1.1395691798535431e-07

a error: 4.208314038113071e-09

v error: 4.214963193114416e-09

## Comparing SGD, SGD+NesterovMomentum, RMSProp, and Adam

The following code will compare optimization with SGD, Momentum, Nesterov Momentum, RMSProp and Adam. In our code, we find that RMSProp, Adam, and SGD + Nesterov Momentum achieve approximately the same training error after a few training epochs.


```

In [15]: 1 learning_rates = {'rmsprop': 2e-4, 'adam': 1e-3}
2
3 for update_rule in ['adam', 'rmsprop']:
4     print('Optimizing with {}'.format(update_rule))
5     model = FullyConnectedNet([100, 100, 100, 100, 100], weight_scale=5e-2)
6
7     solver = Solver(model, small_data,
8                     num_epochs=5, batch_size=100,
9                     update_rule=update_rule,
10                    optim_config={
11                        'learning_rate': learning_rates[update_rule]
12                    },
13                    verbose=False)
14     solvers[update_rule] = solver
15     solver.train()
16     print
17
18 plt.subplot(3, 1, 1)
19 plt.title('Training loss')
20 plt.xlabel('Iteration')
21
22 plt.subplot(3, 1, 2)
23 plt.title('Training accuracy')
24 plt.xlabel('Epoch')
25
26 plt.subplot(3, 1, 3)
27 plt.title('Validation accuracy')
28 plt.xlabel('Epoch')
29
30 for update_rule, solver in solvers.items():
31     plt.subplot(3, 1, 1)
32     plt.plot(solver.loss_history, 'o', label=update_rule)
33
34     plt.subplot(3, 1, 2)
35     plt.plot(solver.train_acc_history, '-o', label=update_rule)
36
37     plt.subplot(3, 1, 3)
38     plt.plot(solver.val_acc_history, '-o', label=update_rule)
39
40 for i in [1, 2, 3]:
41     plt.subplot(3, 1, i)
42     plt.legend(loc='upper center', ncol=4)
43 plt.gcf().set_size_inches(15, 15)
44 plt.show()

```

Optimizing with adam  
Optimizing with rmsprop

/Users/hannah\_wang/anaconda3/lib/python3.6/site-packages/matplotlib/figure.py:98: MatplotlibDeprecationWarning:  
Adding an axes using the same arguments as a previous axes currently reuses the earlier instance. In a future version,  
a new instance will always be created and returned. Meanwhile, this warning can be suppressed, and the future behavior  
ensured, by passing a unique label to each axes instance.  
"Adding an axes using the same arguments as a previous axes "



## Easier optimization

In the following cell, we'll train a 4 layer neural network having 500 units in each hidden layer with the different optimizers, and find that it is far easier to get up to 50+% performance on CIFAR-10. After we implement batchnorm and dropout, we'll ask you to get 60+% on CIFAR-10.



```
In [16]: 1 optimizer = 'adam'
2 best_model = None
3
4 layer_dims = [500, 500, 500]
5 weight_scale = 0.01
6 learning_rate = 1e-3
7 lr_decay = 0.9
8
9 model = FullyConnectedNet(layer_dims, weight_scale=weight_scale,
10                           use_batchnorm=True)
11
12 solver = Solver(model, data,
13                 num_epochs=10, batch_size=100,
14                 update_rule=optimizer,
15                 optim_config={
16                     'learning_rate': learning_rate,
17                 },
18                 lr_decay=lr_decay,
19                 verbose=True, print_every=50)
20 solver.train()
```

```
(Iteration 1 / 4900) loss: 2.317612
(Epoch 0 / 10) train acc: 0.214000; val_acc: 0.222000
(Iteration 51 / 4900) loss: 1.850729
(Iteration 101 / 4900) loss: 1.529407
(Iteration 151 / 4900) loss: 1.734173
(Iteration 201 / 4900) loss: 1.648241
(Iteration 251 / 4900) loss: 1.607426
(Iteration 301 / 4900) loss: 1.404215
(Iteration 351 / 4900) loss: 1.707623
(Iteration 401 / 4900) loss: 1.347259
(Iteration 451 / 4900) loss: 1.336099
(Epoch 1 / 10) train acc: 0.470000; val_acc: 0.456000
(Iteration 501 / 4900) loss: 1.372245
(Iteration 551 / 4900) loss: 1.223931
(Iteration 601 / 4900) loss: 1.269121
(Iteration 651 / 4900) loss: 1.391613
(Iteration 701 / 4900) loss: 1.319634
(Iteration 751 / 4900) loss: 1.299265
(Iteration 801 / 4900) loss: 1.345726
(Iteration 851 / 4900) loss: 1.238437
(Iteration 901 / 4900) loss: 1.301950
(Iteration 951 / 4900) loss: 1.184255
(Epoch 2 / 10) train acc: 0.553000; val_acc: 0.519000
(Iteration 1001 / 4900) loss: 1.447854
(Iteration 1051 / 4900) loss: 1.135678
(Iteration 1101 / 4900) loss: 1.067297
(Iteration 1151 / 4900) loss: 1.279670
(Iteration 1201 / 4900) loss: 1.213916
(Iteration 1251 / 4900) loss: 1.283601
(Iteration 1301 / 4900) loss: 1.303041
(Iteration 1351 / 4900) loss: 1.313844
(Iteration 1401 / 4900) loss: 1.038823
(Iteration 1451 / 4900) loss: 1.161419
(Epoch 3 / 10) train acc: 0.578000; val_acc: 0.524000
(Iteration 1501 / 4900) loss: 1.080736
(Iteration 1551 / 4900) loss: 1.033311
(Iteration 1601 / 4900) loss: 1.256431
(Iteration 1651 / 4900) loss: 1.165053
(Iteration 1701 / 4900) loss: 1.240109
(Iteration 1751 / 4900) loss: 0.956544
(Iteration 1801 / 4900) loss: 1.123703
(Iteration 1851 / 4900) loss: 1.019171
(Iteration 1901 / 4900) loss: 1.116145
(Iteration 1951 / 4900) loss: 1.156048
(Epoch 4 / 10) train acc: 0.591000; val_acc: 0.544000
(Iteration 2001 / 4900) loss: 1.324829
(Iteration 2051 / 4900) loss: 1.150866
(Iteration 2101 / 4900) loss: 0.935841
(Iteration 2151 / 4900) loss: 0.957888
(Iteration 2201 / 4900) loss: 0.936134
(Iteration 2251 / 4900) loss: 0.978844
(Iteration 2301 / 4900) loss: 0.941161
(Iteration 2351 / 4900) loss: 0.915056
(Iteration 2401 / 4900) loss: 1.104329
(Epoch 5 / 10) train acc: 0.672000; val_acc: 0.532000
(Iteration 2451 / 4900) loss: 0.710098
(Iteration 2501 / 4900) loss: 0.913061
(Iteration 2551 / 4900) loss: 1.135810
(Iteration 2601 / 4900) loss: 0.876673
(Iteration 2651 / 4900) loss: 0.989470
(Iteration 2701 / 4900) loss: 0.990899
(Iteration 2751 / 4900) loss: 0.909822
(Iteration 2801 / 4900) loss: 0.756945
(Iteration 2851 / 4900) loss: 0.910173
(Iteration 2901 / 4900) loss: 0.725418
(Epoch 6 / 10) train acc: 0.692000; val_acc: 0.535000
(Iteration 2951 / 4900) loss: 0.615575
```

```

(Iteration 3001 / 4900) loss: 0.878288
(Iteration 3051 / 4900) loss: 1.034637
(Iteration 3101 / 4900) loss: 0.660970
(Iteration 3151 / 4900) loss: 0.891185
(Iteration 3201 / 4900) loss: 0.765091
(Iteration 3251 / 4900) loss: 0.859637
(Iteration 3301 / 4900) loss: 0.739107
(Iteration 3351 / 4900) loss: 0.663448
(Iteration 3401 / 4900) loss: 0.911083
(Epoch 7 / 10) train acc: 0.711000; val_acc: 0.569000
(Iteration 3451 / 4900) loss: 0.755873
(Iteration 3501 / 4900) loss: 0.817034
(Iteration 3551 / 4900) loss: 0.873263
(Iteration 3601 / 4900) loss: 0.684989
(Iteration 3651 / 4900) loss: 0.784117
(Iteration 3701 / 4900) loss: 0.709356
(Iteration 3751 / 4900) loss: 0.724685
(Iteration 3801 / 4900) loss: 0.639091
(Iteration 3851 / 4900) loss: 0.738882
(Iteration 3901 / 4900) loss: 0.473726
(Epoch 8 / 10) train acc: 0.754000; val_acc: 0.554000
(Iteration 3951 / 4900) loss: 0.587603
(Iteration 4001 / 4900) loss: 0.856407
(Iteration 4051 / 4900) loss: 0.672353
(Iteration 4101 / 4900) loss: 0.781707
(Iteration 4151 / 4900) loss: 0.604926
(Iteration 4201 / 4900) loss: 0.596246
(Iteration 4251 / 4900) loss: 0.630911
(Iteration 4301 / 4900) loss: 0.534743
(Iteration 4351 / 4900) loss: 0.716217
(Iteration 4401 / 4900) loss: 0.587710
(Epoch 9 / 10) train acc: 0.802000; val_acc: 0.569000
(Iteration 4451 / 4900) loss: 0.556148
(Iteration 4501 / 4900) loss: 0.547765
(Iteration 4551 / 4900) loss: 0.704546
(Iteration 4601 / 4900) loss: 0.596978
(Iteration 4651 / 4900) loss: 0.642664
(Iteration 4701 / 4900) loss: 0.584378
(Iteration 4751 / 4900) loss: 0.542869
(Iteration 4801 / 4900) loss: 0.612463
(Iteration 4851 / 4900) loss: 0.613176
(Epoch 10 / 10) train acc: 0.818000; val_acc: 0.557000

```

```

In [17]: 1 y_test_pred = np.argmax(model.loss(data['X_test']), axis=1)
          2 y_val_pred = np.argmax(model.loss(data['X_val']), axis=1)
          3 print('Validation set accuracy: {}'.format(np.mean(y_val_pred == data['y_val'])))
          4 print('Test set accuracy: {}'.format(np.mean(y_test_pred == data['y_test'])))

```

```

Validation set accuracy: 0.578
Test set accuracy: 0.569

```

```

In [ ]: 1

```

# Batch Normalization

In this notebook, you will implement the batch normalization layers of a neural network to increase its performance. If you have any confusion, please review the details of batch normalization from the lecture notes.

CS231n has built a solid API for building these modular frameworks and training them, and we will use their very well implemented framework as opposed to "reinventing the wheel." This includes using their Solver, various utility functions, and their layer structure. This also includes `nndl.fc_net`, `nndl.layers`, and `nndl.layer_utils`. As in prior assignments, we thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class ([cs231n.stanford.edu](http://cs231n.stanford.edu)).

```
In [1]: 1  ## Import and setups
2
3  import time
4  import numpy as np
5  import matplotlib.pyplot as plt
6  from nndl.fc_net import *
7  from nndl.layers import *
8  from cs231n.data_utils import get_CIFAR10_data
9  from cs231n.gradient_check import eval_numerical_gradient, eval_numerical_gradient_array
10 from cs231n.solver import Solver
11
12 %matplotlib inline
13 plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
14 plt.rcParams['image.interpolation'] = 'nearest'
15 plt.rcParams['image.cmap'] = 'gray'
16
17 # for auto-reloading external modules
18 # see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
19 %load_ext autoreload
20 %autoreload 2
21
22 def rel_error(x, y):
23     """ returns relative error """
24     return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

```
In [2]: 1  # Load the (preprocessed) CIFAR10 data.
2
3  data = get_CIFAR10_data()
4  for k in data.keys():
5      print('{}: {}'.format(k, data[k].shape))
```

```
X_train: (49000, 3, 32, 32)
y_train: (49000,)
X_val: (1000, 3, 32, 32)
y_val: (1000,)
X_test: (1000, 3, 32, 32)
y_test: (1000,)
```

## Batchnorm forward pass

Implement the training time batchnorm forward pass, `batchnorm_forward`, in `nndl/layers.py`. After that, test your implementation by running the following cell.

```
In [4]: 1 # Check the training-time forward pass by checking means and variances
2 # of features both before and after batch normalization
3
4 # Simulate the forward pass for a two-layer network
5 N, D1, D2, D3 = 200, 50, 60, 3
6 X = np.random.randn(N, D1)
7 W1 = np.random.randn(D1, D2)
8 W2 = np.random.randn(D2, D3)
9 a = np.maximum(0, X.dot(W1)).dot(W2)
10
11 print('Before batch normalization:')
12 print(' means: ', a.mean(axis=0))
13 print(' stds: ', a.std(axis=0))
14
15 # Means should be close to zero and stds close to one
16 print('After batch normalization (gamma=1, beta=0)')
17 a_norm, _ = batchnorm_forward(a, np.ones(D3), np.zeros(D3), {'mode': 'train'})
18 print(' mean: ', a_norm.mean(axis=0))
19 print(' std: ', a_norm.std(axis=0))
20
21 # Now means should be close to beta and stds close to gamma
22 gamma = np.asarray([1.0, 2.0, 3.0])
23 beta = np.asarray([11.0, 12.0, 13.0])
24 a_norm, _ = batchnorm_forward(a, gamma, beta, {'mode': 'train'})
25 print('After batch normalization (nontrivial gamma, beta)')
26 print(' means: ', a_norm.mean(axis=0))
27 print(' stds: ', a_norm.std(axis=0))
```

```
Before batch normalization:
 means: [-9.27589484 21.75456393 17.54660778]
 stds: [29.33259377 40.84649232 29.16922165]
After batch normalization (gamma=1, beta=0)
 mean: [-5.10702591e-17 -5.41927614e-17 -1.74860126e-16]
 std: [0.99999999 1. 0.99999999]
After batch normalization (nontrivial gamma, beta)
 means: [11. 12. 13.]
 stds: [0.99999999 1.99999999 2.99999998]
```

Implement the testing time batchnorm forward pass, `batchnorm_forward`, in `nndl/layers.py`. After that, test your implementation by running the following cell.

```
In [6]: 1 # Check the test-time forward pass by running the training-time
2 # forward pass many times to warm up the running averages, and then
3 # checking the means and variances of activations after a test-time
4 # forward pass.
5
6 N, D1, D2, D3 = 200, 50, 60, 3
7 W1 = np.random.randn(D1, D2)
8 W2 = np.random.randn(D2, D3)
9
10 bn_param = {'mode': 'train'}
11 gamma = np.ones(D3)
12 beta = np.zeros(D3)
13 for t in np.arange(50):
14     X = np.random.randn(N, D1)
15     a = np.maximum(0, X.dot(W1)).dot(W2)
16     batchnorm_forward(a, gamma, beta, bn_param)
17 bn_param['mode'] = 'test'
18 X = np.random.randn(N, D1)
19 a = np.maximum(0, X.dot(W1)).dot(W2)
20 a_norm, _ = batchnorm_forward(a, gamma, beta, bn_param)
21
22 # Means should be close to zero and stds close to one, but will be
23 # noisier than training-time forward passes.
24 print('After batch normalization (test-time):')
25 print(' means: ', a_norm.mean(axis=0))
26 print(' stds: ', a_norm.std(axis=0))
```

```
After batch normalization (test-time):
 means: [-0.04310833 0.03218559 -0.01450178]
 stds: [0.98114234 1.08661348 1.01088546]
```

## Batchnorm backward pass

Implement the backward pass for the batchnorm layer, `batchnorm_backward` in `nndl/layers.py`. Check your implementation by running the following cell.

```
In [7]: 1 # Gradient check batchnorm backward pass
2
3 N, D = 4, 5
4 x = 5 * np.random.randn(N, D) + 12
5 gamma = np.random.randn(D)
6 beta = np.random.randn(D)
7 dout = np.random.randn(N, D)
8
9 bn_param = {'mode': 'train'}
10 fx = lambda x: batchnorm_forward(x, gamma, beta, bn_param)[0]
11 fg = lambda a: batchnorm_forward(x, gamma, beta, bn_param)[0]
12 fb = lambda b: batchnorm_forward(x, gamma, beta, bn_param)[0]
13
14 dx_num = eval_numerical_gradient_array(fx, x, dout)
15 da_num = eval_numerical_gradient_array(fg, gamma, dout)
16 db_num = eval_numerical_gradient_array(fb, beta, dout)
17
18 _, cache = batchnorm_forward(x, gamma, beta, bn_param)
19 dx, dgamma, dbeta = batchnorm_backward(dout, cache)
20 print('dx error: ', rel_error(dx_num, dx))
21 print('dgamma error: ', rel_error(da_num, dgamma))
22 print('dbeta error: ', rel_error(db_num, dbeta))
```

```
dx error: 1.3407317993977258e-09
dgamma error: 3.2369133023643904e-12
dbeta error: 3.275420131545755e-12
```

## Implement a fully connected neural network with batchnorm layers

Modify the `FullyConnectedNet()` class in `nndl/fc_net.py` to incorporate batchnorm layers. You will need to modify the class in the following areas:

- (1) The gammas and betas need to be initialized to 1's and 0's respectively in `__init__`.
- (2) The `batchnorm_forward` layer needs to be inserted between each affine and relu layer (except in the output layer) in a forward pass computation in `loss`. You may find it helpful to write an `affine_batchnorm_relu()` layer in `nndl/layer_utils.py` although this is not necessary.
- (3) The `batchnorm_backward` layer has to be appropriately inserted when calculating gradients.

After you have done the appropriate modifications, check your implementation by running the following cell.

Note, while the relative error for W3 should be small, as we backprop gradients more, you may find the relative error increases. Our relative error for W1 is on the order of  $1e-4$ .

```
In [8]: 1 N, D, H1, H2, C = 2, 15, 20, 30, 10
2 X = np.random.randn(N, D)
3 y = np.random.randint(C, size=(N,))
4
5 for reg in [0, 3.14]:
6     print('Running check with reg = ', reg)
7     model = FullyConnectedNet([H1, H2], input_dim=D, num_classes=C,
8                               reg=reg, weight_scale=5e-2, dtype=np.float64,
9                               use_batchnorm=True)
10
11     loss, grads = model.loss(X, y)
12     print('Initial loss: ', loss)
13
14     for name in sorted(grads):
15         f = lambda _: model.loss(X, y)[0]
16         grad_num = eval_numerical_gradient(f, model.params[name], verbose=False, h=1e-5)
17         print('{} relative error: {}'.format(name, rel_error(grad_num, grads[name])))
18     if reg == 0: print('\n')
```

```
Running check with reg = 0
Initial loss: 2.116495647317618
W1 relative error: 0.003059418529722701
W2 relative error: 3.848515927934831e-06
W3 relative error: 4.093183258233747e-10
b1 relative error: 4.336808689942018e-11
b2 relative error: 1.1102230246251565e-08
b3 relative error: 1.237819602511532e-10
beta1 relative error: 3.265429352574524e-07
beta2 relative error: 7.424791031728775e-09
gamma1 relative error: 3.984791350136109e-07
gamma2 relative error: 3.9423807659745554e-09
```

```
Running check with reg = 3.14
Initial loss: 6.945176771933103
W1 relative error: 1.602456195923461e-05
W2 relative error: 2.03407132663365e-06
W3 relative error: 1.4177185950422472e-08
b1 relative error: 1.1102230246251565e-08
b2 relative error: 2.7755575615628914e-09
b3 relative error: 1.47406506027144e-10
beta1 relative error: 4.766663710604228e-09
beta2 relative error: 2.5601940267967956e-08
gamma1 relative error: 4.697072145384006e-09
gamma2 relative error: 5.502280624431884e-09
```

## Training a deep fully connected network with batch normalization.

To see if batchnorm helps, let's train a deep neural network with and without batch normalization.

```

In [9]: 1 # Try training a very deep net with batchnorm
2 hidden_dims = [100, 100, 100, 100, 100]
3
4 num_train = 1000
5 small_data = {
6     'X_train': data['X_train'][:num_train],
7     'y_train': data['y_train'][:num_train],
8     'X_val': data['X_val'],
9     'y_val': data['y_val'],
10 }
11
12 weight_scale = 2e-2
13 bn_model = FullyConnectedNet(hidden_dims, weight_scale=weight_scale, use_batchnorm=True)
14 model = FullyConnectedNet(hidden_dims, weight_scale=weight_scale, use_batchnorm=False)
15
16 bn_solver = Solver(bn_model, small_data,
17                     num_epochs=10, batch_size=50,
18                     update_rule='adam',
19                     optim_config={
20                         'learning_rate': 1e-3,
21                     },
22                     verbose=True, print_every=200)
23 bn_solver.train()
24
25 solver = Solver(model, small_data,
26                 num_epochs=10, batch_size=50,
27                 update_rule='adam',
28                 optim_config={
29                     'learning_rate': 1e-3,
30                 },
31                 verbose=True, print_every=200)
32 solver.train()

```

```

(Iteration 1 / 200) loss: 2.306698
(Epoch 0 / 10) train acc: 0.139000; val_acc: 0.118000
(Epoch 1 / 10) train acc: 0.340000; val_acc: 0.279000
(Epoch 2 / 10) train acc: 0.424000; val_acc: 0.303000
(Epoch 3 / 10) train acc: 0.509000; val_acc: 0.311000
(Epoch 4 / 10) train acc: 0.572000; val_acc: 0.323000
(Epoch 5 / 10) train acc: 0.619000; val_acc: 0.318000
(Epoch 6 / 10) train acc: 0.653000; val_acc: 0.333000
(Epoch 7 / 10) train acc: 0.732000; val_acc: 0.345000
(Epoch 8 / 10) train acc: 0.754000; val_acc: 0.329000
(Epoch 9 / 10) train acc: 0.800000; val_acc: 0.329000
(Epoch 10 / 10) train acc: 0.807000; val_acc: 0.313000
(Iteration 1 / 200) loss: 2.302677
(Epoch 0 / 10) train acc: 0.165000; val_acc: 0.155000
(Epoch 1 / 10) train acc: 0.241000; val_acc: 0.229000
(Epoch 2 / 10) train acc: 0.290000; val_acc: 0.255000
(Epoch 3 / 10) train acc: 0.343000; val_acc: 0.303000
(Epoch 4 / 10) train acc: 0.361000; val_acc: 0.275000
(Epoch 5 / 10) train acc: 0.415000; val_acc: 0.280000
(Epoch 6 / 10) train acc: 0.488000; val_acc: 0.317000
(Epoch 7 / 10) train acc: 0.447000; val_acc: 0.286000
(Epoch 8 / 10) train acc: 0.524000; val_acc: 0.298000
(Epoch 9 / 10) train acc: 0.550000; val_acc: 0.334000
(Epoch 10 / 10) train acc: 0.658000; val_acc: 0.338000

```


```

In [10]: 1 plt.subplot(3, 1, 1)
2 plt.title('Training loss')
3 plt.xlabel('Iteration')
4
5 plt.subplot(3, 1, 2)
6 plt.title('Training accuracy')
7 plt.xlabel('Epoch')
8
9 plt.subplot(3, 1, 3)
10 plt.title('Validation accuracy')
11 plt.xlabel('Epoch')
12
13 plt.subplot(3, 1, 1)
14 plt.plot(solver.loss_history, 'o', label='baseline')
15 plt.plot(bn_solver.loss_history, 'o', label='batchnorm')
16
17 plt.subplot(3, 1, 2)
18 plt.plot(solver.train_acc_history, '-o', label='baseline')
19 plt.plot(bn_solver.train_acc_history, '-o', label='batchnorm')
20
21 plt.subplot(3, 1, 3)
22 plt.plot(solver.val_acc_history, '-o', label='baseline')
23 plt.plot(bn_solver.val_acc_history, '-o', label='batchnorm')
24
25 for i in [1, 2, 3]:
26     plt.subplot(3, 1, i)
27     plt.legend(loc='upper center', ncol=4)
28 plt.gcf().set_size_inches(15, 15)
29 plt.show()

```

/Users/hannah\_wang/anaconda3/lib/python3.6/site-packages/matplotlib/figure.py:98: MatplotlibDeprecationWarning: Adding an axes using the same arguments as a previous axes currently reuses the earlier instance. In a future version, a new instance will always be created and returned. Meanwhile, this warning can be suppressed, and the future behavior ensured, by passing a unique label to each axes instance.

"Adding an axes using the same arguments as a previous axes "





## Batchnorm and initialization

The following cells run an experiment where for a deep network, the initialization is varied. We do training for when batchnorm layers are and are not included.

```
In [11]: 1 # Try training a very deep net with batchnorm
2 hidden_dims = [50, 50, 50, 50, 50, 50, 50]
3
4 num_train = 1000
5 small_data = {
6     'X_train': data['X_train'][:num_train],
7     'y_train': data['y_train'][:num_train],
8     'X_val': data['X_val'],
9     'y_val': data['y_val'],
10 }
11
12 bn_solvers = {}
13 solvers = {}
14 weight_scales = np.logspace(-4, 0, num=20)
15 for i, weight_scale in enumerate(weight_scales):
16     print('Running weight scale {} / {}'.format(i + 1, len(weight_scales)))
17     bn_model = FullyConnectedNet(hidden_dims, weight_scale=weight_scale, use_batchnorm=True)
18     model = FullyConnectedNet(hidden_dims, weight_scale=weight_scale, use_batchnorm=False)
19
20     bn_solver = Solver(bn_model, small_data,
21                       num_epochs=10, batch_size=50,
22                       update_rule='adam',
23                       optim_config={
24                           'learning_rate': 1e-3,
25                       },
26                       verbose=False, print_every=200)
27     bn_solver.train()
28     bn_solvers[weight_scale] = bn_solver
29
30     solver = Solver(model, small_data,
31                    num_epochs=10, batch_size=50,
32                    update_rule='adam',
33                    optim_config={
34                        'learning_rate': 1e-3,
35                    },
36                    verbose=False, print_every=200)
37     solver.train()
38     solvers[weight_scale] = solver
```

```
Running weight scale 1 / 20
Running weight scale 2 / 20
Running weight scale 3 / 20
Running weight scale 4 / 20
Running weight scale 5 / 20
Running weight scale 6 / 20
Running weight scale 7 / 20
Running weight scale 8 / 20
Running weight scale 9 / 20
Running weight scale 10 / 20
Running weight scale 11 / 20
Running weight scale 12 / 20
Running weight scale 13 / 20
Running weight scale 14 / 20
Running weight scale 15 / 20
Running weight scale 16 / 20
```


```
/Users/hannah_wang/Desktop/hw4/code/nndl/layers.py:438: RuntimeWarning: divide by zero encountered in log
  loss = -np.sum(np.log(probs[np.arange(N), y])) / N
```

```
Running weight scale 17 / 20
Running weight scale 18 / 20
Running weight scale 19 / 20
Running weight scale 20 / 20
```

```

In [12]: 1 # Plot results of weight scale experiment
2 best_train_accs, bn_best_train_accs = [], []
3 best_val_accs, bn_best_val_accs = [], []
4 final_train_loss, bn_final_train_loss = [], []
5
6 for ws in weight_scales:
7     best_train_accs.append(max(solvers[ws].train_acc_history))
8     bn_best_train_accs.append(max(bn_solvers[ws].train_acc_history))
9
10    best_val_accs.append(max(solvers[ws].val_acc_history))
11    bn_best_val_accs.append(max(bn_solvers[ws].val_acc_history))
12
13    final_train_loss.append(np.mean(solvers[ws].loss_history[-100:]))
14    bn_final_train_loss.append(np.mean(bn_solvers[ws].loss_history[-100:]))
15
16    plt.subplot(3, 1, 1)
17    plt.title('Best val accuracy vs weight initialization scale')
18    plt.xlabel('Weight initialization scale')
19    plt.ylabel('Best val accuracy')
20    plt.semilogx(weight_scales, best_val_accs, '-o', label='baseline')
21    plt.semilogx(weight_scales, bn_best_val_accs, '-o', label='batchnorm')
22    plt.legend(ncol=2, loc='lower right')
23
24    plt.subplot(3, 1, 2)
25    plt.title('Best train accuracy vs weight initialization scale')
26    plt.xlabel('Weight initialization scale')
27    plt.ylabel('Best training accuracy')
28    plt.semilogx(weight_scales, best_train_accs, '-o', label='baseline')
29    plt.semilogx(weight_scales, bn_best_train_accs, '-o', label='batchnorm')
30    plt.legend()
31
32    plt.subplot(3, 1, 3)
33    plt.title('Final training loss vs weight initialization scale')
34    plt.xlabel('Weight initialization scale')
35    plt.ylabel('Final training loss')
36    plt.semilogx(weight_scales, final_train_loss, '-o', label='baseline')
37    plt.semilogx(weight_scales, bn_final_train_loss, '-o', label='batchnorm')
38    plt.legend()
39
40    plt.gcf().set_size_inches(10, 15)
41    plt.show()

```



## Question:

In the cell below, summarize the findings of this experiment, and WHY these results make sense.

## Answer:

From the figures above, on the one hand, batchnorm can achieve more accurate results. We can find that when weight initialization scale is smaller than  $10^{-1}$ , the accuracies of model with batchnorm are almost larger than those of model without batchnorm. And the loss of model with batchnorm are almost smaller than those of model without batchnorm. While when weight initialization scale is larger than  $10^{-1}$ , the model with batchnorm has poorer performance than model without batchnorm. On the other hand, we find that the model with batchnorm is much less sensitive to the weight initialization scale. The orange curves have wider stable ranges than blue curves.

As we know, batchnorm can be used as a regularizer to make results more stable. Thus, the model with batchnorm should be less affected by weight initialization scale. And the results we observed meet this theory and make sense.

In [ ]:

1

# Dropout

In this notebook, you will implement dropout. Then we will ask you to train a network with batchnorm and dropout, and achieve over 60% accuracy on CIFAR-10.

CS231n has built a solid API for building these modular frameworks and training them, and we will use their very well implemented framework as opposed to "reinventing the wheel." This includes using their Solver, various utility functions, and their layer structure. This also includes `nndl.fc_net`, `nndl.layers`, and `nndl.layer_utils`. As in prior assignments, we thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class ([cs231n.stanford.edu](http://cs231n.stanford.edu)).

```
In [2]: 1  ## Import and setups
2
3  import time
4  import numpy as np
5  import matplotlib.pyplot as plt
6  from nndl.fc_net import *
7  from nndl.layers import *
8  from cs231n.data_utils import get_CIFAR10_data
9  from cs231n.gradient_check import eval_numerical_gradient, eval_numerical_gradient_array
10 from cs231n.solver import Solver
11
12 %matplotlib inline
13 plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
14 plt.rcParams['image.interpolation'] = 'nearest'
15 plt.rcParams['image.cmap'] = 'gray'
16
17 # for auto-reloading external modules
18 # see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
19 %load_ext autoreload
20 %autoreload 2
21
22 def rel_error(x, y):
23     """ returns relative error """
24     return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

```
In [3]: 1  # Load the (preprocessed) CIFAR10 data.
2
3  data = get_CIFAR10_data()
4  for k in data.keys():
5      print('{}: {}'.format(k, data[k].shape))
```

```
X_train: (49000, 3, 32, 32)
y_train: (49000,)
X_val: (1000, 3, 32, 32)
y_val: (1000,)
X_test: (1000, 3, 32, 32)
y_test: (1000,)
```

## Dropout forward pass

Implement the training and test time dropout forward pass, `dropout_forward`, in `nndl/layers.py`. After that, test your implementation by running the following cell.

```
In [4]: 1 x = np.random.randn(500, 500) + 10
2
3 for p in [0.3, 0.6, 0.75]:
4     out, _ = dropout_forward(x, {'mode': 'train', 'p': p})
5     out_test, _ = dropout_forward(x, {'mode': 'test', 'p': p})
6
7     print('Running tests with p = ', p)
8     print('Mean of input: ', x.mean())
9     print('Mean of train-time output: ', out.mean())
10    print('Mean of test-time output: ', out_test.mean())
11    print('Fraction of train-time output set to zero: ', (out == 0).mean())
12    print('Fraction of test-time output set to zero: ', (out_test == 0).mean())
```

```
Running tests with p = 0.3
Mean of input: 9.997675409022476
Mean of train-time output: 9.982471546063854
Mean of test-time output: 9.997675409022476
Fraction of train-time output set to zero: 0.30116
Fraction of test-time output set to zero: 0.0
Running tests with p = 0.6
Mean of input: 9.997675409022476
Mean of train-time output: 9.953386552374035
Mean of test-time output: 9.997675409022476
Fraction of train-time output set to zero: 0.601796
Fraction of test-time output set to zero: 0.0
Running tests with p = 0.75
Mean of input: 9.997675409022476
Mean of train-time output: 9.96677999564878
Mean of test-time output: 9.997675409022476
Fraction of train-time output set to zero: 0.750764
Fraction of test-time output set to zero: 0.0
```

## Dropout backward pass

Implement the backward pass, `dropout_backward`, in `nndl/layers.py`. After that, test your gradients by running the following cell:

```
In [5]: 1 x = np.random.randn(10, 10) + 10
2 dout = np.random.randn(*x.shape)
3
4 dropout_param = {'mode': 'train', 'p': 0.8, 'seed': 123}
5 out, cache = dropout_forward(x, dropout_param)
6 dx = dropout_backward(dout, cache)
7 dx_num = eval_numerical_gradient_array(lambda xx: dropout_forward(xx, dropout_param)[0], x, dout)
8
9 print('dx relative error: ', rel_error(dx, dx_num))
```

```
dx relative error: 1.892905802538152e-11
```

## Implement a fully connected neural network with dropout layers

Modify the `FullyConnectedNet()` class in `nndl/fc_net.py` to incorporate dropout. A dropout layer should be incorporated after every ReLU layer. Concretely, there shouldn't be a dropout at the output layer since there is no ReLU at the output layer. You will need to modify the class in the following areas:

- (1) In the forward pass, you will need to incorporate a dropout layer after every relu layer.
- (2) In the backward pass, you will need to incorporate a dropout backward pass layer.

Check your implementation by running the following code. Our W1 gradient relative error is on the order of  $1e-6$  (the largest of all the relative errors).

```
In [6]: 1 N, D, H1, H2, C = 2, 15, 20, 30, 10
2 X = np.random.randn(N, D)
3 y = np.random.randint(C, size=(N,))
4
5 for dropout in [0, 0.25, 0.5]:
6     print('Running check with dropout = ', dropout)
7     model = FullyConnectedNet([H1, H2], input_dim=D, num_classes=C,
8                               weight_scale=5e-2, dtype=np.float64,
9                               dropout=dropout, seed=123)
10
11     loss, grads = model.loss(X, y)
12     print('Initial loss: ', loss)
13
14     for name in sorted(grads):
15         f = lambda _: model.loss(X, y)[0]
16         grad_num = eval_numerical_gradient(f, model.params[name], verbose=False, h=1e-5)
17         print('{} relative error: {}'.format(name, rel_error(grad_num, grads[name])))
18     print('\n')
```

```
Running check with dropout = 0
Initial loss: 2.3051948273987857
W1 relative error: 2.5272575344376073e-07
W2 relative error: 1.5034484929313676e-05
W3 relative error: 2.753446833630168e-07
b1 relative error: 2.936957476400148e-06
b2 relative error: 5.051339805546953e-08
b3 relative error: 1.1740467838205477e-10
```

```
Running check with dropout = 0.25
Initial loss: 2.3052077546540826
W1 relative error: 2.613846944812385e-07
W2 relative error: 5.022056536108928e-07
W3 relative error: 4.456316077044505e-08
b1 relative error: 7.39711723790801e-08
b2 relative error: 7.151678402730031e-10
b3 relative error: 1.003974732116764e-10
```

```
Running check with dropout = 0.5
Initial loss: 2.3035667586595423
W1 relative error: 1.1401257458777745e-06
W2 relative error: 1.847669681023635e-07
W3 relative error: 6.5966195253431734e-09
b1 relative error: 7.71639621892128e-08
b2 relative error: 1.1975910493629166e-09
b3 relative error: 1.4558471033827801e-10
```

## Dropout as a regularizer

In class, we claimed that dropout acts as a regularizer by effectively bagging. To check this, we will train two small networks, one with dropout and one without dropout.

```

In [7]: 1 # Train two identical nets, one with dropout and one without
2
3 num_train = 500
4 small_data = {
5     'X_train': data['X_train'][:num_train],
6     'y_train': data['y_train'][:num_train],
7     'X_val': data['X_val'],
8     'y_val': data['y_val'],
9 }
10
11 solvers = {}
12 dropout_choices = [0, 0.6]
13 for dropout in dropout_choices:
14     model = FullyConnectedNet([100, 100, 100], dropout=dropout)
15
16     solver = Solver(model, small_data,
17                     num_epochs=25, batch_size=100,
18                     update_rule='adam',
19                     optim_config={
20                         'learning_rate': 5e-4,
21                     },
22                     verbose=True, print_every=100)
23     solvers[dropout] = solver
24


```

```

(Iteration 1 / 125) loss: 2.300804
(Epoch 0 / 25) train acc: 0.220000; val_acc: 0.168000
(Epoch 1 / 25) train acc: 0.188000; val_acc: 0.147000
(Epoch 2 / 25) train acc: 0.266000; val_acc: 0.200000
(Epoch 3 / 25) train acc: 0.338000; val_acc: 0.262000
(Epoch 4 / 25) train acc: 0.378000; val_acc: 0.278000
(Epoch 5 / 25) train acc: 0.428000; val_acc: 0.297000
(Epoch 6 / 25) train acc: 0.468000; val_acc: 0.323000
(Epoch 7 / 25) train acc: 0.494000; val_acc: 0.287000
(Epoch 8 / 25) train acc: 0.566000; val_acc: 0.328000
(Epoch 9 / 25) train acc: 0.572000; val_acc: 0.322000
(Epoch 10 / 25) train acc: 0.622000; val_acc: 0.324000
(Epoch 11 / 25) train acc: 0.670000; val_acc: 0.279000
(Epoch 12 / 25) train acc: 0.710000; val_acc: 0.338000
(Epoch 13 / 25) train acc: 0.746000; val_acc: 0.319000
(Epoch 14 / 25) train acc: 0.792000; val_acc: 0.307000
(Epoch 15 / 25) train acc: 0.834000; val_acc: 0.297000
(Epoch 16 / 25) train acc: 0.876000; val_acc: 0.327000
(Epoch 17 / 25) train acc: 0.886000; val_acc: 0.320000
(Epoch 18 / 25) train acc: 0.918000; val_acc: 0.314000
(Epoch 19 / 25) train acc: 0.922000; val_acc: 0.290000
(Epoch 20 / 25) train acc: 0.944000; val_acc: 0.306000
(Iteration 101 / 125) loss: 0.156105
(Epoch 21 / 25) train acc: 0.968000; val_acc: 0.302000
(Epoch 22 / 25) train acc: 0.978000; val_acc: 0.302000
(Epoch 23 / 25) train acc: 0.976000; val_acc: 0.289000
(Epoch 24 / 25) train acc: 0.986000; val_acc: 0.285000
(Epoch 25 / 25) train acc: 0.978000; val_acc: 0.311000
(Iteration 1 / 125) loss: 2.298716
(Epoch 0 / 25) train acc: 0.132000; val_acc: 0.146000
(Epoch 1 / 25) train acc: 0.118000; val_acc: 0.131000
(Epoch 2 / 25) train acc: 0.220000; val_acc: 0.214000
(Epoch 3 / 25) train acc: 0.206000; val_acc: 0.180000
(Epoch 4 / 25) train acc: 0.220000; val_acc: 0.193000
(Epoch 5 / 25) train acc: 0.264000; val_acc: 0.229000
(Epoch 6 / 25) train acc: 0.268000; val_acc: 0.203000
(Epoch 7 / 25) train acc: 0.266000; val_acc: 0.212000
(Epoch 8 / 25) train acc: 0.282000; val_acc: 0.236000
(Epoch 9 / 25) train acc: 0.310000; val_acc: 0.255000
(Epoch 10 / 25) train acc: 0.320000; val_acc: 0.267000
(Epoch 11 / 25) train acc: 0.338000; val_acc: 0.273000
(Epoch 12 / 25) train acc: 0.346000; val_acc: 0.278000
(Epoch 13 / 25) train acc: 0.332000; val_acc: 0.279000
(Epoch 14 / 25) train acc: 0.328000; val_acc: 0.284000
(Epoch 15 / 25) train acc: 0.354000; val_acc: 0.271000
(Epoch 16 / 25) train acc: 0.386000; val_acc: 0.277000
(Epoch 17 / 25) train acc: 0.388000; val_acc: 0.297000
(Epoch 18 / 25) train acc: 0.402000; val_acc: 0.280000
(Epoch 19 / 25) train acc: 0.388000; val_acc: 0.274000
(Epoch 20 / 25) train acc: 0.386000; val_acc: 0.274000
(Iteration 101 / 125) loss: 1.919649
(Epoch 21 / 25) train acc: 0.402000; val_acc: 0.272000
(Epoch 22 / 25) train acc: 0.440000; val_acc: 0.286000
(Epoch 23 / 25) train acc: 0.458000; val_acc: 0.295000
(Epoch 24 / 25) train acc: 0.462000; val_acc: 0.311000
(Epoch 25 / 25) train acc: 0.466000; val_acc: 0.297000

```

```
In [8]: 1 # Plot train and validation accuracies of the two models
2
3 train_accs = []
4 val_accs = []
5 for dropout in dropout_choices:
6     solver = solvers[dropout]
7     train_accs.append(solver.train_acc_history[-1])
8     val_accs.append(solver.val_acc_history[-1])
9
10 plt.subplot(3, 1, 1)
11 for dropout in dropout_choices:
12     plt.plot(solvers[dropout].train_acc_history, 'o', label='%.2f dropout' % dropout)
13 plt.title('Train accuracy')
14 plt.xlabel('Epoch')
15 plt.ylabel('Accuracy')
16 plt.legend(ncol=2, loc='lower right')
17
18 plt.subplot(3, 1, 2)
19 for dropout in dropout_choices:
20     plt.plot(solvers[dropout].val_acc_history, 'o', label='%.2f dropout' % dropout)
21 plt.title('Val accuracy')
22 plt.xlabel('Epoch')
23 plt.ylabel('Accuracy')
24 plt.legend(ncol=2, loc='lower right')
25
26 plt.gcf().set_size_inches(15, 15)
27 plt.show()
```



## Question

Based off the results of this experiment, is dropout performing regularization? Explain your answer.

## Answer:

Yes, dropout is performing regularization. From two figures above, we can see that in the second figure, the validation accuracies of model with and without dropout are similar. While in the first figure, the training accuracy of model without dropout is apparently larger than the model with dropout. It indicates that the dropout regularized the overfitting in training data.

## Final part of the assignment

Get over 60% validation accuracy on CIFAR-10 by using the layers you have implemented. You will be graded according to the following equation:

$\min(\text{floor}((X - 32\%)) / 28\%, 1)$  where if you get 60% or higher validation accuracy, you get full points.



In [19]:

```
1 # ===== #
2 # YOUR CODE HERE:
3 # Implement a FC-net that achieves at least 60% validation accuracy
4 # on CIFAR-10.
5 # ===== #
6
7 # set parameters
8 hidden_dims = [600, 600, 600, 600]
9 learning_rate = 2e-3
10 weight_scale = 0.01
11 lr_decay = 0.95
12 dropout = 0.5
13 update_rule = 'adam'
14
15 # create FullyConnectedNet
16 model = FullyConnectedNet(hidden_dims=hidden_dims, weight_scale=weight_scale,
17                             dropout=dropout, use_batchnorm=True, reg=0.0)
18
19 # solve
20 solver = Solver(model, data,
21                  num_epochs=80, batch_size=100,
22                  update_rule=update_rule,
23                  optim_config={
24                      'learning_rate': learning_rate,
25                  },
26                  lr_decay=lr_decay,
27                  verbose=True, print_every=100)
28 solver.train()
29
30 # print out the validation accuracy
31 y_test_pred = np.argmax(model.loss(data['X_test']), axis=1)
32 y_val_pred = np.argmax(model.loss(data['X_val']), axis=1)
33 print('Validation set accuracy: {}'.format(np.mean(y_val_pred == data['y_val'])))
34 print('Test set accuracy: {}'.format(np.mean(y_test_pred == data['y_test'])))
35
36 # ===== #
37 # END YOUR CODE HERE
38 # ===== #
39
```

```
(Iteration 1 / 39200) loss: 2.361412
(Epoch 0 / 80) train acc: 0.185000; val_acc: 0.197000
(Iteration 101 / 39200) loss: 1.853467
(Iteration 201 / 39200) loss: 1.860087
(Iteration 301 / 39200) loss: 1.786663
(Iteration 401 / 39200) loss: 1.698297
(Epoch 1 / 80) train acc: 0.442000; val_acc: 0.442000
(Iteration 501 / 39200) loss: 1.532157
(Iteration 601 / 39200) loss: 1.779091
(Iteration 701 / 39200) loss: 1.590496
(Iteration 801 / 39200) loss: 1.645538
(Iteration 901 / 39200) loss: 1.716053
(Epoch 2 / 80) train acc: 0.457000; val_acc: 0.471000
(Iteration 1001 / 39200) loss: 1.560481
(Iteration 1101 / 39200) loss: 1.402321
(Iteration 1201 / 39200) loss: 1.559676
(Iteration 1301 / 39200) loss: 1.514315
(Iteration 1401 / 39200) loss: 1.365991
(Epoch 3 / 80) train acc: 0.503000; val_acc: 0.500000
(Iteration 1501 / 39200) loss: 1.577194
(Iteration 1601 / 39200) loss: 1.380106
(Iteration 1701 / 39200) loss: 1.449796
(Iteration 1801 / 39200) loss: 1.444384
(Iteration 1901 / 39200) loss: 1.465490
(Epoch 4 / 80) train acc: 0.519000; val_acc: 0.518000
(Iteration 2001 / 39200) loss: 1.448264
(Iteration 2101 / 39200) loss: 1.407761
(Iteration 2201 / 39200) loss: 1.635408
(Iteration 2301 / 39200) loss: 1.491990
(Iteration 2401 / 39200) loss: 1.442680
(Epoch 5 / 80) train acc: 0.536000; val_acc: 0.526000
(Iteration 2501 / 39200) loss: 1.322390
(Iteration 2601 / 39200) loss: 1.230148
(Iteration 2701 / 39200) loss: 1.390268
(Iteration 2801 / 39200) loss: 1.354062
(Iteration 2901 / 39200) loss: 1.368472
(Epoch 6 / 80) train acc: 0.522000; val_acc: 0.520000
(Iteration 3001 / 39200) loss: 1.486219
(Iteration 3101 / 39200) loss: 1.269545
(Iteration 3201 / 39200) loss: 1.411731
(Iteration 3301 / 39200) loss: 1.544854
(Iteration 3401 / 39200) loss: 1.208551
(Epoch 7 / 80) train acc: 0.559000; val_acc: 0.538000
(Iteration 3501 / 39200) loss: 1.357756
(Iteration 3601 / 39200) loss: 1.460392
(Iteration 3701 / 39200) loss: 1.351261
(Iteration 3801 / 39200) loss: 1.317291
(Iteration 3901 / 39200) loss: 1.375946
```

(Epoch 8 / 80) train acc: 0.593000; val\_acc: 0.523000  
(Iteration 4001 / 39200) loss: 1.535556  
(Iteration 4101 / 39200) loss: 1.370858  
(Iteration 4201 / 39200) loss: 1.463170  
(Iteration 4301 / 39200) loss: 1.335995  
(Iteration 4401 / 39200) loss: 1.238724  
(Epoch 9 / 80) train acc: 0.603000; val\_acc: 0.543000  
(Iteration 4501 / 39200) loss: 1.143931  
(Iteration 4601 / 39200) loss: 1.250761  
(Iteration 4701 / 39200) loss: 1.309563  
(Iteration 4801 / 39200) loss: 1.253994  
(Epoch 10 / 80) train acc: 0.601000; val\_acc: 0.554000  
(Iteration 4901 / 39200) loss: 1.295024  
(Iteration 5001 / 39200) loss: 1.377004  
(Iteration 5101 / 39200) loss: 1.171046  
(Iteration 5201 / 39200) loss: 1.235158  
(Iteration 5301 / 39200) loss: 1.346153  
(Epoch 11 / 80) train acc: 0.610000; val\_acc: 0.562000  
(Iteration 5401 / 39200) loss: 1.265616  
(Iteration 5501 / 39200) loss: 1.114638  
(Iteration 5601 / 39200) loss: 1.352518  
(Iteration 5701 / 39200) loss: 1.313555  
(Iteration 5801 / 39200) loss: 1.357678  
(Epoch 12 / 80) train acc: 0.609000; val\_acc: 0.571000  
(Iteration 5901 / 39200) loss: 1.306446  
(Iteration 6001 / 39200) loss: 1.259089  
(Iteration 6101 / 39200) loss: 1.183187  
(Iteration 6201 / 39200) loss: 1.309481  
(Iteration 6301 / 39200) loss: 1.226226  
(Epoch 13 / 80) train acc: 0.604000; val\_acc: 0.572000  
(Iteration 6401 / 39200) loss: 1.293294  
(Iteration 6501 / 39200) loss: 1.186847  
(Iteration 6601 / 39200) loss: 1.171038  
(Iteration 6701 / 39200) loss: 1.268189  
(Iteration 6801 / 39200) loss: 1.128496  
(Epoch 14 / 80) train acc: 0.622000; val\_acc: 0.564000  
(Iteration 6901 / 39200) loss: 1.195287  
(Iteration 7001 / 39200) loss: 1.199769  
(Iteration 7101 / 39200) loss: 1.265404  
(Iteration 7201 / 39200) loss: 1.083662  
(Iteration 7301 / 39200) loss: 1.171192  
(Epoch 15 / 80) train acc: 0.655000; val\_acc: 0.559000  
(Iteration 7401 / 39200) loss: 1.132624  
(Iteration 7501 / 39200) loss: 1.482053  
(Iteration 7601 / 39200) loss: 1.218198  
(Iteration 7701 / 39200) loss: 1.353706  
(Iteration 7801 / 39200) loss: 1.156821  
(Epoch 16 / 80) train acc: 0.653000; val\_acc: 0.571000  
(Iteration 7901 / 39200) loss: 1.110690  
(Iteration 8001 / 39200) loss: 1.061102  
(Iteration 8101 / 39200) loss: 1.280265  
(Iteration 8201 / 39200) loss: 1.259501  
(Iteration 8301 / 39200) loss: 1.400559  
(Epoch 17 / 80) train acc: 0.663000; val\_acc: 0.558000  
(Iteration 8401 / 39200) loss: 1.351456  
(Iteration 8501 / 39200) loss: 1.185197  
(Iteration 8601 / 39200) loss: 1.114226  
(Iteration 8701 / 39200) loss: 1.235766  
(Iteration 8801 / 39200) loss: 1.042444  
(Epoch 18 / 80) train acc: 0.645000; val\_acc: 0.564000  
(Iteration 8901 / 39200) loss: 1.190117  
(Iteration 9001 / 39200) loss: 1.151252  
(Iteration 9101 / 39200) loss: 1.242444  
(Iteration 9201 / 39200) loss: 1.226552  
(Iteration 9301 / 39200) loss: 1.112046  
(Epoch 19 / 80) train acc: 0.686000; val\_acc: 0.573000  
(Iteration 9401 / 39200) loss: 1.085735  
(Iteration 9501 / 39200) loss: 1.150893  
(Iteration 9601 / 39200) loss: 1.142801  
(Iteration 9701 / 39200) loss: 1.243212  
(Epoch 20 / 80) train acc: 0.688000; val\_acc: 0.580000  
(Iteration 9801 / 39200) loss: 1.100593  
(Iteration 9901 / 39200) loss: 1.195339  
(Iteration 10001 / 39200) loss: 1.455515  
(Iteration 10101 / 39200) loss: 1.102199  
(Iteration 10201 / 39200) loss: 1.211633  
(Epoch 21 / 80) train acc: 0.689000; val\_acc: 0.567000  
(Iteration 10301 / 39200) loss: 1.194243  
(Iteration 10401 / 39200) loss: 1.110677  
(Iteration 10501 / 39200) loss: 1.151281  
(Iteration 10601 / 39200) loss: 1.167924  
(Iteration 10701 / 39200) loss: 1.079824  
(Epoch 22 / 80) train acc: 0.693000; val\_acc: 0.569000  
(Iteration 10801 / 39200) loss: 1.076431  
(Iteration 10901 / 39200) loss: 1.072421  
(Iteration 11001 / 39200) loss: 1.313280  
(Iteration 11101 / 39200) loss: 1.213015  
(Iteration 11201 / 39200) loss: 1.004086

(Epoch 23 / 80) train acc: 0.691000; val\_acc: 0.576000  
(Iteration 11301 / 39200) loss: 0.811971  
(Iteration 11401 / 39200) loss: 0.999775  
(Iteration 11501 / 39200) loss: 1.087524  
(Iteration 11601 / 39200) loss: 0.891453  
(Iteration 11701 / 39200) loss: 0.944478  
(Epoch 24 / 80) train acc: 0.712000; val\_acc: 0.576000  
(Iteration 11801 / 39200) loss: 1.135463  
(Iteration 11901 / 39200) loss: 1.051916  
(Iteration 12001 / 39200) loss: 1.058195  
(Iteration 12101 / 39200) loss: 1.094761  
(Iteration 12201 / 39200) loss: 1.218822  
(Epoch 25 / 80) train acc: 0.700000; val\_acc: 0.587000  
(Iteration 12301 / 39200) loss: 1.095869  
(Iteration 12401 / 39200) loss: 0.981534  
(Iteration 12501 / 39200) loss: 1.040588  
(Iteration 12601 / 39200) loss: 1.101972  
(Iteration 12701 / 39200) loss: 0.979700  
(Epoch 26 / 80) train acc: 0.727000; val\_acc: 0.580000  
(Iteration 12801 / 39200) loss: 0.884404  
(Iteration 12901 / 39200) loss: 1.039024  
(Iteration 13001 / 39200) loss: 0.881299  
(Iteration 13101 / 39200) loss: 1.123838  
(Iteration 13201 / 39200) loss: 1.139810  
(Epoch 27 / 80) train acc: 0.726000; val\_acc: 0.585000  
(Iteration 13301 / 39200) loss: 0.981287  
(Iteration 13401 / 39200) loss: 0.906821  
(Iteration 13501 / 39200) loss: 0.915128  
(Iteration 13601 / 39200) loss: 1.141380  
(Iteration 13701 / 39200) loss: 0.944810  
(Epoch 28 / 80) train acc: 0.730000; val\_acc: 0.588000  
(Iteration 13801 / 39200) loss: 0.835803  
(Iteration 13901 / 39200) loss: 0.983092  
(Iteration 14001 / 39200) loss: 1.012665  
(Iteration 14101 / 39200) loss: 1.056424  
(Iteration 14201 / 39200) loss: 0.922939  
(Epoch 29 / 80) train acc: 0.750000; val\_acc: 0.573000  
(Iteration 14301 / 39200) loss: 0.889565  
(Iteration 14401 / 39200) loss: 1.063439  
(Iteration 14501 / 39200) loss: 0.716205  
(Iteration 14601 / 39200) loss: 0.810242  
(Epoch 30 / 80) train acc: 0.722000; val\_acc: 0.580000  
(Iteration 14701 / 39200) loss: 1.299068  
(Iteration 14801 / 39200) loss: 0.808461  
(Iteration 14901 / 39200) loss: 0.870385  
(Iteration 15001 / 39200) loss: 1.024490  
(Iteration 15101 / 39200) loss: 1.034015  
(Epoch 31 / 80) train acc: 0.754000; val\_acc: 0.586000  
(Iteration 15201 / 39200) loss: 1.134013  
(Iteration 15301 / 39200) loss: 0.978965  
(Iteration 15401 / 39200) loss: 0.951913  
(Iteration 15501 / 39200) loss: 1.156696  
(Iteration 15601 / 39200) loss: 0.845961  
(Epoch 32 / 80) train acc: 0.750000; val\_acc: 0.582000  
(Iteration 15701 / 39200) loss: 0.966734  
(Iteration 15801 / 39200) loss: 0.742249  
  
(Iteration 15901 / 39200) loss: 0.811808  
(Iteration 16001 / 39200) loss: 1.023509  
(Iteration 16101 / 39200) loss: 0.903440  
(Epoch 33 / 80) train acc: 0.727000; val\_acc: 0.586000  
(Iteration 16201 / 39200) loss: 0.947621  
(Iteration 16301 / 39200) loss: 0.775905  
(Iteration 16401 / 39200) loss: 1.019085  
(Iteration 16501 / 39200) loss: 0.961113  
(Iteration 16601 / 39200) loss: 0.977766  
(Epoch 34 / 80) train acc: 0.742000; val\_acc: 0.583000  
(Iteration 16701 / 39200) loss: 0.783518  
(Iteration 16801 / 39200) loss: 1.039373  
(Iteration 16901 / 39200) loss: 0.858510  
(Iteration 17001 / 39200) loss: 0.866176  
(Iteration 17101 / 39200) loss: 0.983566  
(Epoch 35 / 80) train acc: 0.755000; val\_acc: 0.588000  
(Iteration 17201 / 39200) loss: 1.048066  
(Iteration 17301 / 39200) loss: 0.989369  
(Iteration 17401 / 39200) loss: 1.015185  
(Iteration 17501 / 39200) loss: 1.068445  
(Iteration 17601 / 39200) loss: 0.851720  
(Epoch 36 / 80) train acc: 0.755000; val\_acc: 0.585000  
(Iteration 17701 / 39200) loss: 1.036607  
(Iteration 17801 / 39200) loss: 0.897850  
(Iteration 17901 / 39200) loss: 0.908984  
(Iteration 18001 / 39200) loss: 0.939887  
(Iteration 18101 / 39200) loss: 0.835395  
(Epoch 37 / 80) train acc: 0.767000; val\_acc: 0.578000  
(Iteration 18201 / 39200) loss: 1.042282  
(Iteration 18301 / 39200) loss: 0.758919  
(Iteration 18401 / 39200) loss: 1.002052  
(Iteration 18501 / 39200) loss: 0.841024

(Iteration 18601 / 39200) loss: 0.882245  
(Epoch 38 / 80) train acc: 0.765000; val\_acc: 0.594000  
(Iteration 18701 / 39200) loss: 0.958950  
(Iteration 18801 / 39200) loss: 0.851094  
(Iteration 18901 / 39200) loss: 0.708583  
(Iteration 19001 / 39200) loss: 0.770913  
(Iteration 19101 / 39200) loss: 0.733707  
(Epoch 39 / 80) train acc: 0.778000; val\_acc: 0.590000  
(Iteration 19201 / 39200) loss: 0.990309  
(Iteration 19301 / 39200) loss: 0.958403  
(Iteration 19401 / 39200) loss: 0.992356  
(Iteration 19501 / 39200) loss: 1.005794  
(Epoch 40 / 80) train acc: 0.752000; val\_acc: 0.588000  
(Iteration 19601 / 39200) loss: 1.032480  
(Iteration 19701 / 39200) loss: 1.035892  
(Iteration 19801 / 39200) loss: 0.906792  
(Iteration 19901 / 39200) loss: 0.895575  
(Iteration 20001 / 39200) loss: 0.825461  
(Epoch 41 / 80) train acc: 0.765000; val\_acc: 0.591000  
(Iteration 20101 / 39200) loss: 0.876272  
(Iteration 20201 / 39200) loss: 0.902269  
(Iteration 20301 / 39200) loss: 0.834744  
(Iteration 20401 / 39200) loss: 0.961422  
(Iteration 20501 / 39200) loss: 0.975309  
(Epoch 42 / 80) train acc: 0.795000; val\_acc: 0.585000  
(Iteration 20601 / 39200) loss: 1.038000  
(Iteration 20701 / 39200) loss: 0.837005  
(Iteration 20801 / 39200) loss: 0.717900  
(Iteration 20901 / 39200) loss: 0.850868  
(Iteration 21001 / 39200) loss: 1.017165  
(Epoch 43 / 80) train acc: 0.752000; val\_acc: 0.585000  
(Iteration 21101 / 39200) loss: 0.905929  
(Iteration 21201 / 39200) loss: 1.025768  
(Iteration 21301 / 39200) loss: 0.923585  
(Iteration 21401 / 39200) loss: 0.892869  
(Iteration 21501 / 39200) loss: 0.887772  
(Epoch 44 / 80) train acc: 0.777000; val\_acc: 0.590000  
(Iteration 21601 / 39200) loss: 0.925660  
(Iteration 21701 / 39200) loss: 0.861675  
(Iteration 21801 / 39200) loss: 0.809872  
(Iteration 21901 / 39200) loss: 0.920505  
(Iteration 22001 / 39200) loss: 1.009144  
(Epoch 45 / 80) train acc: 0.786000; val\_acc: 0.591000  
(Iteration 22101 / 39200) loss: 0.856028  
(Iteration 22201 / 39200) loss: 0.866785  
(Iteration 22301 / 39200) loss: 0.830510  
(Iteration 22401 / 39200) loss: 1.279860  
(Iteration 22501 / 39200) loss: 0.971482  
(Epoch 46 / 80) train acc: 0.806000; val\_acc: 0.593000  
(Iteration 22601 / 39200) loss: 0.905765  
(Iteration 22701 / 39200) loss: 1.066834  
(Iteration 22801 / 39200) loss: 0.794572  
(Iteration 22901 / 39200) loss: 1.013506  
(Iteration 23001 / 39200) loss: 1.105925  
(Epoch 47 / 80) train acc: 0.797000; val\_acc: 0.591000  
(Iteration 23101 / 39200) loss: 0.908778  
(Iteration 23201 / 39200) loss: 0.864585  
(Iteration 23301 / 39200) loss: 0.902531  
(Iteration 23401 / 39200) loss: 0.800342  
(Iteration 23501 / 39200) loss: 0.721073  
(Epoch 48 / 80) train acc: 0.781000; val\_acc: 0.595000  
(Iteration 23601 / 39200) loss: 0.840992  
(Iteration 23701 / 39200) loss: 1.122623  
(Iteration 23801 / 39200) loss: 1.032104  
(Iteration 23901 / 39200) loss: 0.909329  
(Iteration 24001 / 39200) loss: 0.642137  
(Epoch 49 / 80) train acc: 0.799000; val\_acc: 0.593000  
(Iteration 24101 / 39200) loss: 1.074038  
(Iteration 24201 / 39200) loss: 0.888440  
(Iteration 24301 / 39200) loss: 0.732228  
(Iteration 24401 / 39200) loss: 0.830554  
(Epoch 50 / 80) train acc: 0.803000; val\_acc: 0.588000  
(Iteration 24501 / 39200) loss: 0.690772  
(Iteration 24601 / 39200) loss: 1.048687  
(Iteration 24701 / 39200) loss: 1.012022  
(Iteration 24801 / 39200) loss: 0.850425  
(Iteration 24901 / 39200) loss: 0.838619  
(Epoch 51 / 80) train acc: 0.796000; val\_acc: 0.588000  
(Iteration 25001 / 39200) loss: 0.654319  
(Iteration 25101 / 39200) loss: 0.797890  
(Iteration 25201 / 39200) loss: 0.852705  
(Iteration 25301 / 39200) loss: 0.917384  
(Iteration 25401 / 39200) loss: 0.845211  
(Epoch 52 / 80) train acc: 0.813000; val\_acc: 0.595000  
(Iteration 25501 / 39200) loss: 0.753940  
(Iteration 25601 / 39200) loss: 0.997589  
(Iteration 25701 / 39200) loss: 0.775196  
(Iteration 25801 / 39200) loss: 0.917092

(Iteration 25901 / 39200) loss: 0.954966  
(Epoch 53 / 80) train acc: 0.785000; val\_acc: 0.598000  
(Iteration 26001 / 39200) loss: 0.828941  
(Iteration 26101 / 39200) loss: 0.922264  
(Iteration 26201 / 39200) loss: 0.748944  
(Iteration 26301 / 39200) loss: 0.830192  
(Iteration 26401 / 39200) loss: 0.816747  
(Epoch 54 / 80) train acc: 0.783000; val\_acc: 0.591000  
(Iteration 26501 / 39200) loss: 0.682134  
(Iteration 26601 / 39200) loss: 0.801736  
(Iteration 26701 / 39200) loss: 0.749627  
(Iteration 26801 / 39200) loss: 1.052627  
(Iteration 26901 / 39200) loss: 0.990648  
(Epoch 55 / 80) train acc: 0.816000; val\_acc: 0.597000  
(Iteration 27001 / 39200) loss: 0.950973  
(Iteration 27101 / 39200) loss: 0.734878  
(Iteration 27201 / 39200) loss: 0.747308  
(Iteration 27301 / 39200) loss: 0.884391  
(Iteration 27401 / 39200) loss: 0.918879  
(Epoch 56 / 80) train acc: 0.797000; val\_acc: 0.594000  
(Iteration 27501 / 39200) loss: 0.853080  
(Iteration 27601 / 39200) loss: 0.990341  
(Iteration 27701 / 39200) loss: 0.884415  
(Iteration 27801 / 39200) loss: 0.795938  
(Iteration 27901 / 39200) loss: 0.919736  
(Epoch 57 / 80) train acc: 0.829000; val\_acc: 0.592000  
(Iteration 28001 / 39200) loss: 0.795244  
(Iteration 28101 / 39200) loss: 0.804721  
(Iteration 28201 / 39200) loss: 0.958718  
(Iteration 28301 / 39200) loss: 1.037316  
(Iteration 28401 / 39200) loss: 0.912411  
(Epoch 58 / 80) train acc: 0.795000; val\_acc: 0.598000  
(Iteration 28501 / 39200) loss: 0.786430  
(Iteration 28601 / 39200) loss: 0.824360  
(Iteration 28701 / 39200) loss: 0.857696  
(Iteration 28801 / 39200) loss: 0.716881  
(Iteration 28901 / 39200) loss: 0.745188  
(Epoch 59 / 80) train acc: 0.827000; val\_acc: 0.599000  
(Iteration 29001 / 39200) loss: 0.836852  
(Iteration 29101 / 39200) loss: 0.952089  
(Iteration 29201 / 39200) loss: 0.784935  
(Iteration 29301 / 39200) loss: 0.845400  
(Epoch 60 / 80) train acc: 0.826000; val\_acc: 0.586000  
(Iteration 29401 / 39200) loss: 0.868991  
(Iteration 29501 / 39200) loss: 0.695192  
(Iteration 29601 / 39200) loss: 0.853898  
(Iteration 29701 / 39200) loss: 0.978458  
(Iteration 29801 / 39200) loss: 0.727721  
(Epoch 61 / 80) train acc: 0.802000; val\_acc: 0.587000  
(Iteration 29901 / 39200) loss: 0.705634  
(Iteration 30001 / 39200) loss: 0.829131  
(Iteration 30101 / 39200) loss: 0.955335  
(Iteration 30201 / 39200) loss: 0.861911  
(Iteration 30301 / 39200) loss: 0.960021  
(Epoch 62 / 80) train acc: 0.823000; val\_acc: 0.583000  
(Iteration 30401 / 39200) loss: 1.059026  
(Iteration 30501 / 39200) loss: 0.820306  
(Iteration 30601 / 39200) loss: 0.889187  
(Iteration 30701 / 39200) loss: 0.888176  
(Iteration 30801 / 39200) loss: 0.701103  
(Epoch 63 / 80) train acc: 0.816000; val\_acc: 0.587000  
(Iteration 30901 / 39200) loss: 0.750650  
(Iteration 31001 / 39200) loss: 0.976211  
(Iteration 31101 / 39200) loss: 0.729728  
(Iteration 31201 / 39200) loss: 0.829218  
(Iteration 31301 / 39200) loss: 0.711588  
(Epoch 64 / 80) train acc: 0.818000; val\_acc: 0.588000  
(Iteration 31401 / 39200) loss: 0.781439  
(Iteration 31501 / 39200) loss: 0.746261  
  
(Iteration 31601 / 39200) loss: 0.707213  
(Iteration 31701 / 39200) loss: 0.833365  
(Iteration 31801 / 39200) loss: 0.814129  
(Epoch 65 / 80) train acc: 0.831000; val\_acc: 0.578000  
(Iteration 31901 / 39200) loss: 0.823512  
(Iteration 32001 / 39200) loss: 0.797918  
(Iteration 32101 / 39200) loss: 0.741748  
(Iteration 32201 / 39200) loss: 0.908855  
(Iteration 32301 / 39200) loss: 1.023562  
(Epoch 66 / 80) train acc: 0.820000; val\_acc: 0.587000  
(Iteration 32401 / 39200) loss: 1.022008  
(Iteration 32501 / 39200) loss: 0.963455  
(Iteration 32601 / 39200) loss: 0.807658  
(Iteration 32701 / 39200) loss: 0.779625  
(Iteration 32801 / 39200) loss: 1.084706  
(Epoch 67 / 80) train acc: 0.814000; val\_acc: 0.590000  
(Iteration 32901 / 39200) loss: 0.747726  
(Iteration 33001 / 39200) loss: 0.897594  
(Iteration 33101 / 39200) loss: 0.958651

```

(Iteration 33201 / 39200) loss: 0.731197
(Iteration 33301 / 39200) loss: 1.067751
(Epoch 68 / 80) train acc: 0.810000; val_acc: 0.593000
(Iteration 33401 / 39200) loss: 0.932779
(Iteration 33501 / 39200) loss: 0.934577
(Iteration 33601 / 39200) loss: 0.699594
(Iteration 33701 / 39200) loss: 0.728579
(Iteration 33801 / 39200) loss: 0.832864
(Epoch 69 / 80) train acc: 0.801000; val_acc: 0.591000
(Iteration 33901 / 39200) loss: 0.767292
(Iteration 34001 / 39200) loss: 0.882040
(Iteration 34101 / 39200) loss: 0.705112
(Iteration 34201 / 39200) loss: 0.948548
(Epoch 70 / 80) train acc: 0.813000; val_acc: 0.586000
(Iteration 34301 / 39200) loss: 0.759878
(Iteration 34401 / 39200) loss: 0.910183
(Iteration 34501 / 39200) loss: 0.665622
(Iteration 34601 / 39200) loss: 0.957756
(Iteration 34701 / 39200) loss: 0.611733
(Epoch 71 / 80) train acc: 0.810000; val_acc: 0.594000
(Iteration 34801 / 39200) loss: 0.929759
(Iteration 34901 / 39200) loss: 0.701143
(Iteration 35001 / 39200) loss: 0.796047
(Iteration 35101 / 39200) loss: 0.673842
(Iteration 35201 / 39200) loss: 0.706945
(Epoch 72 / 80) train acc: 0.825000; val_acc: 0.587000
(Iteration 35301 / 39200) loss: 0.809401
(Iteration 35401 / 39200) loss: 0.721291
(Iteration 35501 / 39200) loss: 0.876259
(Iteration 35601 / 39200) loss: 0.836721
(Iteration 35701 / 39200) loss: 0.792897
(Epoch 73 / 80) train acc: 0.835000; val_acc: 0.589000
(Iteration 35801 / 39200) loss: 1.015434
(Iteration 35901 / 39200) loss: 0.775106
(Iteration 36001 / 39200) loss: 0.770852
(Iteration 36101 / 39200) loss: 0.818432
(Iteration 36201 / 39200) loss: 0.903966
(Epoch 74 / 80) train acc: 0.815000; val_acc: 0.597000
(Iteration 36301 / 39200) loss: 0.890942
(Iteration 36401 / 39200) loss: 0.761177
(Iteration 36501 / 39200) loss: 0.618575
(Iteration 36601 / 39200) loss: 0.836164
(Iteration 36701 / 39200) loss: 0.791301
(Epoch 75 / 80) train acc: 0.832000; val_acc: 0.592000
(Iteration 36801 / 39200) loss: 0.911602
(Iteration 36901 / 39200) loss: 0.706361
(Iteration 37001 / 39200) loss: 0.727411
(Iteration 37101 / 39200) loss: 0.721145
(Iteration 37201 / 39200) loss: 0.926907
(Epoch 76 / 80) train acc: 0.836000; val_acc: 0.594000
(Iteration 37301 / 39200) loss: 0.798053
(Iteration 37401 / 39200) loss: 0.893760
(Iteration 37501 / 39200) loss: 0.773815
(Iteration 37601 / 39200) loss: 0.768303
(Iteration 37701 / 39200) loss: 0.823549
(Epoch 77 / 80) train acc: 0.820000; val_acc: 0.590000
(Iteration 37801 / 39200) loss: 0.743476
(Iteration 37901 / 39200) loss: 0.719722
(Iteration 38001 / 39200) loss: 0.703971
(Iteration 38101 / 39200) loss: 0.725379
(Iteration 38201 / 39200) loss: 1.030217
(Epoch 78 / 80) train acc: 0.813000; val_acc: 0.591000
(Iteration 38301 / 39200) loss: 0.901584
(Iteration 38401 / 39200) loss: 1.047116
(Iteration 38501 / 39200) loss: 0.754022
(Iteration 38601 / 39200) loss: 0.612967
(Iteration 38701 / 39200) loss: 0.597829
(Epoch 79 / 80) train acc: 0.830000; val_acc: 0.590000
(Iteration 38801 / 39200) loss: 0.771769
(Iteration 38901 / 39200) loss: 0.576011
(Iteration 39001 / 39200) loss: 0.735640
(Iteration 39101 / 39200) loss: 0.777838
(Epoch 80 / 80) train acc: 0.830000; val_acc: 0.590000
Validation set accuracy: 0.604
Test set accuracy: 0.597

```

```

1 import numpy as np
2 import pdb
3
4 """
5 This code was originally written for CS 231n at Stanford University
6 (cs231n.stanford.edu). It has been modified in various areas for use in the
7 ECE 239AS class at UCLA. This includes the descriptions of what code to
8 implement as well as some slight potential changes in variable names to be
9 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung
10 for
11 permission to use this code. To see the original version, please visit
12 cs231n.stanford.edu.
13 """
14 def affine_forward(x, w, b):
15     """
16     Computes the forward pass for an affine (fully-connected) layer.
17
18     The input x has shape (N, d_1, ..., d_k) and contains a minibatch of N
19     examples, where each example x[i] has shape (d_1, ..., d_k). We will
20     reshape each input into a vector of dimension D = d_1 * ... * d_k, and
21     then transform it to an output vector of dimension M.
22
23     Inputs:
24     - x: A numpy array containing input data, of shape (N, d_1, ..., d_k)
25     - w: A numpy array of weights, of shape (D, M)
26     - b: A numpy array of biases, of shape (M,)
27
28     Returns a tuple of:
29     - out: output, of shape (N, M)
30     - cache: (x, w, b)
31     """
32
33     # ===== #
34     # YOUR CODE HERE:
35     # Calculate the output of the forward pass. Notice the dimensions
36     # of w are D x M, which is the transpose of what we did in earlier
37     # assignments.
38     # ===== #
39
40     x_reshape = x.reshape((x.shape[0], w.shape[0])) # N * D
41     out = x_reshape.dot(w) + b.reshape((1, b.shape[0])) # N * M
42
43     # ===== #
44     # END YOUR CODE HERE
45     # ===== #
46
47     cache = (x, w, b)
48     return out, cache
49
50 def affine_backward(dout, cache):
51     """
52     Computes the backward pass for an affine layer.
53
54     Inputs:
55     - dout: Upstream derivative, of shape (N, M)
56     - cache: Tuple of:
57       - x: Input data, of shape (N, d_1, ..., d_k)
58       - w: Weights, of shape (D, M)
59
60     Returns a tuple of:
61     - dx: Gradient with respect to x, of shape (N, d_1, ..., d_k)
62     - dw: Gradient with respect to w, of shape (D, M)
63     - db: Gradient with respect to b, of shape (M,)
64     """
65
66     x, w, b = cache
67     dx, dw, db = None, None, None
68
69     # ===== #
70     # YOUR CODE HERE:
71     # Calculate the gradients for the backward pass.
72     # ===== #
73
74     x_reshape = np.reshape(x, (x.shape[0], w.shape[0]))
75     dx_reshape = dout.dot(w.T)
76     dx = np.reshape(dx_reshape, x.shape) # N * D
77     dw = x_reshape.T.dot(dout) # D * M
78     db = dout.T.dot(np.ones(x.shape[0])) # M * 1
79
80     # ===== #
81     # END YOUR CODE HERE
82     # ===== #
83

```

```

84     return dx, dw, db
85
86 def relu_forward(x):
87     """
88     Computes the forward pass for a layer of rectified linear units (ReLUs).
89
90     Input:
91     - x: Inputs, of any shape
92
93     Returns a tuple of:
94     - out: Output, of the same shape as x
95     - cache: x
96     """
97     # ===== #
98     # YOUR CODE HERE:
99     #   Implement the ReLU forward pass.
100    # ===== #
101
102    out = np.maximum(0, x)
103
104    # ===== #
105    # END YOUR CODE HERE
106    # ===== #
107
108    cache = x
109    return out, cache
110
111 def relu_backward(dout, cache):
112     """
113     Computes the backward pass for a layer of rectified linear units (ReLUs).
114
115     Input:
116     - dout: Upstream derivatives, of any shape
117     - cache: Input x, of same shape as dout
118
119     Returns:
120     - dx: Gradient with respect to x
121     """
122
123     x = cache
124
125     # ===== #
126     # YOUR CODE HERE:
127     #   Implement the ReLU backward pass
128     # ===== #
129
130     dx = (x > 0) * (dout)
131
132     # ===== #
133     # END YOUR CODE HERE
134     # ===== #
135
136     return dx
137
138 def batchnorm_forward(x, gamma, beta, bn_param):
139     """
140     Forward pass for batch normalization.
141
142     During training the sample mean and (uncorrected) sample variance are
143     computed from minibatch statistics and used to normalize the incoming data.
144     During training we also keep an exponentially decaying running mean of the
145     mean
146     and variance of each feature, and these averages are used to normalize data
147     at test-time.
148
149     At each timestep we update the running averages for mean and variance using
150     an exponential decay based on the momentum parameter:
151
152     running_mean = momentum * running_mean + (1 - momentum) * sample_mean
153     running_var = momentum * running_var + (1 - momentum) * sample_var
154
155     Note that the batch normalization paper suggests a different test-time
156     behavior: they compute sample mean and variance for each feature using a
157     large number of training images rather than using a running average. For
158     this implementation we have chosen to use running averages instead since
159     they do not require an additional estimation step; the torch7
160     implementation
161     of batch normalization also uses running averages.
162
163     Input:
164     - x: Data of shape (N, D)
165     - gamma: Scale parameter of shape (D,)
166     - beta: Shift parameter of shape (D,)
167     - bn_param: Dictionary with the following keys:

```



```

166 - mode: 'train' or 'test'; required
167 - eps: Constant for numeric stability
168 - momentum: Constant for running mean / variance.
169 - running_mean: Array of shape (D,) giving running mean of features
170 - running_var Array of shape (D,) giving running variance of features
171
172 Returns a tuple of:
173 - out: of shape (N, D)
174 - cache: A tuple of values needed in the backward pass
175 """
176 mode = bn_param['mode']
177 eps = bn_param.get('eps', 1e-5)
178 momentum = bn_param.get('momentum', 0.9)
179
180 N, D = x.shape
181 running_mean = bn_param.get('running_mean', np.zeros(D, dtype=x.dtype))
182 running_var = bn_param.get('running_var', np.zeros(D, dtype=x.dtype))
183
184 out, cache = None, None
185 if mode == 'train':
186
187     # ===== #
188     # YOUR CODE HERE:
189     #   A few steps here:
190     #   (1) Calculate the running mean and variance of the minibatch.
191     #   (2) Normalize the activations with the batch mean and variance.
192     #   (3) Scale and shift the normalized activations. Store this
193     #       as the variable 'out'
194     #   (4) Store any variables you may need for the backward pass in
195     #       the 'cache' variable.
196     # ===== #
197
198     minibatch_mean = np.mean(x, axis=0)
199     minibatch_var = np.var(x, axis=0)
200     x_normalize = (x - minibatch_mean) / np.sqrt(minibatch_var + eps)
201     out = gamma * x_normalize + beta
202
203     running_mean = momentum * running_mean + (1 - momentum) * minibatch_mean
204     running_var = momentum * running_var + (1 - momentum) * minibatch_var
205     bn_param['running_mean'] = running_mean
206     bn_param['running_var'] = running_var
207
208     cache = {
209         'minibatch_var': minibatch_var,
210         'x_centralize': (x - minibatch_mean),
211         'x_normalize': x_normalize,
212         'gamma': gamma,
213         'eps': eps
214     }
215
216     # ===== #
217     # END YOUR CODE HERE
218     # ===== #
219
220 elif mode == 'test':
221
222     # ===== #
223     # YOUR CODE HERE:
224     #   Calculate the testing time normalized activations. Normalize using
225     #   the running mean and variance, and then scale and shift
226     #   appropriately.
227     #   Store the output as 'out'.
228     # ===== #
229
230     out = gamma * (x - running_mean) / np.sqrt(running_var + eps) + beta
231
232     # ===== #
233     # END YOUR CODE HERE
234     # ===== #
235
236 else:
237     raise ValueError('Invalid forward batchnorm mode "%s"' % mode)
238
239 # Store the updated running means back into bn_param
240 bn_param['running_mean'] = running_mean
241 bn_param['running_var'] = running_var
242
243 return out, cache
244
245 def batchnorm_backward(dout, cache):
246     """
247     Backward pass for batch normalization.
248
249     For this implementation, you should write out a computation graph for

```

```

249 batch normalization on paper and propagate gradients backward through
250 intermediate nodes.
251
252 Inputs:
253 - dout: Upstream derivatives, of shape (N, D)
254 - cache: Variable of intermediates from batchnorm_forward.
255
256 Returns a tuple of:
257 - dx: Gradient with respect to inputs x, of shape (N, D)
258 - dgamma: Gradient with respect to scale parameter gamma, of shape (D,)
259 - dbeta: Gradient with respect to shift parameter beta, of shape (D,)
260 """
261 dx, dgamma, dbeta = None, None, None
262
263 # ===== #
264 # YOUR CODE HERE:
265 # Implement the batchnorm backward pass, calculating dx, dgamma, and
266 dbeta.
267 # ===== #
268
269 # get parameters from cache
270 N = dout.shape[0]
271 minibatch_var = cache.get('minibatch_var')
272 x_centralize = cache.get('x_centralize')
273 x_normalize = cache.get('x_normalize')
274 gamma = cache.get('gamma')
275 eps = cache.get('eps')
276
277 # calculate dx
278 dxhat = dout * gamma
279 dxmu1 = dxhat / np.sqrt(minibatch_var + eps)
280 sqrt_var = np.sqrt(minibatch_var + eps)
281 dsqrt_var = -np.sum(dxhat * x_centralize, axis=0) / (sqrt_var**2)
282 dvar = dsqrt_var * 0.5 / sqrt_var
283 dxmu2 = 2 * x_centralize * dvar * np.ones_like(dout) / N
284 dx1 = dxmu1 + dxmu2
285 dx2 = -np.sum(dx1, axis=0) * np.ones_like(dout) / N
286 dx = dx1 + dx2
287
288 # calculate dbeta and dgamma
289 dbeta = np.sum(dout, axis=0)
290 dgamma = np.sum(dout * x_normalize, axis=0)
291
292 # ===== #
293 # END YOUR CODE HERE
294 # ===== #
295
296 return dx, dgamma, dbeta
297
298 def dropout_forward(x, dropout_param):
299     """
300     Performs the forward pass for (inverted) dropout.
301
302     Inputs:
303     - x: Input data, of any shape
304     - dropout_param: A dictionary with the following keys:
305       - p: Dropout parameter. We drop each neuron output with probability p.
306       - mode: 'test' or 'train'. If the mode is train, then perform dropout;
307       if the mode is test, then just return the input.
308       - seed: Seed for the random number generator. Passing seed makes this
309       function deterministic, which is needed for gradient checking but not
310       in
311       real networks.
312
313     Outputs:
314     - out: Array of the same shape as x.
315     - cache: A tuple (dropout_param, mask). In training mode, mask is the
316     dropout
317     mask that was used to multiply the input; in test mode, mask is None.
318     """
319     p, mode = dropout_param['p'], dropout_param['mode']
320     if 'seed' in dropout_param:
321         np.random.seed(dropout_param['seed'])
322
323     mask = None
324     out = None
325
326     if mode == 'train':
327         # ===== #
328         # YOUR CODE HERE:
329         # Implement the inverted dropout forward pass during training time.
330         # Store the masked and scaled activations in out, and store the
331         # dropout mask as the variable mask.
332         # ===== #

```

```

330
331     mask = (np.random.random_sample(x.shape) >= p) / (1 - p)
332     out = x * mask
333
334     # ===== #
335     # END YOUR CODE HERE
336     # ===== #
337
338 elif mode == 'test':
339
340     # ===== #
341     # YOUR CODE HERE:
342     # Implement the inverted dropout forward pass during test time.
343     # ===== #
344
345     out = x
346
347     # ===== #
348     # END YOUR CODE HERE
349     # ===== #
350
351     cache = (dropout_param, mask)
352     out = out.astype(x.dtype, copy=False)
353
354     return out, cache
355
356 def dropout_backward(dout, cache):
357     """
358     Perform the backward pass for (inverted) dropout.
359
360     Inputs:
361     - dout: Upstream derivatives, of any shape
362     - cache: (dropout_param, mask) from dropout_forward.
363     """
364     dropout_param, mask = cache
365     mode = dropout_param['mode']
366
367     dx = None
368     if mode == 'train':
369         # ===== #
370         # YOUR CODE HERE:
371         # Implement the inverted dropout backward pass during training time.
372         # ===== #
373
374         dx = dout * mask
375
376         # ===== #
377         # END YOUR CODE HERE
378         # ===== #
379     elif mode == 'test':
380         # ===== #
381         # YOUR CODE HERE:
382         # Implement the inverted dropout backward pass during test time.
383         # ===== #
384
385         dx = dout
386
387         # ===== #
388         # END YOUR CODE HERE
389         # ===== #
390     return dx
391
392 def svm_loss(x, y):
393     """
394     Computes the loss and gradient using for multiclass SVM classification.
395
396     Inputs:
397     - x: Input data, of shape (N, C) where x[i, j] is the score for the jth
398     class for the ith input.
399     - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
400     0 <= y[i] < C
401
402     Returns a tuple of:
403     - loss: Scalar giving the loss
404     - dx: Gradient of the loss with respect to x
405     """
406     N = x.shape[0]
407     correct_class_scores = x[np.arange(N), y]
408     margins = np.maximum(0, x - correct_class_scores[:, np.newaxis] + 1.0)
409     margins[np.arange(N), y] = 0
410     loss = np.sum(margins) / N
411     num_pos = np.sum(margins > 0, axis=1)
412     dx = np.zeros_like(x)

```

```

413 dx[margins > 0] = 1
414 dx[np.arange(N), y] -= num_pos
415 dx /= N
416 return loss, dx
417
418
419 def softmax_loss(x, y):
420     """
421     Computes the loss and gradient for softmax classification.
422
423     Inputs:
424     - x: Input data, of shape (N, C) where x[i, j] is the score for the jth
425       class for the ith input.
426     - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
427       0 <= y[i] < C
428
429     Returns a tuple of:
430     - loss: Scalar giving the loss
431     - dx: Gradient of the loss with respect to x
432     """
433
434     probs = np.exp(x - np.max(x, axis=1, keepdims=True))
435     probs /= np.sum(probs, axis=1, keepdims=True)
436     N = x.shape[0]
437     loss = -np.sum(np.log(probs[np.arange(N), y])) / N
438     dx = probs.copy()
439     dx[np.arange(N), y] -= 1
440     dx /= N
441     return loss, dx
442

```

```

1 import numpy as np
2 import pdb
3
4 from .layers import *
5 from .layer_utils import *
6
7 """
8 This code was originally written for CS 231n at Stanford University
9 (cs231n.stanford.edu). It has been modified in various areas for use in the
10 ECE 239AS class at UCLA. This includes the descriptions of what code to
11 implement as well as some slight potential changes in variable names to be
12 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung
13 for
14 permission to use this code. To see the original version, please visit
15 cs231n.stanford.edu.
16 """
17 class FullyConnectedNet(object):
18     """
19     A fully-connected neural network with an arbitrary number of hidden layers,
20     ReLU nonlinearities, and a softmax loss function. This will also implement
21     dropout and batch normalization as options. For a network with L layers,
22     the architecture will be
23
24     {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax
25
26     where batch normalization and dropout are optional, and the {...} block is
27     repeated L - 1 times.
28
29     Similar to the TwoLayerNet above, learnable parameters are stored in the
30     self.params dictionary and will be learned using the Solver class.
31     """
32
33     def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
34                 dropout=0, use_batchnorm=False, reg=0.0,
35                 weight_scale=1e-2, dtype=np.float32, seed=None):
36         """
37         Initialize a new FullyConnectedNet.
38
39         Inputs:
40         - hidden_dims: A list of integers giving the size of each hidden layer.
41         - input_dim: An integer giving the size of the input.
42         - num_classes: An integer giving the number of classes to classify.
43         - dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0
44         then
45         the network should not use dropout at all.
46         - use_batchnorm: Whether or not the network should use batch
47         normalization.
48         - reg: Scalar giving L2 regularization strength.
49         - weight_scale: Scalar giving the standard deviation for random
50         initialization of the weights.
51         - dtype: A numpy datatype object; all computations will be performed
52         using
53         this datatype. float32 is faster but less accurate, so you should use
54         float64 for numeric gradient checking.
55         - seed: If not None, then pass this random seed to the dropout layers.
56         This
57         will make the dropout layers deterministic so we can gradient check the
58         model.
59         """
60         self.use_batchnorm = use_batchnorm
61         self.use_dropout = dropout > 0
62         self.reg = reg
63         self.num_layers = 1 + len(hidden_dims)
64         self.dtype = dtype
65         self.params = {}
66
67         # ===== #
68         # YOUR CODE HERE:
69         # Initialize all parameters of the network in the self.params
70         dictionary.
71         # The weights and biases of layer 1 are W1 and b1; and in general the
72         # weights and biases of layer i are Wi and bi. The
73         # biases are initialized to zero and the weights are initialized
74         # so that each parameter has mean 0 and standard deviation
75         weight_scale.
76         #
77         # BATCHNORM: Initialize the gammas of each layer to 1 and the beta
78         # parameters to zero. The gamma and beta parameters for layer 1 should
79         # be self.params['gamma1'] and self.params['beta1']. For layer 2, they
80         # should be gamma2 and beta2, etc. Only use batchnorm if
81         self.use_batchnorm
82         # is true and DO NOT batch normalize the output scores.
83         # ===== #

```

```

77
78     cur_dim = input_dim
79     for idx, hidden_dim in enumerate(hidden_dims):
80         # initialize weights and bias
81         self.params['W' + str(idx + 1)] = np.random.randn(cur_dim, hidden_dim)
82 * weight_scale
83         self.params['b' + str(idx + 1)] = np.zeros(hidden_dim)
84
85         # initialize gammas and betas
86         if self.use_batchnorm:
87             self.params['gamma' + str(idx + 1)] = np.ones(hidden_dim)
88             self.params['beta' + str(idx + 1)] = np.zeros(hidden_dim)
89
90         cur_dim = hidden_dim
91
92         self.params['W' + str(self.num_layers)] = np.random.randn(cur_dim,
93 num_classes) * weight_scale
94         self.params['b' + str(self.num_layers)] = np.zeros(num_classes)
95
96         # ===== #
97         # END YOUR CODE HERE
98         # ===== #
99
100         # When using dropout we need to pass a dropout_param dictionary to each
101         # dropout layer so that the layer knows the dropout probability and the
102         # mode
103         # (train / test). You can pass the same dropout_param to each dropout
104         # layer.
105         self.dropout_param = {}
106         if self.use_dropout:
107             self.dropout_param = {'mode': 'train', 'p': dropout}
108             if seed is not None:
109                 self.dropout_param['seed'] = seed
110
111         # With batch normalization we need to keep track of running means and
112         # variances, so we need to pass a special bn_param object to each batch
113         # normalization layer. You should pass self.bn_params[0] to the forward
114         # pass
115         # of the first batch normalization layer, self.bn_params[1] to the
116         # forward
117         # pass of the second batch normalization layer, etc.
118         self.bn_params = []
119         if self.use_batchnorm:
120             self.bn_params = [{'mode': 'train'} for i in np.arange(self.num_layers
121 - 1)]
122
123         # Cast all parameters to the correct datatype
124         for k, v in self.params.items():
125             self.params[k] = v.astype(dtype)
126
127 def loss(self, X, y=None):
128     """
129     Compute loss and gradient for the fully-connected net.
130
131     Input / output: Same as TwoLayerNet above.
132     """
133     X = X.astype(self.dtype)
134     mode = 'test' if y is None else 'train'
135
136     # Set train/test mode for batchnorm params and dropout param since they
137     # behave differently during training and testing.
138     if self.dropout_param is not None:
139         self.dropout_param['mode'] = mode
140     if self.use_batchnorm:
141         for bn_param in self.bn_params:
142             bn_param['mode'] = mode
143
144     scores = None
145
146     # ===== #
147     # YOUR CODE HERE:
148     # Implement the forward pass of the FC net and store the output
149     # scores as the variable "scores".
150     #
151     # BATCHNORM: If self.use_batchnorm is true, insert a bathnorm layer
152     # between the affine_forward and relu_forward layers. You may
153     # also write an affine_batchnorm_relu() function in layer_utils.py.
154     #
155     # DROPOUT: If dropout is non-zero, insert a dropout layer after
156     # every ReLU layer.
157     # ===== #
158
159     # initialize caches

```

```

154     fc_cache = {}
155     relu_cache = {}
156     batchnorm_cache = {}
157     dropout_cache = {}
158
159     # flatten image
160     X = np.reshape(X, [X.shape[0], -1])
161
162     # go through all layers
163     for i in range(self.num_layers - 1):
164         # fc layer
165         fc_out, fc_cache[str(i + 1)] = affine_forward(X, self.params['W' +
str(i + 1)], self.params['b' + str(i + 1)])
166
167         # batchnorm layer
168         relu_input = fc_out
169         if self.use_batchnorm:
170             batchnorm_out, batchnorm_cache[str(i + 1)] =
batchnorm_forward(fc_out, self.params['gamma' + str(i + 1)],
self.params['beta' + str(i + 1)], self.bn_params[i])
171             relu_input = batchnorm_out
172             relu_out, relu_cache[str(i + 1)] = relu_forward(relu_input)
173
174         # dropout layer
175         if self.use_dropout:
176             relu_out, dropout_cache[str(i + 1)] = dropout_forward(relu_out,
self.dropout_param)
177
178         # update X
179         X = relu_out.copy()
180
181         # output FC layer with no relu
182         scores, final_cache = affine_forward(X, self.params['W' +
str(self.num_layers)], self.params['b' + str(self.num_layers)])
183
184         # ===== #
185         # END YOUR CODE HERE
186         # ===== #
187
188         # If test mode return early
189         if mode == 'test':
190             return scores
191
192         loss, grads = 0.0, {}
193         # ===== #
194         # YOUR CODE HERE:
195         # Implement the backwards pass of the FC net and store the gradients
196         # in the grads dict, so that grads[k] is the gradient of self.params[k]
197         # Be sure your L2 regularization includes a 0.5 factor.
198         #
199         # BATCHNORM: Incorporate the backward pass of the batchnorm.
200         #
201         # DROPOUT: Incorporate the backward pass of dropout.
202         # ===== #
203
204         # initialize
205         loss, dx = softmax_loss(scores, y)
206         loss += 0.5 * self.reg * (np.sum(np.square(self.params['W' +
str(self.num_layers)])))
207         dx_back, dw_back, db_back = affine_backward(dx, final_cache)
208         grads['W' + str(self.num_layers)] = dw_back + self.reg * self.params['W'
+ str(self.num_layers)]
209         grads['b' + str(self.num_layers)] = db_back
210
211         # go backward all layers and update weights, bias, gammas and betas
212         for i in range(self.num_layers - 1, 0, -1):
213             # dropout layer
214             if self.use_dropout:
215                 dx_back = dropout_backward(dx_back, dropout_cache[str(i)])
216                 dx_relu = relu_backward(dx_back, relu_cache[str(i)])
217
218             # batchnorm layer
219             affine_backward_input = dx_relu
220             if self.use_batchnorm:
221                 dx_bn, dgamma, dbeta = batchnorm_backward(dx_relu,
batchnorm_cache[str(i)])
222                 grads['gamma' + str(i)] = dgamma
223                 grads['beta' + str(i)] = dbeta
224                 affine_backward_input = dx_bn
225                 dx_back, dw_back, db_back = affine_backward(affine_backward_input,
fc_cache[str(i)])
226
227                 grads['W' + str(i)] = dw_back + self.reg * self.params['W' + str(i)]
228                 grads['b' + str(i)] = db_back

```

```
229         loss += 0.5 * self.reg * (np.sum(np.square(self.params['W' + str(i)])))
230
231     # ===== #
232     # END YOUR CODE HERE
233     # ===== #
234
235     return loss, grads
236
```



```

1 import numpy as np
2
3 """
4 This code was originally written for CS 231n at Stanford University
5 (cs231n.stanford.edu). It has been modified in various areas for use in the
6 ECE 239AS class at UCLA. This includes the descriptions of what code to
7 implement as well as some slight potential changes in variable names to be
8 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung
9 for
10 permission to use this code. To see the original version, please visit
11 cs231n.stanford.edu.
12 """
13
14 """
15 This file implements various first-order update rules that are commonly used
16 for
17 training neural networks. Each update rule accepts current weights and the
18 gradient of the loss with respect to those weights and produces the next set
19 of
20 weights. Each update rule has the same interface:
21
22 def update(w, dw, config=None):
23
24     Inputs:
25     - w: A numpy array giving the current weights.
26     - dw: A numpy array of the same shape as w giving the gradient of the
27         loss with respect to w.
28     - config: A dictionary containing hyperparameter values such as learning
29         rate,
30         momentum, etc. If the update rule requires caching values over many
31         iterations, then config will also hold these cached values.
32
33     Returns:
34     - next_w: The next point after the update.
35     - config: The config dictionary to be passed to the next iteration of the
36         update rule.
37
38     NOTE: For most update rules, the default learning rate will probably not
39     perform
40     well; however the default values of the other hyperparameters should work
41     well
42     for a variety of different problems.
43
44     For efficiency, update rules may perform in-place updates, mutating w and
45     setting next_w equal to w.
46 """
47
48 def sgd(w, dw, config=None):
49     """
50     Performs vanilla stochastic gradient descent.
51
52     config format:
53     - learning_rate: Scalar learning rate.
54     """
55     if config is None: config = {}
56     config.setdefault('learning_rate', 1e-2)
57
58     w -= config['learning_rate'] * dw
59     return w, config
60
61 def sgd_momentum(w, dw, config=None):
62     """
63     Performs stochastic gradient descent with momentum.
64
65     config format:
66     - learning_rate: Scalar learning rate.
67     - momentum: Scalar between 0 and 1 giving the momentum value.
68         Setting momentum = 0 reduces to sgd.
69     - velocity: A numpy array of the same shape as w and dw used to store a
70         moving
71         average of the gradients.
72     """
73     if config is None: config = {}
74     config.setdefault('learning_rate', 1e-2)
75     config.setdefault('momentum', 0.9) # set momentum to 0.9 if it wasn't there
76     v = config.get('velocity', np.zeros_like(w)) # gets velocity, else sets
77     it to zero.
78
79     # ===== #
80     # YOUR CODE HERE:
81     # Implement the momentum update formula. Return the updated weights
82     # as next_w, and store the updated velocity as v.

```

```

77 # ===== #
78
79 v = config['momentum'] * v - config['learning_rate'] * dw
80 next_w = w + v
81
82 # ===== #
83 # END YOUR CODE HERE
84 # ===== #
85
86 config['velocity'] = v
87
88 return next_w, config
89
90 def sgd_nesterov_momentum(w, dw, config=None):
91     """
92     Performs stochastic gradient descent with Nesterov momentum.
93
94     config format:
95     - learning_rate: Scalar learning rate.
96     - momentum: Scalar between 0 and 1 giving the momentum value.
97       Setting momentum = 0 reduces to sgd.
98     - velocity: A numpy array of the same shape as w and dw used to store a
99       moving
100       average of the gradients.
101     """
102     if config is None: config = {}
103     config.setdefault('learning_rate', 1e-2)
104     config.setdefault('momentum', 0.9) # set momentum to 0.9 if it wasn't there
105     v = config.get('velocity', np.zeros_like(w)) # gets velocity, else sets
106     it to zero.
107
108     # ===== #
109     # YOUR CODE HERE:
110     # Implement the momentum update formula. Return the updated weights
111     # as next_w, and store the updated velocity as v.
112     # ===== #
113
114     v_old = v
115     v = config['momentum'] * v - config['learning_rate'] * dw
116     w += v + config['momentum'] * (v - v_old)
117     next_w = w
118
119     # ===== #
120     # END YOUR CODE HERE
121     # ===== #
122
123     config['velocity'] = v
124
125     return next_w, config
126
127 def rmsprop(w, dw, config=None):
128     """
129     Uses the RMSProp update rule, which uses a moving average of squared
130     gradient
131     values to set adaptive per-parameter learning rates.
132
133     config format:
134     - learning_rate: Scalar learning rate.
135     - decay_rate: Scalar between 0 and 1 giving the decay rate for the squared
136       gradient cache.
137     - epsilon: Small scalar used for smoothing to avoid dividing by zero.
138     - beta: Moving average of second moments of gradients.
139     """
140     if config is None: config = {}
141     config.setdefault('learning_rate', 1e-2)
142     config.setdefault('decay_rate', 0.99)
143     config.setdefault('epsilon', 1e-8)
144     config.setdefault('a', np.zeros_like(w))
145
146     next_w = None
147
148     # ===== #
149     # YOUR CODE HERE:
150     # Implement RMSProp. Store the next value of w as next_w. You need
151     # to also store in config['a'] the moving average of the second
152     # moment gradients, so they can be used for future gradients. Concretely,
153     # config['a'] corresponds to "a" in the lecture notes.
154     # ===== #
155
156     config['a'] = config['decay_rate'] * config['a'] + (1 -
157 config['decay_rate']) * (dw**2)
158     next_w = w - config['learning_rate'] * dw / (np.sqrt(config['a']) +
159 config['epsilon'])

```

```

156 # ===== #
157 # END YOUR CODE HERE
158 # ===== #
159
160 return next_w, config
161
162
163 def adam(w, dw, config=None):
164     """
165     Uses the Adam update rule, which incorporates moving averages of both the
166     gradient and its square and a bias correction term.
167
168     config format:
169     - learning_rate: Scalar learning rate.
170     - beta1: Decay rate for moving average of first moment of gradient.
171     - beta2: Decay rate for moving average of second moment of gradient.
172     - epsilon: Small scalar used for smoothing to avoid dividing by zero.
173     - m: Moving average of gradient.
174     - v: Moving average of squared gradient.
175     - t: Iteration number.
176     """
177     if config is None: config = {}
178     config.setdefault('learning_rate', 1e-3)
179     config.setdefault('beta1', 0.9)
180     config.setdefault('beta2', 0.999)
181     config.setdefault('epsilon', 1e-8)
182     config.setdefault('v', np.zeros_like(w))
183     config.setdefault('a', np.zeros_like(w))
184     config.setdefault('t', 0)
185
186     next_w = None
187
188     # ===== #
189     # YOUR CODE HERE:
190     # Implement Adam. Store the next value of w as next_w. You need
191     # to also store in config['a'] the moving average of the second
192     # moment gradients, and in config['v'] the moving average of the
193     # first moments. Finally, store in config['t'] the increasing time.
194     # ===== #
195
196     beta1 = config['beta1']
197     beta2 = config['beta2']
198     t = config['t'] + 1
199
200     v = beta1 * config['v'] + (1 - beta1) * dw
201     a = beta2 * config['a'] + (1 - beta2) * (dw**2)
202     v_corrected = v / (1 - beta1**t)
203     a_corrected = a / (1 - beta2**t)
204     next_w = w - config['learning_rate'] * v_corrected / (np.sqrt(a_corrected)
+ config['epsilon'])
205
206     config['v'] = v
207     config['a'] = a
208     config['t'] = t
209
210     # ===== #
211     # END YOUR CODE HERE
212     # ===== #
213
214     return next_w, config
215
216
217
218
219
220

```